

Exploiting CVE-2018-8611

Windows Kernel Transaction Manager (KTM) Race Condition

Cedric Halbronn - OffensiveCon 2020



About

- Exploit Development Group (EDG), NCC Group
- Write exploits to help consultants do their job
- Focus on patched vulnerabilities

Cedric Halbronn

- Presenting
- @saidelike, cedric.halbronn@nccgroup.com

Aaron Adams

- Co-researcher, unable to attend
- @fidgetingbits, aaron.adams@nccgroup.com

This talk

- Discuss an interesting race condition affecting Microsoft Kernel Transaction Manager (KTM)
 - Found used in the wild by [Kaspersky](#)
 - Boris Larin (@Oct0xor) and Igor Soumenkov (@2igosha)
- Exploited by us early 2019
 - Never got to see the original exploit or details
- Minimal details from Kaspersky at the time
 - Race condition in KTM
 - Exploitable from inside browser sandbox
 - Works on Windows 10
 - A few hints for triggering the race
- Presented by Aaron at POC2019 in November

Notable KTM-related security findings

- 2010 - [CVE-2010-1889](#) - Tavis Ormandy - invalid free
- 2015 - [MS15-038](#) - James Forshaw - type confusion
- 2017 - [CVE-2017-8481](#) - j00ru - stack memory disclosure
- 2018 - [CVE-2018-8611](#) - Oct0xor/2igosha - Kaspersky blog
- 2019 - [Proton Bot malware uses KTM](#)
 - Used transacted versions of common functions to evade API inspection

Oct0xor/2igosha - Kaspersky blog

To abuse this vulnerability exploit first creates a named pipe and opens it for read and write. Then it creates a pair of new transaction manager objects, resource manager objects, transaction objects and creates a big number of enlistment objects for what we will call "Transaction #2". Enlistment is a special object that is used for association between a transaction and a resource manager. When the transaction state changes associated resource manager is notified by the KTM. After that it creates one more enlistment object only now it does so for "Transaction #1" and commits all the changes made during this transaction.

After all the initial preparations have been made exploit proceeds to the second part of vulnerability trigger. It creates multiple threads and binds them to a single CPU core. One of created threads calls `NtQueryInformationResourceManager` in a loop, while second thread tries to execute `NtRecoverResourceManager` once. But the vulnerability itself is triggered in the third thread. This thread uses a trick of execution `NtQueryInformationThread` to obtain information on the latest executed syscall for the second thread. Successful execution of `NtRecoverResourceManager` will mean that race condition has occurred and further execution of `WriteFile` on previously created named pipe will lead to memory corruption.

```
rax=0000000000000000 rbx=fffffe10e6bfb7dd8 rcx=fffff8016ffd2180
rdx=0000000000000000 rsi=fffffe10e6bfb7db0 rdi=4141414141414141
rip=fffff80b8a9d4b40 rsp=fffffbf0046229a00 rbp=fffffbf0046229b80
r8=fffffe10e68064080 r9=fffffbf0046229760 r10=fffffe10e6c11f308
r11=0000000000000000 r12=fffffe10e6bfb7ec0 r13=0000000000000100
r14=41414141414140b9 r15=0000000000000000
iopl=0 nv up ei pl nz na po cy
cs=0010 ss=0018 ds=002b es=002b fs=0053 gs=002b efl=00010207
tm!TmRecoverResourceManagerExt+0x100:
fffff80b`8a9d4b40 418b86ac000000 mov    eax,dword ptr [r14+0ACh] ds:002b:41414141`41414165=????????
```

Proof of concept: execution of WriteFile with buffer set to 0x41

Me when reading write-up

How to draw an owl

1.



2.



1. Draw some circles

2. Draw the rest of the fucking owl

Tooling

- Virtualization: [VMWare Workstation](#)
- Binary analysis: [IDA Pro](#), [Hex-Rays Decompiler](#)
- Binary diffing: [Diaphora](#)
- Collaboration: [IDArling](#)
- Debugging: [WinDbg](#) (ring0), [virtualkd](#), [x64dbg](#) (ring3)
- Additional plugins/tools: [ret-sync](#), [HexRaysPyTools](#)
- Structure analysis: [Vergilius Project](#), [ReactOS](#) source
- Syscall numbers: [windows-syscalls](#) by j00ru
- Slides: [Remarkjs](#)

Diffing - functions

Line	Address	Name	Address 2	Name 2	Ratio	BBlocks 1	BBlocks 2	Description
00015	140452b70	LpcRequestWaitReplyPortEx	140452b60	LpcRequestWaitReplyPortEx	0.890	1	1	Perfect match, same name
00009	1403dabd0	PfpRepurposeNameLoggingTrace	1403da8a0	PfpRepurposeNameLoggingTrace	0.890	1	1	Perfect match, same name
00019	140561540	PnpWaitForDevicesToStart	140562540	PnpWaitForDevicesToStart	0.880	1	1	Perfect match, same name
00014	140432b40	LpcRequestWaitReplyPort	140432850	LpcRequestWaitReplyPort	0.860	1	1	Perfect match, same name
00003	14033da80	TmCommitComplete	14033d760	TmCommitComplete	0.860	1	1	Perfect match, same name
00002	14033da14	TmPrepareComplete	14033d6f4	TmPrepareComplete	0.860	1	1	Perfect match, same name
00004	14033fe14	TmReadOnlyEnlistment	14033fa84	TmReadOnlyEnlistment	0.770	1	1	Perfect match, same name
00026	140574680	TmpEnlistmentInitialization	140575680	TmpEnlistmentInitialization	0.680	1	1	Perfect match, same name
00025	1405745b0	TmpTransactionManagerInitialization	1405755b0	TmpTransactionManagerInitialization	0.670	1	1	Perfect match, same name
00012	1403ead50	TmpFindTransactionManager	1403eaa20	TmpFindTransactionManager	0.670	1	1	Perfect match, same name
00000	14030b5c4	ObInsertObject	14030b5d4	ObInsertObject	0.670	1	1	Perfect match, same name
00013	14042ebf0	TmRollbackComplete	14042e900	TmRollbackComplete	0.610	1	1	Perfect match, same name
00001	140321998	TmRecoverResourceManager	140474940	TmRecoverResourceManager	0.610	38	39	Perfect match, same name
00028	14050bd20	VerifierExEnterPriorityRegionAndAcquir...	14050cd20	VerifierExEnterCriticalRegionAndAcquir...	0.500	1	1	Nodes, edges, complexity and mnemonics with small differences

Diffing - assembly

Diff assembler TmRecoverResourceManager - TmRecoverResourceManager

```
66 mov    [rsp+0B8h+var_68], rdi
67 test   byte ptr [rsi+0ACh], 4
68 jz    short loc_140321A50
69loc_140321a5e:
70 jmp    short loc_140321A24
71loc_140321a50:
72 lea    rbx, [rsi+40h]
73 and   [rsp+0B8h+TimeOut], 0
74 xor    r9d, r9d; Alertable
75 xor    r8d, r8d; WaitMode
76 xor    edx, edx; WaitReason
77 mov    rcx, rbx; Object
78 call   KeWaitForSingleObject
79 bts   dword ptr [rsi+0ACh], 7
80 xor    edx, edx; Wait
81 mov    rcx, rbx; Mutex
82 call   KeReleaseMutex
83 jmp    loc_140321A24
84loc_140321a51:
```



```
46 test   byte ptr [r12+0ACh], 4
47 jz    short loc_1404745E7
48loc_1404745E8:
49 jmp    short loc_1404745C6
50loc_1404745E7:
51 and   [rsp+0B8h+TimeOut], 0
52 xor    r9d, r9d; Alertable
53 xor    r8d, r8d; WaitMode
54 xor    edx, edx; WaitReason
55 lea    rcx, [r12+40h]; Object
56 call   KeWaitForSingleObject
57 bts   dword ptr [r12+0ACh], 7
58 xor    edx, edx; Wait
59 lea    rcx, [r12+40h]; Mutex
60 call   KeReleaseMutex
61 jmp    short loc_1404745C6
62loc_140474a17:
63 mov    rbx, [r13+0]
64 mov    [rsp+0B8h+var_70], rbx
65 mov    r14d, dword ptr [rsp+0B8h+undefined_value]
66loc_140474a25:
67 cmp    rbx, r13
68 jz    loc_140474BCA
69loc_140474a26:
70 test   byte ptr [rbx+24h], 4
71 jz    short loc_140474A3E
72loc_140474a34:
73 mov    rbx, [rbx]
74 mov    [rsp+0B8h+var_70], rbx
75 jmp    short loc_140474a25
76loc_140474a36:
77 lea    rcx, [rbx-88h]; Object
78 call   ObReferenceObject
79 and   [rsp+0B8h+TimeOut], 0
80 xor    r9d, r9d; Alertable
81 xor    r8d, r8d; WaitMode
82 xor    edx, edx; WaitReason
83 lea    rcx, [rdi+48h]; Object
84 call   KeWaitForSingleObject
85 xor    bl, bl
86 mov    byte ptr [rsp+0B8h+dwEnlistmentFlag_4_bit_], bl
87 mov    ecx, [rdi+24h]
88 test   cl, cl
89 jns    short loc_140321B73
90loc_140321b04:
91 mov    r8d, 1
92 and   ecx, r8d
93 jz    short loc_140321B37
94loc_140321bf0:
95 mov    rax, [rdi+18h]
96 mov    edx, [rax+0C0h]
97 cmp    edx, 3
98 jz    short loc_140321B23
99loc_140321b21:
100 cmp   edx, 4
101 jns    short loc_140321B37
102 mov    bl, r8b
103 mov    byte ptr [rsp+0B8h+dwEnlistmentFlag_4_bit_], bl
104 mov    r15d, 800h
105 mov    r15d, 800h
106 jmp    short loc_140321B6E
107loc_140321b57:
108 test   ecx, ecx
109 jns    short loc_140321B48
110loc_140321b5b:
111 mov    rax, [rdi+18h]
112 cmp    dword ptr [rax+0C0h], 5
113 jz    short loc_140321B5C
114loc_140321b54:
```

Diffing - Hex-Rays pre-cleanup

```
91         v17);
92     v15 = v18;
93     if ( *(_BYTE *) (v9 + 172) & 4 )
94     v15 = 1;
95     v18 = v15;
96     ObfDereferenceObject(v7 - 17);
97     KeWaitForSingleObject((char *)v1 + 40, Executive, 0, 0, 0i64);
98     if ( *(_DWORD *)v1 + 7 ) != 2 )
99     goto LABEL_34;
100    v2 = v18;

101 }
102 else
103 {
104     ObfDereferenceObject(v7 - 17);
105 }
106 if ( v2 )
107 {
108     v7 = (_QWORD *) * (_QWORD *)v1 + 34;
109     v2 = 0;
110     v18 = 0;
111 }
112 else
113 {
114LABEL_12:

88         v16);
89     ObfDereferenceObject(v6 - 17);
90     KeWaitForSingleObject((char *)v1 + 40, Executive, 0, 0, 0i64);
91     if ( *(_DWORD *)v1 + 7 ) != 2 )
92     goto LABEL_32;
93     v14 = *(_QWORD *)v1 + 45;
94     if ( !v14 || *(_DWORD *) (v14 + 64) != 3 )
95     goto LABEL_31;
96     v6 = (_QWORD *) * (_QWORD *)v1 + 34;
97 }
98 else
99 {
100     ObfDereferenceObject(v6 - 17);

101LABEL_12:
```

Agenda

- What is KTM?
- Patch analysis
- Triggering the bug
- Finding a write primitive
- Building a read primitive
- Privilege escalation
- Recent bonus info

Windows Kernel Transaction Manager (KTM)

KTM - What is it?

- MSDN documentation
 - [KTM Portal](#)

The screenshot shows the left sidebar of the MSDN KTM Portal. At the top is a search bar labeled "Filter by title". Below it is a main navigation item "Kernel Transaction Manager" which is highlighted with a grey background. Underneath this are two sections: "About KTM" and "KTM Reference". "About KTM" contains links for "About KTM", "What is a Transaction?", "Working With Transactions", "Writing a Resource Manager", "Interoperability With Other Windows", "Features", and "KTM Security and Access Rights". "KTM Reference" contains links for "KTM Reference", "Kernel Transaction Manager Enumerations", "Kernel Transaction Manager Functions", "Kernel Transaction Manager Structures", and "Kernel Transaction Manager Constants".

Kernel Transaction Manager

05/31/2018 • 2 minutes to read •

Purpose

The Kernel Transaction Manager (KTM) enables the development of applications that use transactions. The transaction engine itself is within the kernel, but transactions can be developed for kernel- or user-mode transactions, and within a single host or among distributed hosts.

The KTM is used to implement Transactional NTFS (TxF) and Transactional Registry (TxR). TxF allows transacted file system operations within the NTFS file system. TxR allows transacted registry operations. KTM enables client applications to coordinate file system and registry operations with a transaction.

To develop an application that coordinates transactions with resources other than TxF or TxR, you must first develop a Win32 transaction-aware service, also called a resource manager.

Managed and COM+ applications should use their native transaction managers.

Where applicable

KTM can be used with applications and resource managers hosted on Windows Vista or Windows Server 2008.

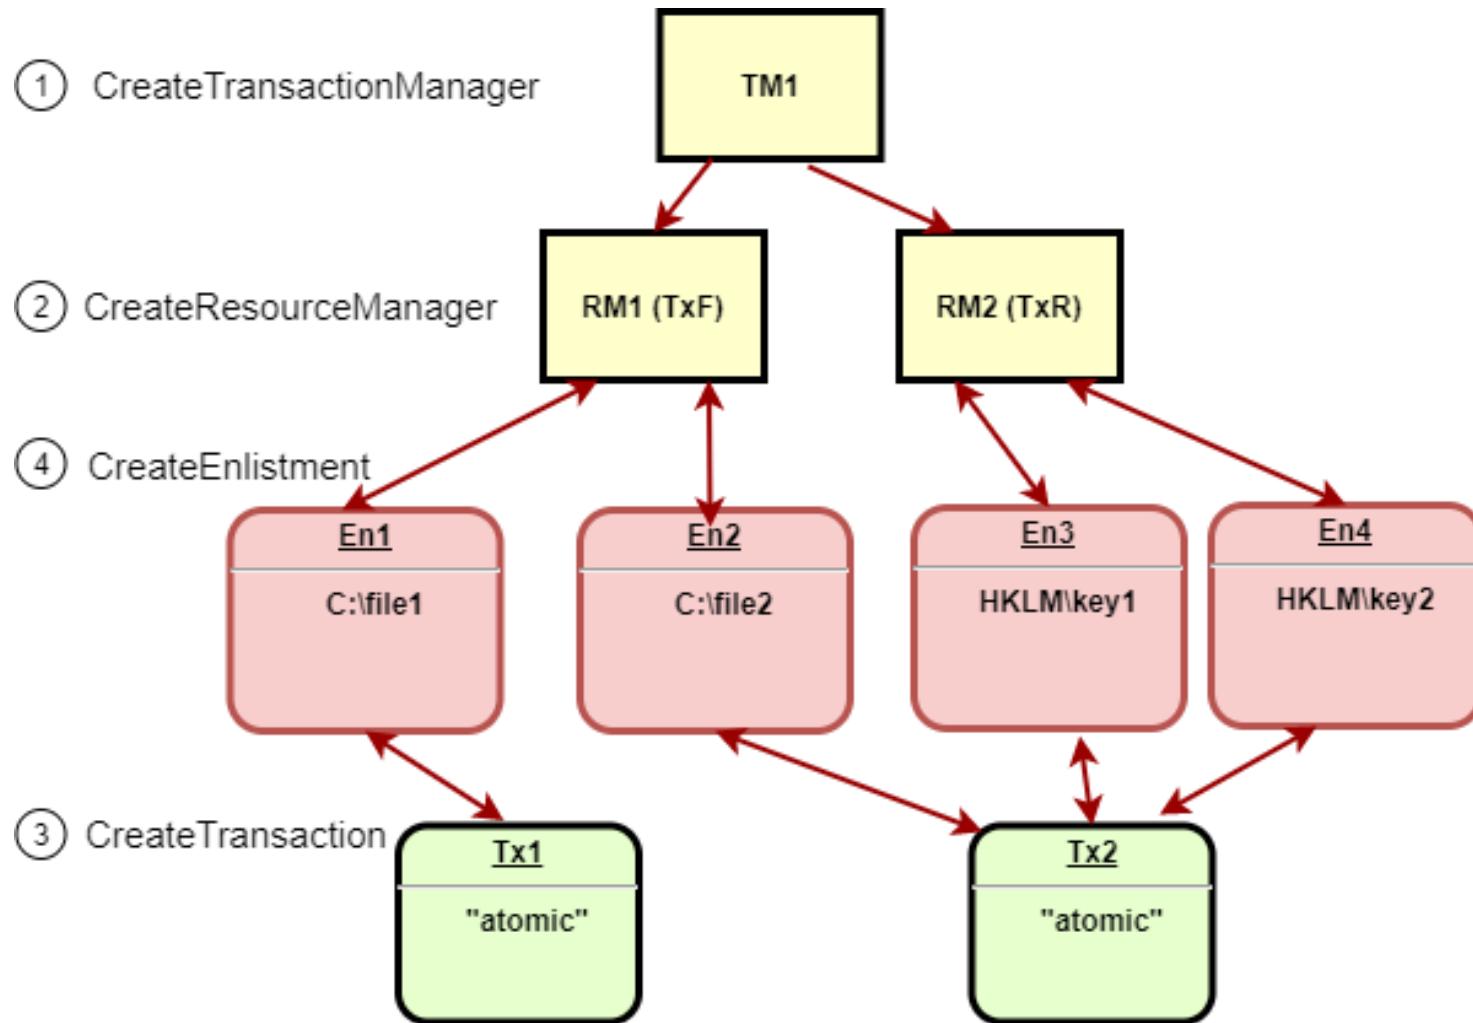
KTM - What is it?

- Kernel service added in Windows Vista (~2006)
 - Windows 7 and earlier: ntoskrnl.exe
 - Windows 8 and later: tm.sys
- Provide "ACID" functionality: atomic, consistent, isolated, and durable
- KTM service used by two major Windows components
 - Transactional Registry
 - Transactional NTFS
- A few dozen APIs/system calls exposed to userland

Important objects

- KTM service has **4** fundamental kernel objects
 - All referenced counted objects created by `ObCreateObject()`
- **Transaction Manager (TM)**
 - Manages a log of transactions associated with one or more resource managers
- **Resource Manager (RM)**
 - Manages enlistments related to a specific managed resource doing work for a Transaction
- **Transaction (Tx)**
 - Tracks a series of sub actions making up a single atomic operation
- **Enlistment (En)**
 - Some code responsible for doing work related to a Transaction

Important objects



Transaction Manager (TM)

- Created using [CreateTransactionManager\(\)](#).
 - Usually first to exist

```
HANDLE CreateTransactionManager(  
    IN LPSECURITY_ATTRIBUTES lpTransactionAttributes,  
    LPWSTR             LogFileName,  
    IN ULONG            CreateOptions,  
    IN ULONG            CommitStrength  
) ;
```

- Allocates a [KTM](#) structure on the non-paged pool
 - TmTm pool tag
- A resource manager must be associated with some TM
- Optional log for transactions
 - A **volatile** TM is one that uses no log file
 - Set TRANSACTION_MANAGER_VOLATILE flag in CreateOptions parameter
 - Logs have limited size - problematic for exploitation

_KTM

- Most fields omitted

```
//0x3c0 bytes (sizeof)
struct _KTM
{
    ULONG cookie;                                //0x0
    struct _KMUTANT Mutex;                      //0x8
    enum KTM_STATE State;                        //0x40
    [...]
    ULONG Flags;                                 //0x80
    [...]
    struct _KRESOURCEMANAGER* TmRm;            //0x2a8
    [...]
};
```

Resource Manager (RM)

- Created using [CreateResourceManager\(\)](#).

```
HANDLE CreateResourceManager(
    IN LPSECURITY_ATTRIBUTES lpResourceManagerAttributes,
    IN LPGUID             ResourceManagerId,
    IN DWORD              CreateOptions,
    IN HANDLE             TmHandle,
    LPWSTR                Description
);
```

- Must be passed a TM handle
- Optional Description parameter
- Allocates a [KRESOURCEMANAGER](#) structure on the non-paged pool
 - TmRm pool tag

_KRESOURCEMANAGER

```
//0x250 bytes (sizeof)
struct _KRESOURCEMANAGER
{
    struct _KEVENT NotificationAvailable;                                //0x0
    ULONG cookie;                                                       //0x18
    enum _KRESOURCEMANAGER_STATE State;                                 //0x1c
    ULONG Flags;                                                       //0x20
    struct _KMUTANT Mutex;                                              //0x28
    [...]
    struct _KQUEUE NotificationQueue;                                    //0x98
    struct _KMUTANT NotificationMutex;                                  //0xd8
    struct _LIST_ENTRY EnlistmentHead;                                 //0x110
    ULONG EnlistmentCount;                                            //0x120
    LONG (*NotificationRoutine)(struct _KENLISTMENT* arg1, VOID* arg2, VOID* arg3,
                               ULONG arg4, union _LARGE_INTEGER* arg5, ULONG arg6, VOID* arg7); //0x128
    [...]
    struct _KTM* Tm;                                                 //0x168
    struct _UNICODE_STRING Description;                                //0x170
    [...]
};
```

_KRESOURCEMANAGER fields

- Tm - Pointer to the associated transaction manager
- Description - Unicode description of resource manager
- Mutex - Locks RM. Other code cannot
 - Parse the resource manager's enlistments list
 - Read Description
 - etc.
- EnlistmentHead - List of associated enlistments with resource manager
- NotificationQueue - Notification events
 - Queried from ring3 to read enlistment state change events

Transaction (Tx)

- Created using [CreateTransaction\(\)](#) function

```
HANDLE CreateTransaction(
    IN LPSECURITY_ATTRIBUTES lpTransactionAttributes,
    IN LPGUID             UOW,
    IN DWORD               CreateOptions,
    IN DWORD               IsolationLevel,
    IN DWORD               IsolationFlags,
    IN DWORD               Timeout,
    LPWSTR                Description
);
```

- Creates a [KTRANSACTION](#) structure on the non-paged pool using
 - TmTx pool tag
- Represents whole piece of work to be done
- Resource managers enlist in this transaction to complete the work

_KTRANSACTION

```
//0x2d8 bytes (sizeof)
struct _KTRANSACTION
{
    struct _KEVENT OutcomeEvent;                                //0x0
    ULONG cookie;                                              //0x18
    struct _KMUTANT Mutex;                                    //0x20
    [...]
    struct _GUID UOW;                                         //0xb0
    enum _KTRANSACTION_STATE State;                           //0xc0
    ULONG Flags;                                               //0xc4
    struct _LIST_ENTRY EnlistmentHead;                         //0xc8
    ULONG EnlistmentCount;                                     //0xd8
    [...]
    union _LARGE_INTEGER Timeout;                            //0x128
    struct _UNICODE_STRING Description;                      //0x130
    [...]
    struct _KTM* Tm;                                         //0x200
    [...]
};
```

Enlistments (En)

- Created using [CreateEnlistment\(\)](#).

```
hEn = CreateEnlistment(
    NULL,      // lpEnlistmentAttributes
    hRM,       // ResourceManagerHandle - Existing resource manager handle
    hTx,       // TransactionHandle - Existing transaction handle
    0x39ffff0f, // NotificationMask - Special value to receive all possible notifications
    0,         // CreateOptions
    NULL       // EnlistmentKey
);
```

- Allocates a [**KENLISTMENT**](#) structure on the non-paged pool
 - TmEn pool tag
- Each has an assigned GUID
- Must be associated with both a resource manager and a transaction
- Typically a transaction will have multiple enlistments

_KENLISTMENT

```
//0x1e0 bytes (sizeof)
struct _KENLISTMENT
{
    ULONG cookie;                                //0x0
    struct _KTMOBJECT_NAMESPACE_LINK NamespaceLink; //0x8
    struct _GUID EnlistmentId;                  //0x30
    struct _KMUTANT Mutex;                      //0x40
    struct _LIST_ENTRY NextSameTx;               //0x78
    struct _LIST_ENTRY NextSameRm;                //0x88
    struct _KRESOURCEMANAGER* ResourceManager;    //0x98
    struct _KTRANSACTION* Transaction;           //0xa0
    enum _KENLISTMENT_STATE State;               //0xa8
    ULONG Flags;                                 //0xac
    ULONG NotificationMask;                     //0xb0
    [...]
};
```

_KENLISTMENT fields of interest

- Transaction - The transaction that the enlistment is actually doing work for
- Flags - Indicates the type and state of the enlistment
- Mutex - Locks the enlistment and prevents other code from manipulating it
- State - The current state of the enlistment in relation to the transaction
- NotificationMask - Which notifications should be queued to the resource manager related to this enlistment
- NextSameRm - A linked list of enlistments associated with the same resource manager
 - This is the list entry whose head is _KRESOURCEMANAGER.EnlistmentHead

_KENLISTMENT flags

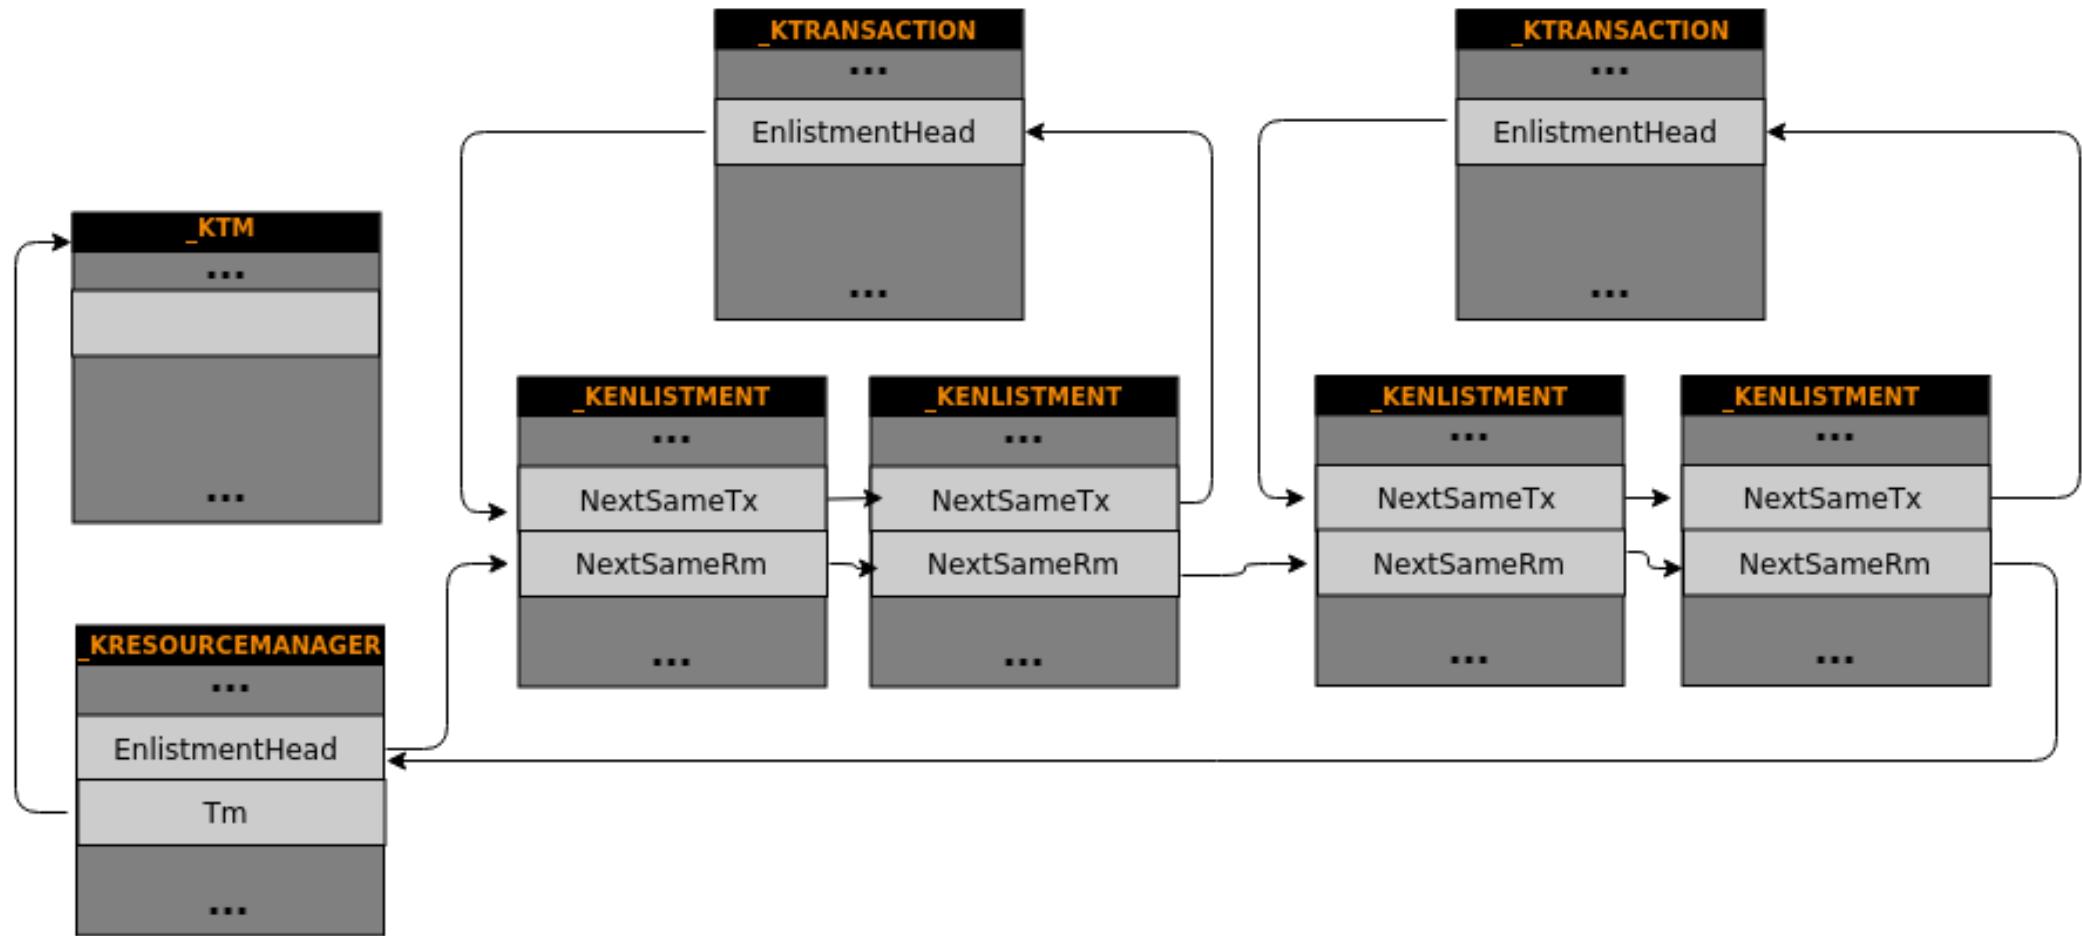
- The Flags field uses undocumented flags

```
enum KENLISTMENT_FLAGS {  
    KENLISTMENT_SUPERIOR          = 0x01,  
    KENLISTMENT_RECOVERABLE       = 0x02,  
    KENLISTMENT_FINALIZED         = 0x04,  
    KENLISTMENT_FINAL_NOTIFICATION = 0x08,  
    KENLISTMENT_OUTCOME_REQUIRED  = 0x10,  
    KENLISTMENT_HAS_SUPERIOR_SUB  = 0x20,  
    KENLISTMENT_IS_NOTIFIABLE     = 0x80,  
    KENLISTMENT_DELETED           = 0x80000000  
};
```

How to finalize and free an enlistment?

- Enlistments are a reference counted object
- Call some code path that triggers `TmpFinalizeEnlistment()` to lower ref count
 - A Prepared enlistment upon moving to Committed state will be finalized
 - Use `CommitComplete()` function on enlistment handle
- Then `CloseHandle()` to remove our final userland reference
- Either frees immediately, or upon any other KTM kernel code doing final dereference

Structure relationship overview



Transaction and Enlistment States

- Transaction not complete until all enlistments have committed
- Transaction cannot be committed until all of enlistments transition through a series of synchronized states
- A transaction with only one enlistment is the exception
- Typical state transitions

PrePreparing -> PrePrepared -> Preparing -> Prepared -> Committed

_KENLISTMENT_STATE

Documented

```
enum _KENLISTMENT_STATE
{
    //...
    KEnlistmentPreparing = 257,
    KEnlistmentPrepared = 258,
    KEnlistmentCommitted = 260,
    //...
    KEnlistmentPreparing = 257,
    //...
    KEnlistmentPrePreparing = 266,
    //...
    KEnlistmentPrePrepared = 273,
};
```

Notifications

- Dictated by enlistment NotificationMask option at creation
- Each RM has a set of associated Tx notifications that occur on milestone events, such as an En switching from one state to another
- Notifications can be read using [GetNotificationResourceManager\(\)](#).
- The events are queued/retrieved using FIFO

```
BOOL GetNotificationResourceManager(
    IN HANDLE             ResourceManagerHandle,
    OUT PTRANSACTION_NOTIFICATION TransactionNotification,
    IN ULONG              NotificationLength,
    IN DWORD              dwMilliseconds,
    OUT PULONG             ReturnLength
);
```

- TRANSACTION_NOTIFICATION struct contains a TRANSACTION_NOTIFICATION_RECOVERY_ARGUMENT
 - Tells us which En a notification is associated with

Recovery

- If a Tx fails or is interrupted for whatever reason, it can be possible to recover
- Recovery in part possible by calling RecoverResourceManager()

```
BOOL RecoverResourceManager(  
    IN HANDLE ResourceManagerHandle  
);
```

- During this recovery phase, each enlistment associated with transactions in specific states will receive a notification
- Allows the enlisted workers to synchronize on what they were doing for the transaction

Understanding CVE-2018-8611

Diffing - functions

Line	Address	Name	Address 2	Name 2	Ratio	BBlocks 1	BBlocks 2	Description
00015	140452b70	LpcRequestWaitReplyPortEx	140452b60	LpcRequestWaitReplyPortEx	0.890	1	1	Perfect match, same name
00009	1403dabd0	PfpRepurposeNameLoggingTrace	1403da8a0	PfpRepurposeNameLoggingTrace	0.890	1	1	Perfect match, same name
00019	140561540	PnpWaitForDevicesToStart	140562540	PnpWaitForDevicesToStart	0.880	1	1	Perfect match, same name
00014	140432b40	LpcRequestWaitReplyPort	140432850	LpcRequestWaitReplyPort	0.860	1	1	Perfect match, same name
00003	14033da80	TmCommitComplete	14033d760	TmCommitComplete	0.860	1	1	Perfect match, same name
00002	14033da14	TmPrepareComplete	14033d6f4	TmPrepareComplete	0.860	1	1	Perfect match, same name
00004	14033fe14	TmReadOnlyEnlistment	14033fa84	TmReadOnlyEnlistment	0.770	1	1	Perfect match, same name
00026	140574680	TmpEnlistmentInitialization	140575680	TmpEnlistmentInitialization	0.680	1	1	Perfect match, same name
00025	1405745b0	TmpTransactionManagerInitialization	1405755b0	TmpTransactionManagerInitialization	0.670	1	1	Perfect match, same name
00012	1403ead50	TmpFindTransactionManager	1403eaa20	TmpFindTransactionManager	0.670	1	1	Perfect match, same name
00000	14030b5c4	ObInsertObject	14030b5d4	ObInsertObject	0.670	1	1	Perfect match, same name
00013	14042ebf0	TmRollbackComplete	14042e900	TmRollbackComplete	0.610	1	1	Perfect match, same name
00001	140321998	TmRecoverResourceManager	140474940	TmRecoverResourceManager	0.610	38	39	Perfect match, same name
00028	14050bd20	VerifierExEnterPriorityRegionAndAcquir...	14050cd20	VerifierExEnterCriticalRegionAndAcquir...	0.500	1	1	Nodes, edges, complexity and mnemonics with small differ...

Diffing - assembly

Diff assembler TmRecoverResourceManager - TmRecoverResourceManager

```
66 mov    [rsp+0B8h+var_68], rdi
67 test   byte ptr [rsi+0ACh], 4
68 jz    short loc_140321A50
69loc_140321a5e:
70 jmp    short loc_140321A24
71loc_140321a50:
72 lea    rbx, [rsi+40h]
73 and   [rsp+0B8h+TimeOut], 0
74 xor    r9d, r9d; Alertable
75 xor    r8d, r8d; WaitMode
76 xor    edx, edx; WaitReason
77 mov    rcx, rbx; Object
78 call   KeWaitForSingleObject
79 bts   dword ptr [rsi+0ACh], 7
80 xor    edx, edx; Wait
81 mov    rcx, rbx; Mutex
82 call   KeReleaseMutex
83 jmp    loc_140321A24
84loc_140321a51:
```



```
46 test   byte ptr [r12+0ACh], 4
47 jz    short loc_1404745E7
48loc_1404745E8:
49 jmp    short loc_1404745C6
50loc_1404745E7:
51 and   [rsp+0B8h+TimeOut], 0
52 xor    r9d, r9d; Alertable
53 xor    r8d, r8d; WaitMode
54 xor    edx, edx; WaitReason
55 lea    rcx, [r12+40h]; Object
56 call   KeWaitForSingleObject
57 bts   dword ptr [r12+0ACh], 7
58 xor    edx, edx; Wait
59 lea    rcx, [r12+40h]; Mutex
60 call   KeReleaseMutex
61 jmp    short loc_1404745C6
62loc_140474417:
63 mov    rbx, [r13+0]
64 mov    [rsp+0B8h+var_70], rbx
65 mov    r14d, dword ptr [rsp+0B8h+undefined_value]
66loc_140474423:
67 cmp    rbx, r13
68 jz    loc_140474BC4
69loc_14047442E:
70 test   byte ptr [rbx+24h], 4
71 jz    short loc_140474A3E
72loc_140474434:
73 mov    rbx, [rbx]
74 mov    [rsp+0B8h+var_70], rbx
75 jmp    short loc_140474423
76loc_14047443E:
77 lea    rcx, [rbx-88h]; Object
78 call   ObReferenceObject
79 and   [rsp+0B8h+TimeOut], 0
80 xor    r9d, r9d; Alertable
81 xor    r8d, r8d; WaitMode
82 xor    edx, edx; WaitReason
83 lea    rcx, [r14+48h]; Object
84 call   KeWaitForSingleObject
85 xor    si1, si1
86 mov    [rsp+0B8h+var_78], si1
87 mov    ecx, [rbx+24h]
88 test   cl, cl
89 jns    short loc_140474ADB
90loc_140474470:
```



```
91 and   ecx, 1
92 jz    short loc_140474A9E
93loc_140474475:
94 mov    rax, [rbx+18h]
95 mov    edx, [rax+0C0h]
96 cmp    edx, 3
97 jz    short loc_140474A89
98loc_140474484:
99 cmp    edx, 4
100 jns    short loc_140474A9E
101loc_140474489:
102 mov    si1, 1
103 mov    [rsp+0B8h+var_78], si1
104 mov    r14d, 800h
105 mov    dword ptr [rsp+0B8h+undefined_value], r14d
106 jmp    short loc_140474ADB
107loc_14047449E:
108 test   ecx, ecx
109 jns    short loc_140474A9F
110loc_1404744A3:
111 mov    rax, [rbx+18h]
112 cmp    dword ptr [rax+0C0h], 5
113 jz    short loc_140474AC3
114loc_1404744A5:
```

Diffing - Hex-Rays pre-cleanup

```
91         v17);
92     v15 = v18;
93     if ( *(_BYTE *) (v9 + 172) & 4 )
94     v15 = 1;
95     v18 = v15;
96     ObfDereferenceObject(v7 - 17);
97     KeWaitForSingleObject((char *)v1 + 40, Executive, 0, 0, 0i64);
98     if ( *(_DWORD *)v1 + 7 ) != 2 )
99     goto LABEL_34;
100    v2 = v18;

101 }
102 else
103 {
104     ObfDereferenceObject(v7 - 17);
105 }
106 if ( v2 )
107 {
108     v7 = (_QWORD *) * (_QWORD *)v1 + 34;
109     v2 = 0;
110     v18 = 0;
111 }
112 else
113 {
114LABEL_12:

88         v16);
89     ObfDereferenceObject(v6 - 17);
90     KeWaitForSingleObject((char *)v1 + 40, Executive, 0, 0, 0i64);
91     if ( *(_DWORD *)v1 + 7 ) != 2 )
92     goto LABEL_32;
93     v14 = *(_QWORD *)v1 + 45;
94     if ( !v14 || *(_DWORD *) (v14 + 64) != 3 )
95     goto LABEL_31;
96     v6 = (_QWORD *) * (_QWORD *)v1 + 34;
97 }
98 else
99 {
100     ObfDereferenceObject(v6 - 17);

101LABEL_12:
```

Diffing - Hex-Rays post-cleanup

```
83          0x20u,
84          &cur_enlistment_guid);
85      if ( ADJ(pEnlistment_shifted)->Flags & KENLISTMENT_FINALIZED )
86      bEnlistmentIsFinalized = 1;
87      ObfDereferenceObject(ADJ(pEnlistment_shifted));
88      KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
89      if ( pResMgr->State != KResourceManagerOnline )
90          goto b_release_mutex;

91      }
92      else
93      {
94          ObfDereferenceObject(ADJ(pEnlistment_shifted));
95      }
96      if ( bEnlistmentIsFinalized )
97      {
98          pEnlistment_shifted = EnlistmentHead_addr->Flink;
99          bEnlistmentIsFinalized = 0;
100         bEnlistmentIsFinalized = 0;
101     }
102     else
103     {
104         pEnlistment_shifted = ADJ(pEnlistment_shifted)->NextSameRm.Flink;
105     }
106 }
```

```
83          0x20u,
84          &cur_enlistment_guid);
85      ObfDereferenceObject(ADJ(pEnlistment_shifted));
86      KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
87      if ( pResMgr->State != KResourceManagerOnline )
88          goto b_release_mutex;
89      Tm_ = pResMgr->Tm;
90      if ( !Tm_ || Tm_->State != KKtmOnline )
91      {
92          ret = STATUS_TRANSACTIONMANAGER_NOT_ONLINE;
93          goto b_release_mutex;
94      }
95      pEnlistment_shifted = EnlistmentHead_addr->Flink;
96      }
97      else
98      {
99          ObfDereferenceObject(ADJ(pEnlistment_shifted));
100         pEnlistment_shifted = ADJ(pEnlistment_shifted)->NextSameRm.Flink;
101     }
102 }
```

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment));
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if (!isSuperior && state == KTransactionCommitted)
                || state == KTransactionInDoubt
                || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

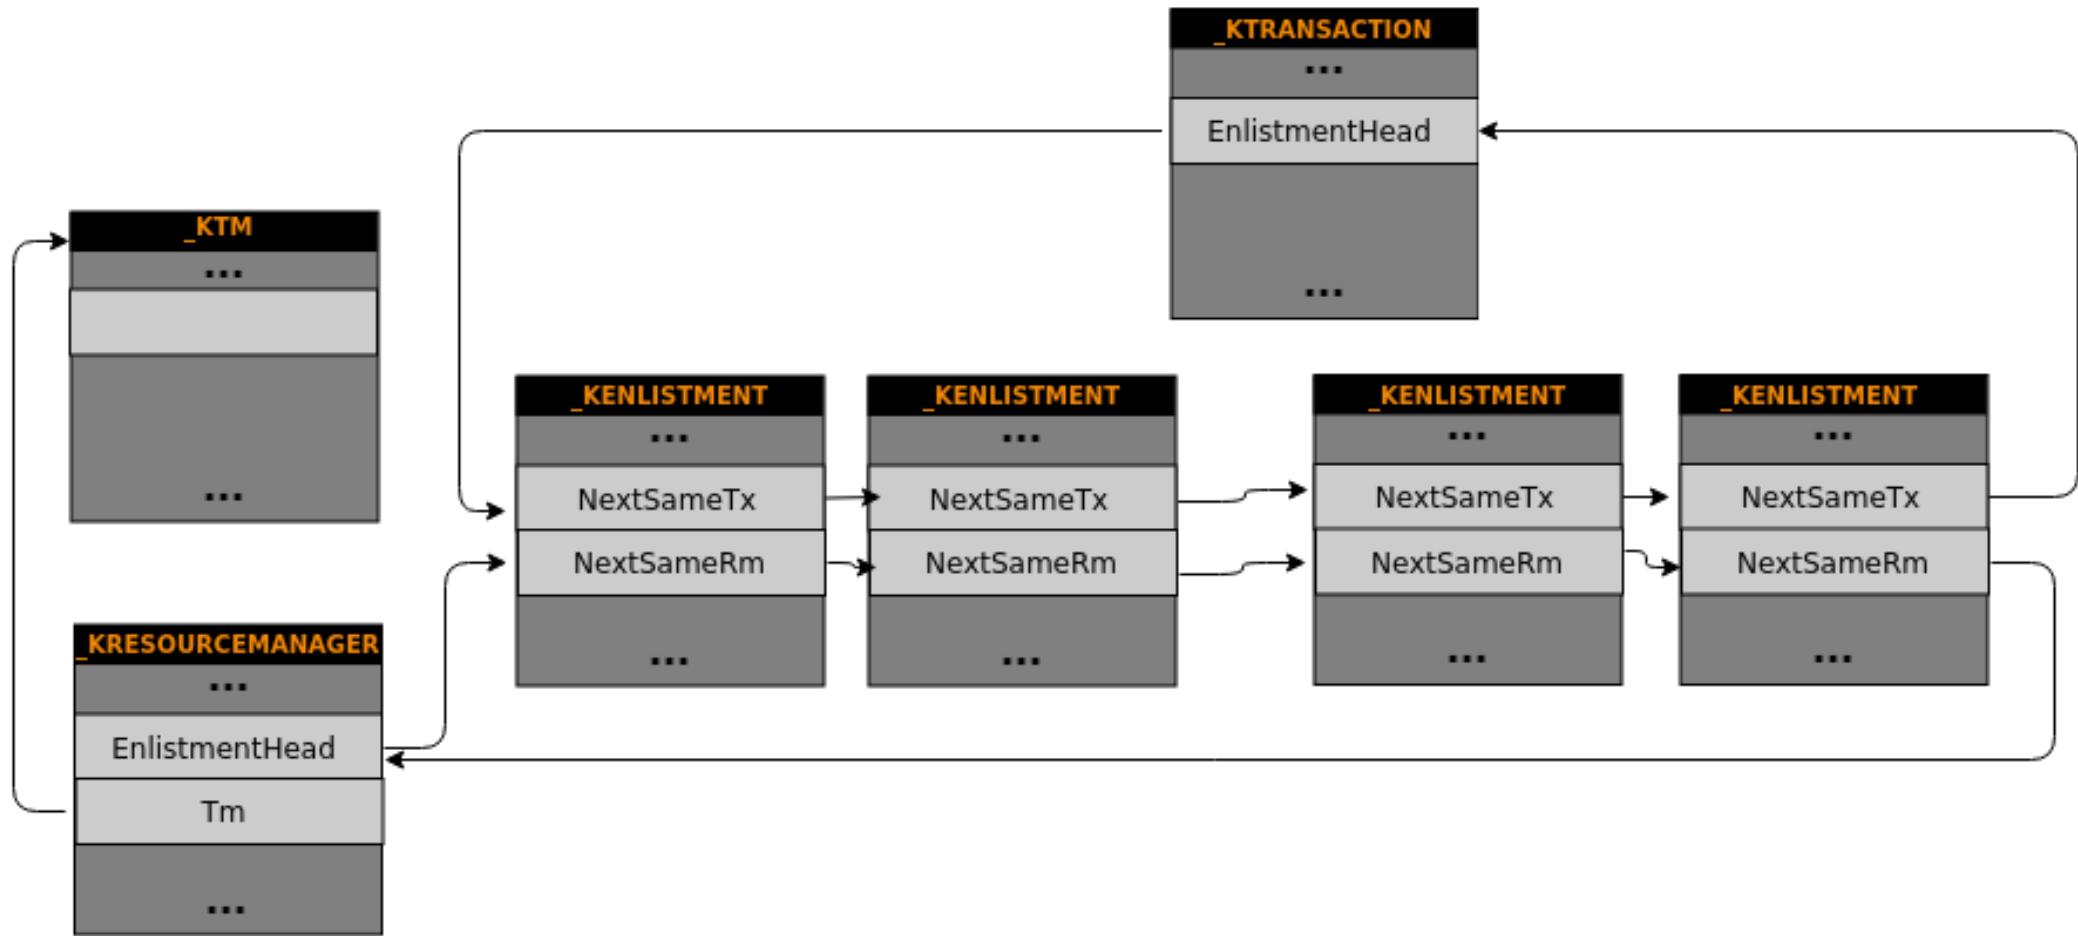
            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

```

Vulnerable
TmRecoverResourceManager() loop



```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment);
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if (!isSuperior && state == KTransactionCommitted)
                || state == KTransactionInDoubt
                || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

```

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {

```

Current enlistment points to
`_KRESOURCEMANAGER`
head to exit loop



```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment);
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior && state == KTransactionCommitted)
                        || state == KTransactionInDoubt
                        || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

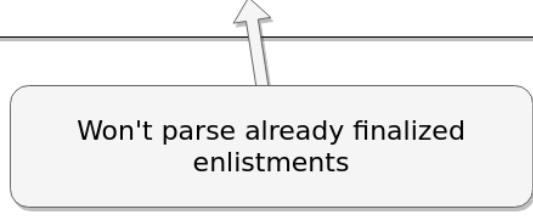
```

```

if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
    pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
} else {

```

Won't parse already finalized
enlistments



```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment);
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior && state == KTransactionCommitted)
                        || state == KTransactionInDoubt
                        || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );
            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

```

```

ObfReferenceObject(pEnlistment));
KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0, 0i64);
bSendNotification = 0;

```

Bump the enlistment ref count and lock the current enlistment

Ref count bump prevents deletion upon finalization while sending notification

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment);
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if (!isSuperior && state == KTransactionCommitted)
                || state == KTransactionInDoubt
                || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
    }
    // ...
    KeReleaseMutex(&pEnlistment->Mutex, 0);

    if ( bSendNotification ) {
        KeReleaseMutex(&pResMgr->Mutex, 0);
        ret = TmpSetNotificationResourceManager( ... );

        if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
            bEnlistmentIsFinalized = 1;
        }

        ObfDereferenceObject(pEnlistment);
        KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
        //...
    } else {
        ObfDereferenceObject(pEnlistment);
    }

    if ( bEnlistmentIsFinalized ) {
        pEnlistment_shifted = EnlistmentHead_addr->Flink;
        bEnlistmentIsFinalized = 0;
    } else {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    }
}

```

```

if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
    // ...
    isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
    state = pEnlistment->Transaction->State;
    if ( ... ) {
        // ...
    } else if ((!isSuperior && state == KTransactionCommitted)
        || state == KTransactionInDoubt
        || state == KTransactionPrepared ) {
        bSendNotification = 1;
        NotificationMask = TRANSACTION_NOTIFY_RECOVER;
    }
    pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
}

```

Each enlistment only gets notified once per loop iteration

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment);
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
    }
    if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
        // ...
        isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
        state = pEnlistment->Transaction->State;
        if ( ... ) {
            // ...
        } else if (!isSuperior && state == KTransactionCommitted)
            || state == KTransactionInDoubt
            || state == KTransactionPrepared ) {
            bSendNotification = 1;
            NotificationMask = TRANSACTION_NOTIFY_RECOVER;
        }
        pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
    }
    // ...
    KeReleaseMutex(&pEnlistment->Mutex, 0);

    if ( bSendNotification ) {
        KeReleaseMutex(&pResMgr->Mutex, 0);
        ret = TmpSetNotificationResourceManager( ... );

        if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
            bEnlistmentIsFinalized = 1;
        }

        ObfDereferenceObject(pEnlistment);
        KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
        //...
    } else {
        ObfDereferenceObject(pEnlistment);
    }

    if ( bEnlistmentIsFinalized ) {
        pEnlistment_shifted = EnlistmentHead_addr->Flink;
        bEnlistmentIsFinalized = 0;
    } else {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    }
}

```

```

if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
    // ...
    isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
    state = pEnlistment->Transaction->State;
    if ( ... ) {
        // ...
    } else if ((!isSuperior && state == KTransactionCommitted)
        || state == KTransactionInDoubt
        || state == KTransactionPrepared ) {
        bSendNotification = 1;
        NotificationMask = TRANSACTION_NOTIFY_RECOVER;
    }
    pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
}

```

Send an enlistment notification
for specific transaction states

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment));
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior & state == KTransactionCommitted)
                        || state == KTransactionInDoubt
                        || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );
        }

        if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
            bEnlistmentIsFinalized = 1;
        }

        ObfDereferenceObject(pEnlistment);
        KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0i64);
        //...
    } else {
        ObfDereferenceObject(pEnlistment);
    }

    if ( bEnlistmentIsFinalized ) {
        pEnlistment_shifted = EnlistmentHead_addr->Flink;
        bEnlistmentIsFinalized = 0;
    } else {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    }
}

```

// ...
KeReleaseMutex(&pEnlistment->Mutex, 0);

if (bSendNotification) {
 KeReleaseMutex(&pResMgr->Mutex, 0);
 ret = TmpSetNotificationResourceManager(...);

Unlock resource manager mutex!
Finalizing enlistments is now
possible, which can lead to deletion
if refcount = 0

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment));
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior & state == KTransactionCommitted)
                || state == KTransactionInDoubt
                || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&ResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&ResMgr->Mutex, Executive, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

```

```

if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
    bEnlistmentIsFinalized = 1;
}

```

Attempt to prevent a use-after-free

Will not use finalized enlistment here
if boolean is set

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment));
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior & state == KTransactionCommitted)
                        || state == KTransactionInDoubt
                        || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&ResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }
            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&ResMgr->Mutex, Executive, 0, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }

        if ( bEnlistmentIsFinalized ) {
            pEnlistment_shifted = EnlistmentHead_addr->Flink;
            bEnlistmentIsFinalized = 0;
        } else {
            pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
        }
    }
}

```

ObfDereferenceObject(pEnlistment);
 KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
 //...

Lower ref count. If enlistment is finalized
 before relocking mutex, pEnlistment
 points to freed memory

Prone to race condition abuse. Can
 congest this mutex from userland.

```

pEnlistment_shifted = EnlistmentHead_addr->Flink;
while ( pEnlistment_shifted != EnlistmentHead_addr ) {
    pEnlistment = ADJ(pEnlistment_shifted)
    if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
        pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
    } else {
        ObfReferenceObject(pEnlistment));
        KeWaitForSingleObject(&pEnlistment->Mutex, Executive, 0, 0, 0i64);
        bSendNotification = 0;
        if ( (pEnlistment->Flags & KENLISTMENT_IS_NOTIFIABLE) != 0 ) {
            // ...
            isSuperior = pEnlistment->Flags & KENLISTMENT_SUPERIOR;
            state = pEnlistment->Transaction->State;
            if ( ... ) {
                // ...
            } else if ( (!isSuperior & state == KTransactionCommitted)
                || state == KTransactionInDoubt
                || state == KTransactionPrepared ) {
                bSendNotification = 1;
                NotificationMask = TRANSACTION_NOTIFY_RECOVER;
            }
            pEnlistment->Flags &= ~KENLISTMENT_IS_NOTIFIABLE;
        }
        // ...
        KeReleaseMutex(&pEnlistment->Mutex, 0);

        if ( bSendNotification ) {
            KeReleaseMutex(&pResMgr->Mutex, 0);
            ret = TmpSetNotificationResourceManager( ... );

            if ( pEnlistment->Flags & KENLISTMENT_FINALIZED ) {
                bEnlistmentIsFinalized = 1;
            }

            ObfDereferenceObject(pEnlistment);
            KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
            //...
        } else {
            ObfDereferenceObject(pEnlistment);
        }
    }
}

```

```

if ( bEnlistmentIsFinalized ) {
    pEnlistment_shifted = EnlistmentHead_addr->Flink;
    bEnlistmentIsFinalized = 0;
} else {
    pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
}

```

Safe use of resource managers head pointer if race lost

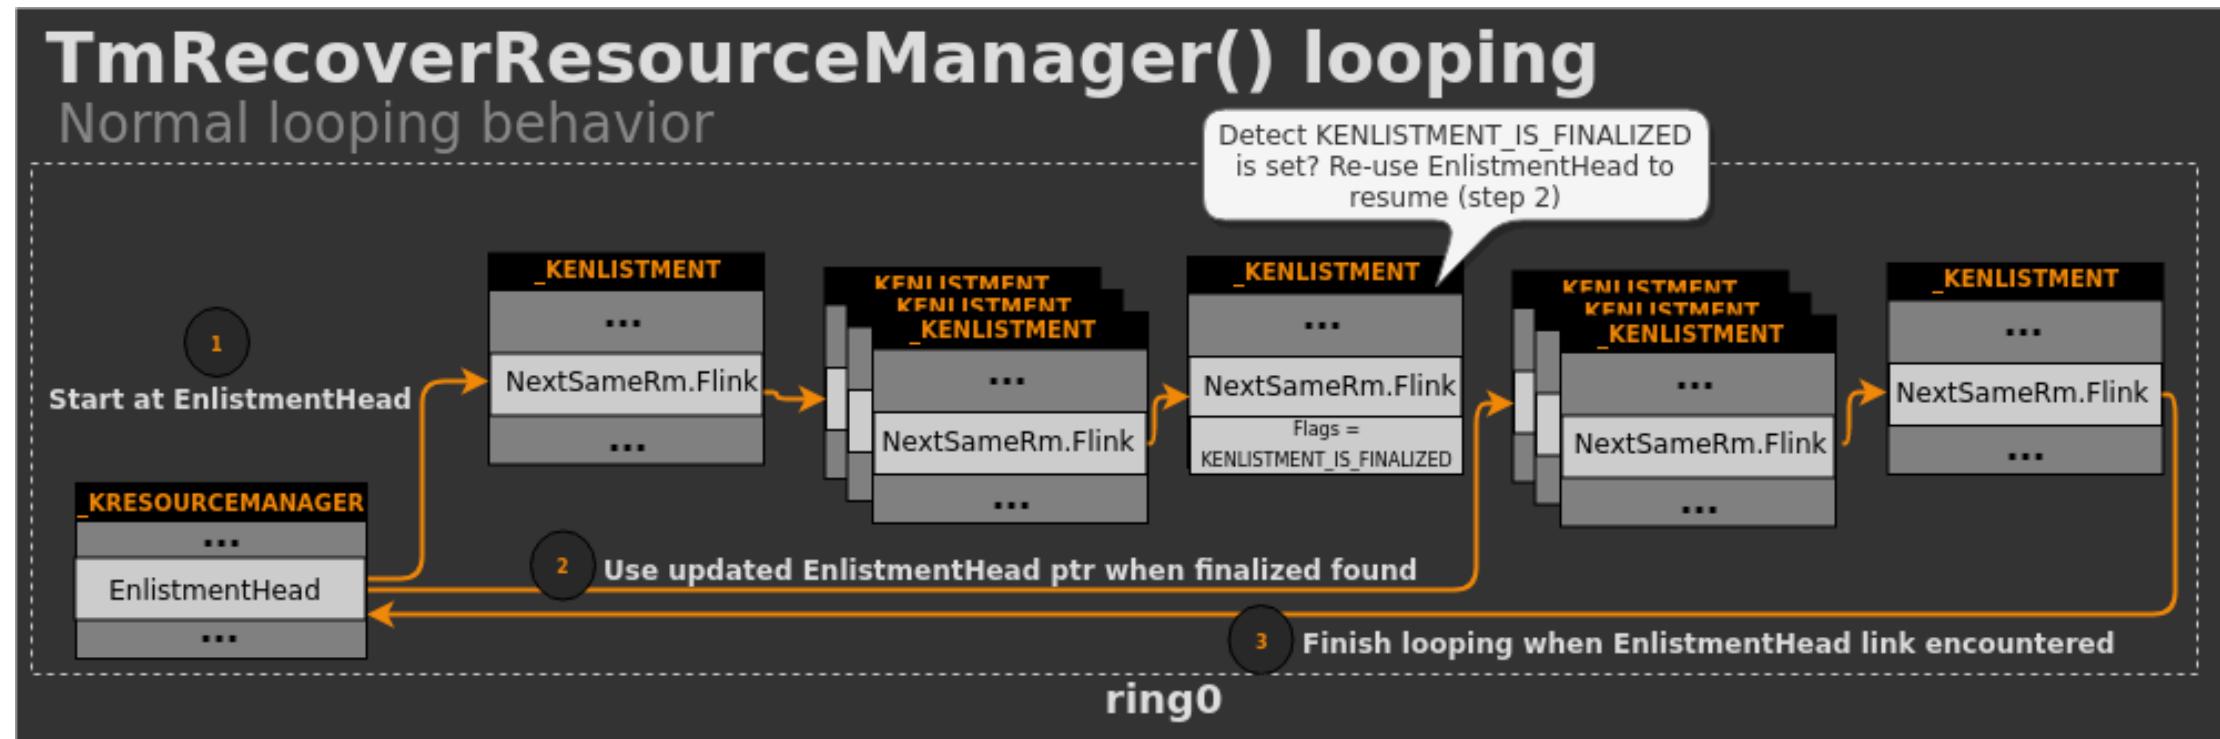
```

if ( bEnlistmentIsFinalized ) {
    pEnlistment_shifted = EnlistmentHead_addr->Flink;
    bEnlistmentIsFinalized = 0;
} else {
    pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
}

```

Used after free if race condition is won

What does TmRecoverResourceManager() normally do?



Vulnerability analysis key points

- A recovering _KRESOURCEMANAGER is unlocked in order to queue a notification
- Code retains pointer to associated _KENLISTMENT, but no lock
- Sends notifications about said _KENLISTMENT
- Attempts to tell if _KENLISTMENT is finalized, but in a racable location
- Drops the reference count by 1, which allows it to become freed if already finalized
- Relocks _KRESOURCEMANAGER
- Tests for a boolean that wasn't set if race condition occurs
- Uses retained _KENLISTMENT pointer
- _KENLISTMENT could now be freed

Triggering CVE-2018-8611

Faking a race win

- Use WinDbg to force race window open
- Patch KeWaitForSingleObject() so we guarantee pEnlistment is freed
 - Patch is just an infinite loop

```
//...
ObfDereferenceObject(pEnlistment);
KeWaitForSingleObject(&pResMgr->Mutex, Executive, 0, 0, 0i64);
//...
} else {
    ObfDereferenceObject(pEnlistment);
}

if ( bEnlistmentIsFinalized ) {
    pEnlistment_shifted = EnlistmentHead_addr->Flink;
    bEnlistmentIsFinalized = 0;
} else {
    pEnlistment_shifted = pEnlistment->NextSameRm.Flink;
}
```

- After freeing all _KENLISTMENTS test if pEnlistment->NextSameRm references freed memory

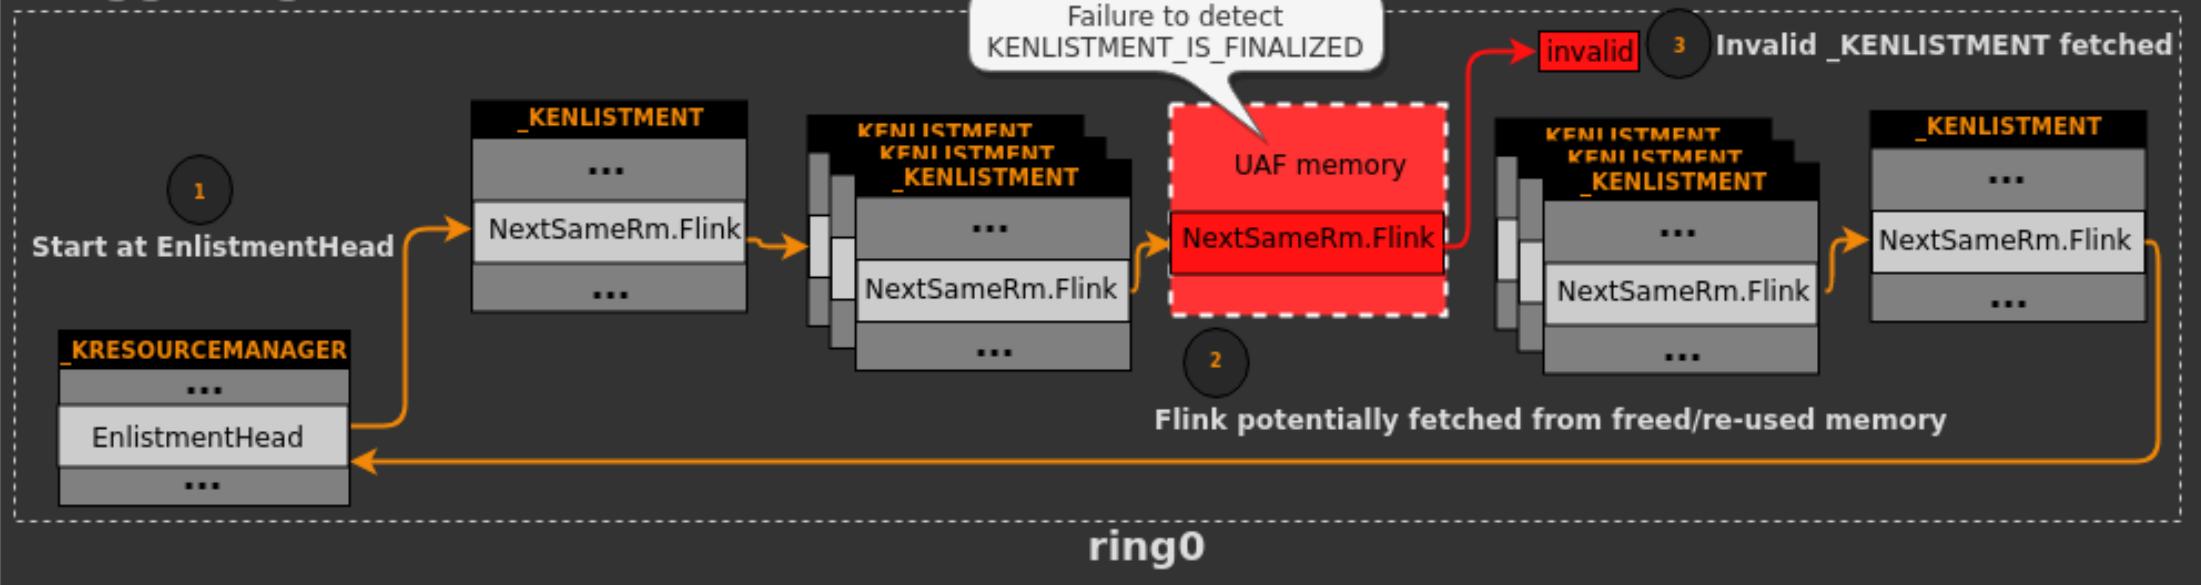
Which _KENLISTMENT to free?

- If we spam a lot of _KENLISTMENT and try to repeatably race...
 - How do we know which one to free?
 - Can't just free them all every time, as we want to maximize attempts
- GetNotificationResourceManager() tells us what a enlistment has been touched by the loop!
- Vulnerable function unlocks the RM specifically to send a notification
 - Correlate the notification to the enlistment, and free it
- Remove infinite loop after we triggered free from userland
- If UAF triggers, it confirms our understanding of the bug
- Run with Driver Verifier to easily confirm

Exploitable loop state

TmRecoverResourceManager() looping

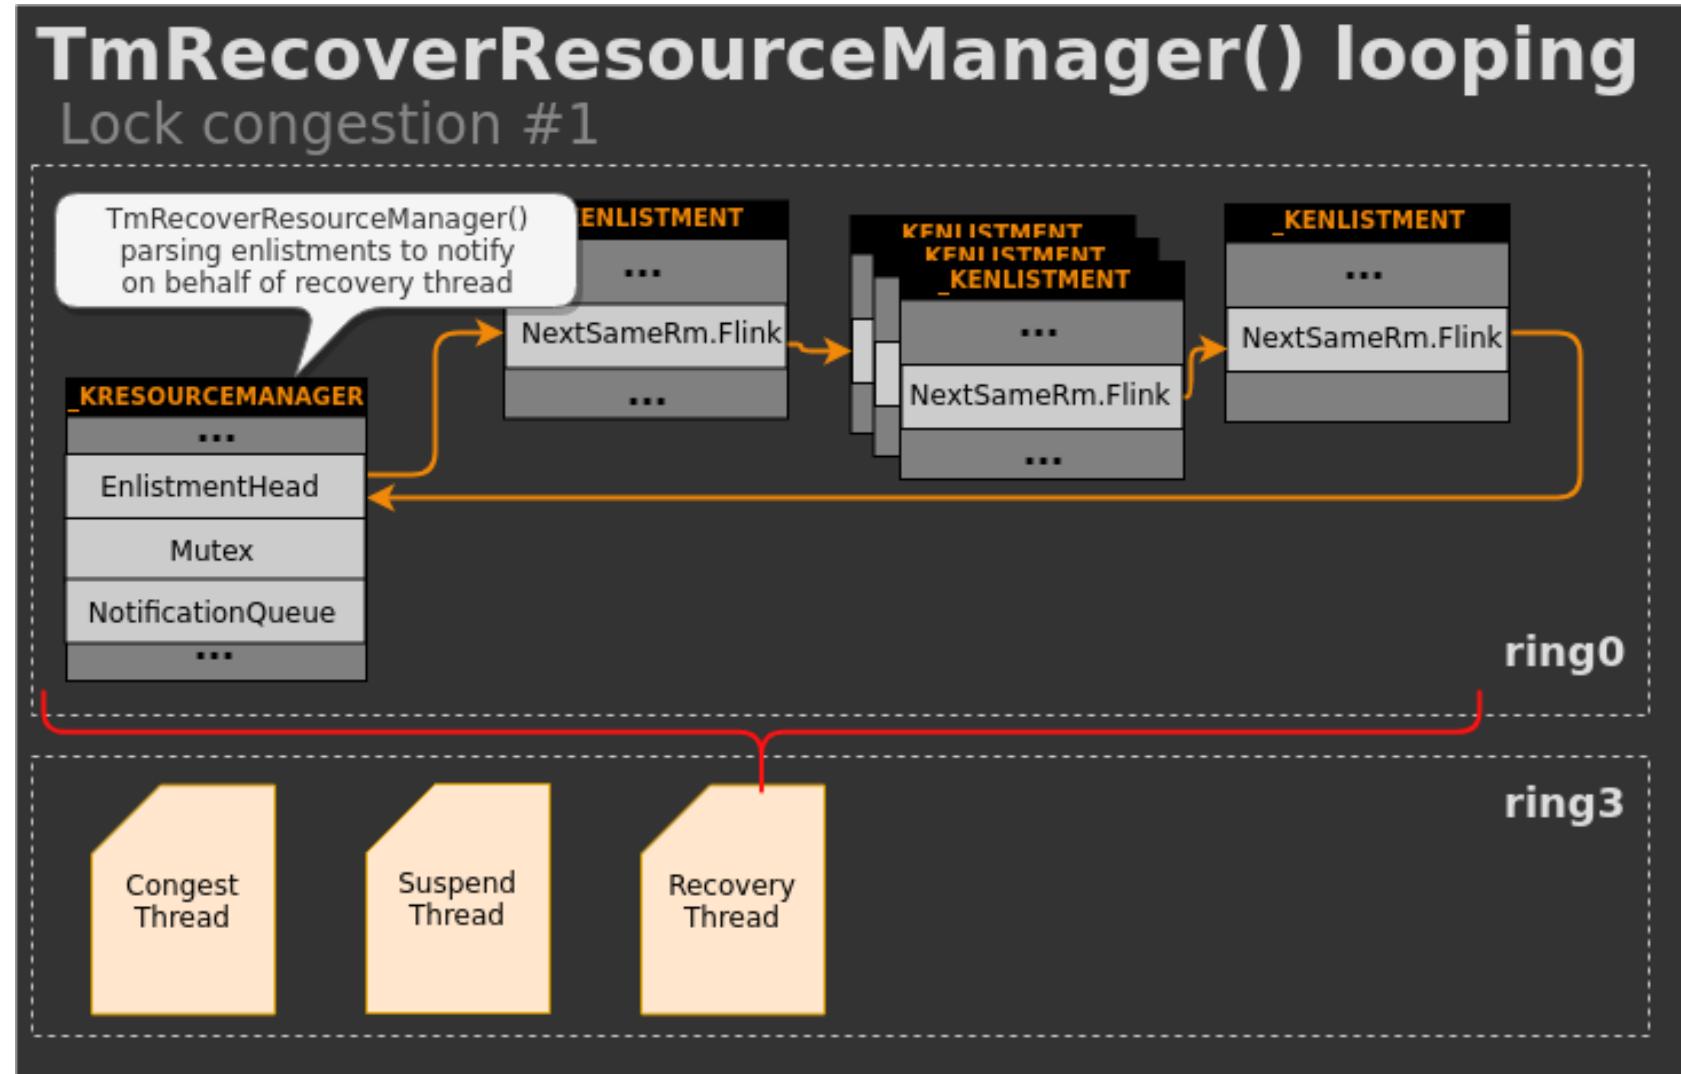
Triggering the vulnerable condition



Actually winning the race

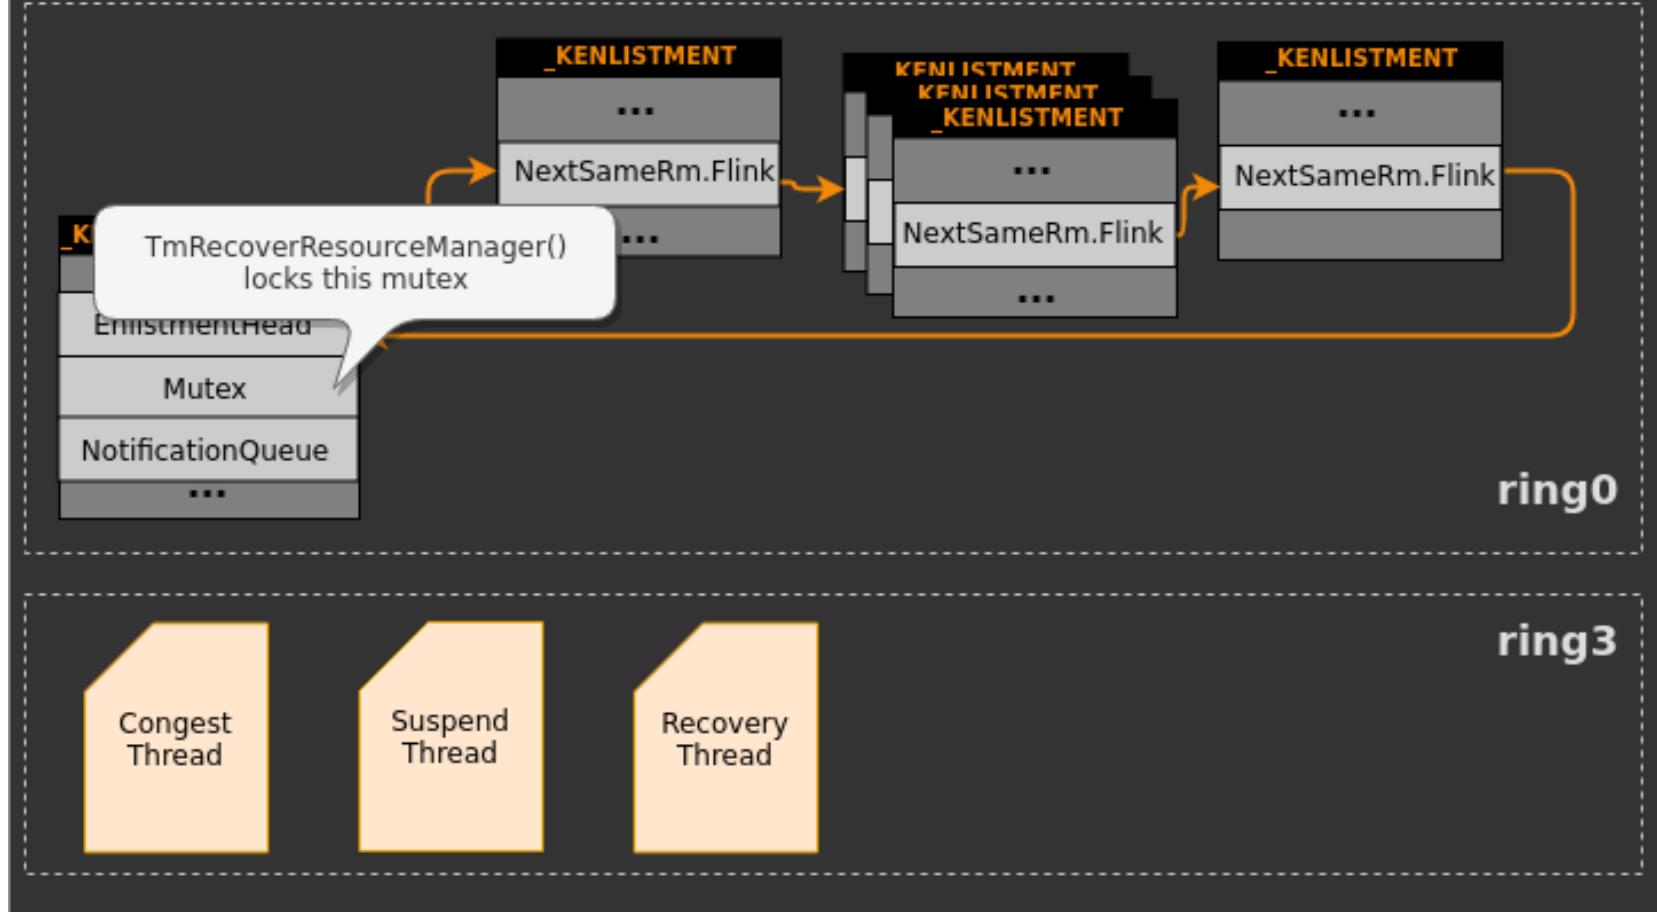
- How do we win this race without patching KeWaitForSingleObject()
 - Was hinted in the Kaspersky blog (though still not obvious to us for quite some time)
 - Suspend the thread stuck in the TmRecoverResourceManager() causing it to effectively block until woken up
- A thread will become blocked on some natural blocking point
 - Like waiting to lock the congested RM mutex
- Congest RM lock to increase likelihood of thread suspending where we want
 - Have a higher priority thread constantly triggering syscall that locks RM
 - Ex: Query the RM description (NtQueryInformationResourceManager)
- Result
 - If thread suspended at right place, we have all the time in the world to free the enlistment
 - If not, no UAF happens, and we keep trying

Lock congestion



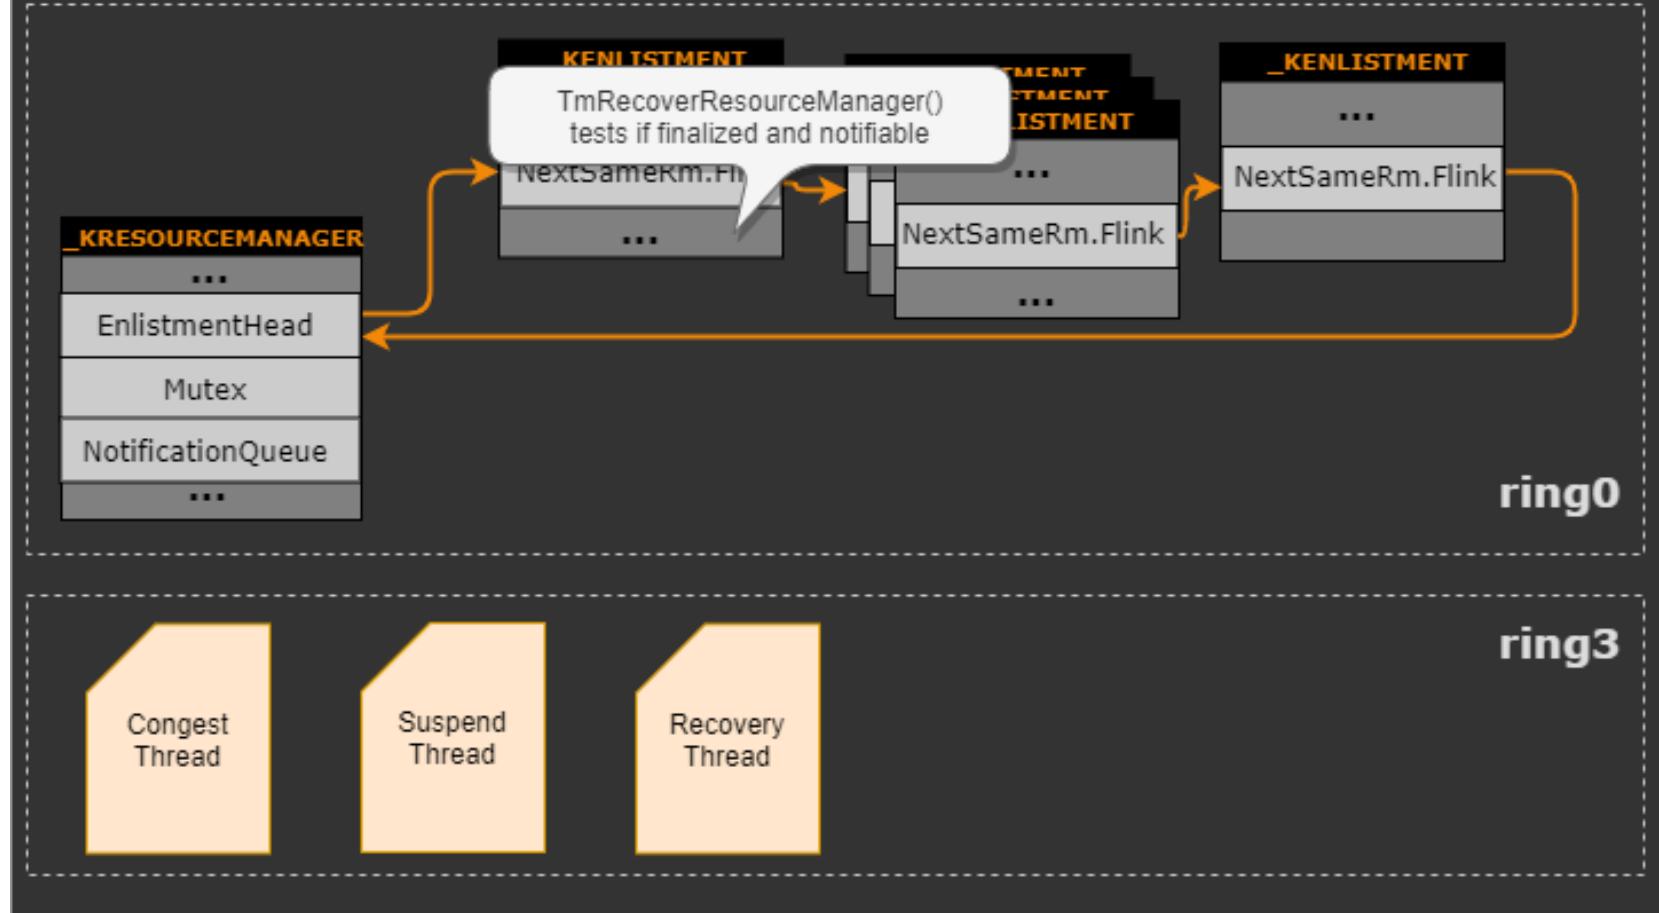
Lock congestion

TmRecoverResourceManager() looping Lock congestion #2



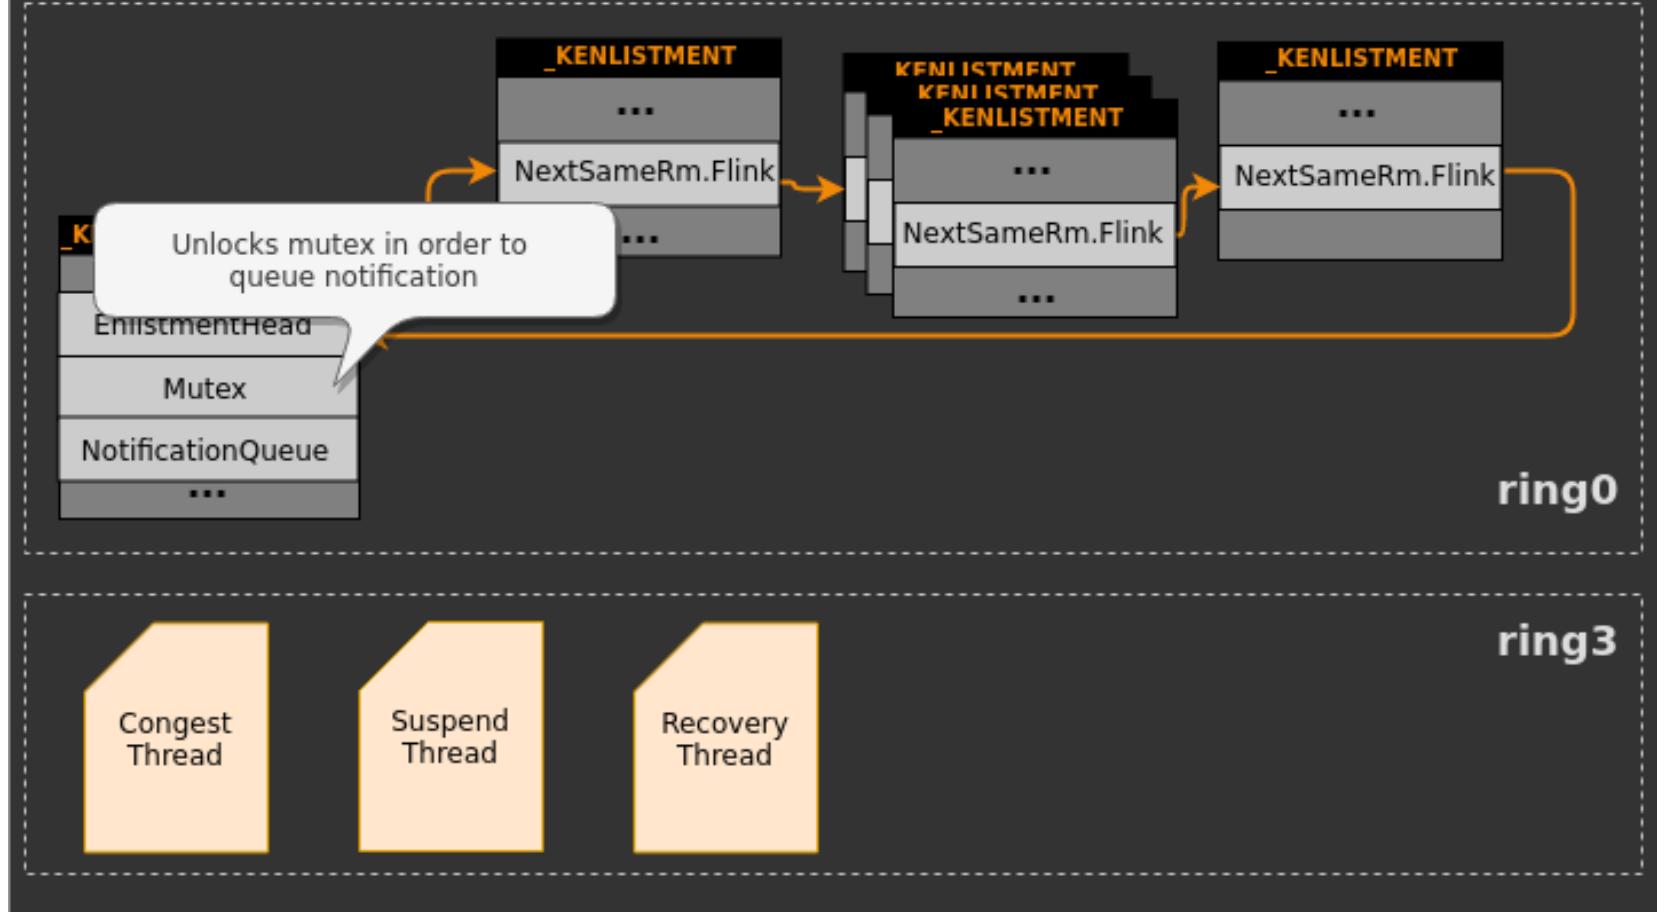
Lock congestion

TmRecoverResourceManager() looping Lock congestion #3



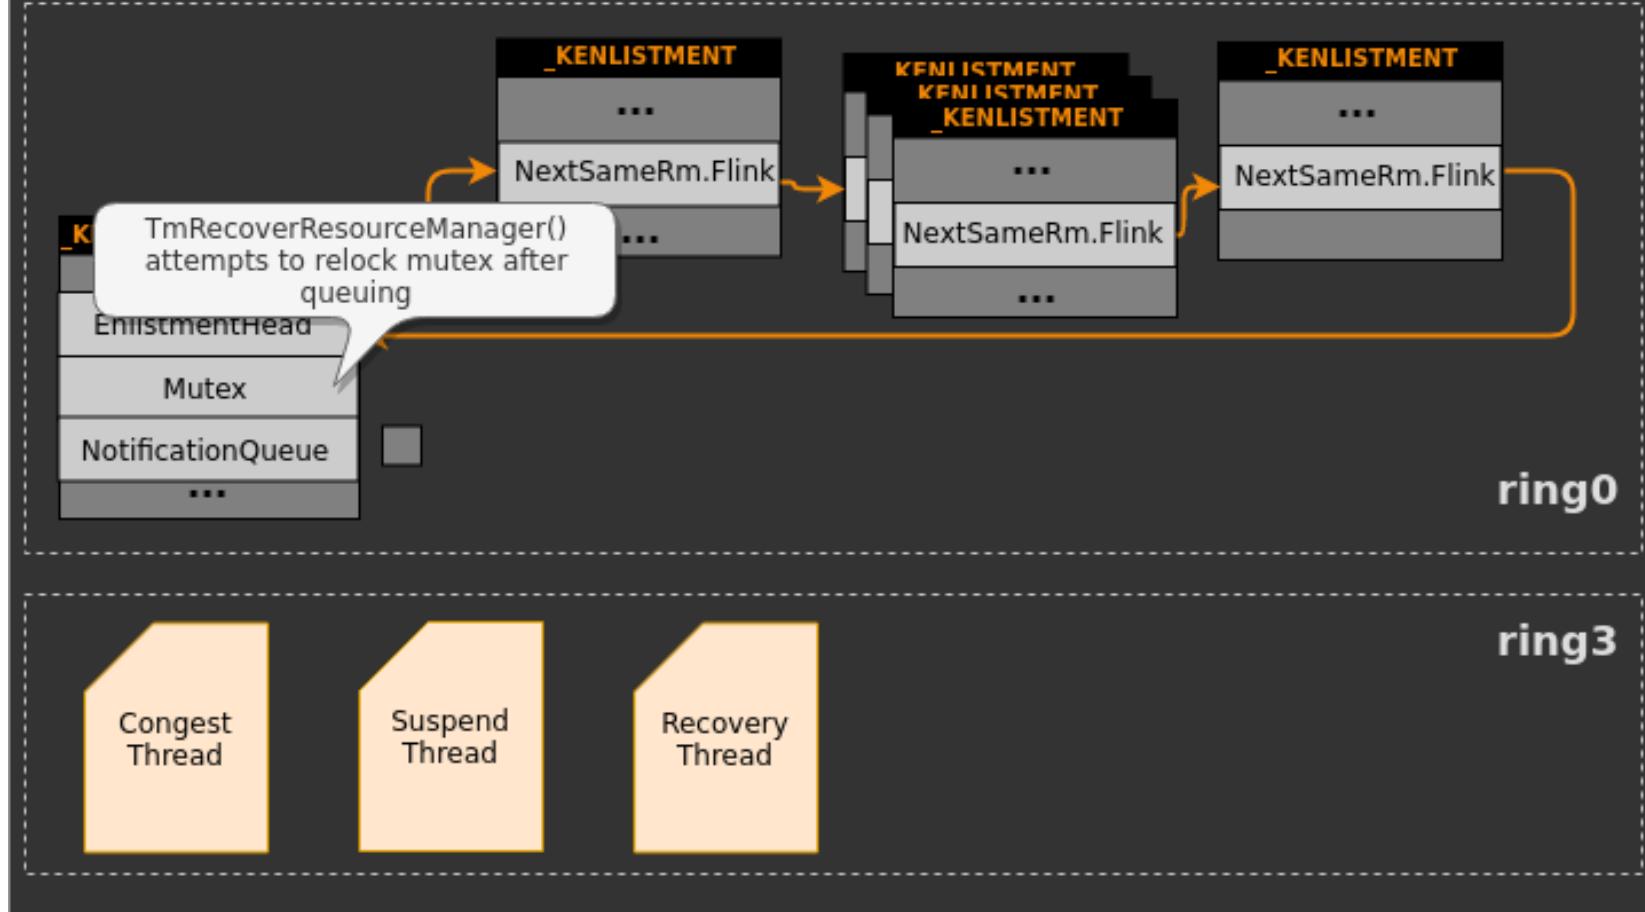
Lock congestion

TmRecoverResourceManager() looping Lock congestion #4



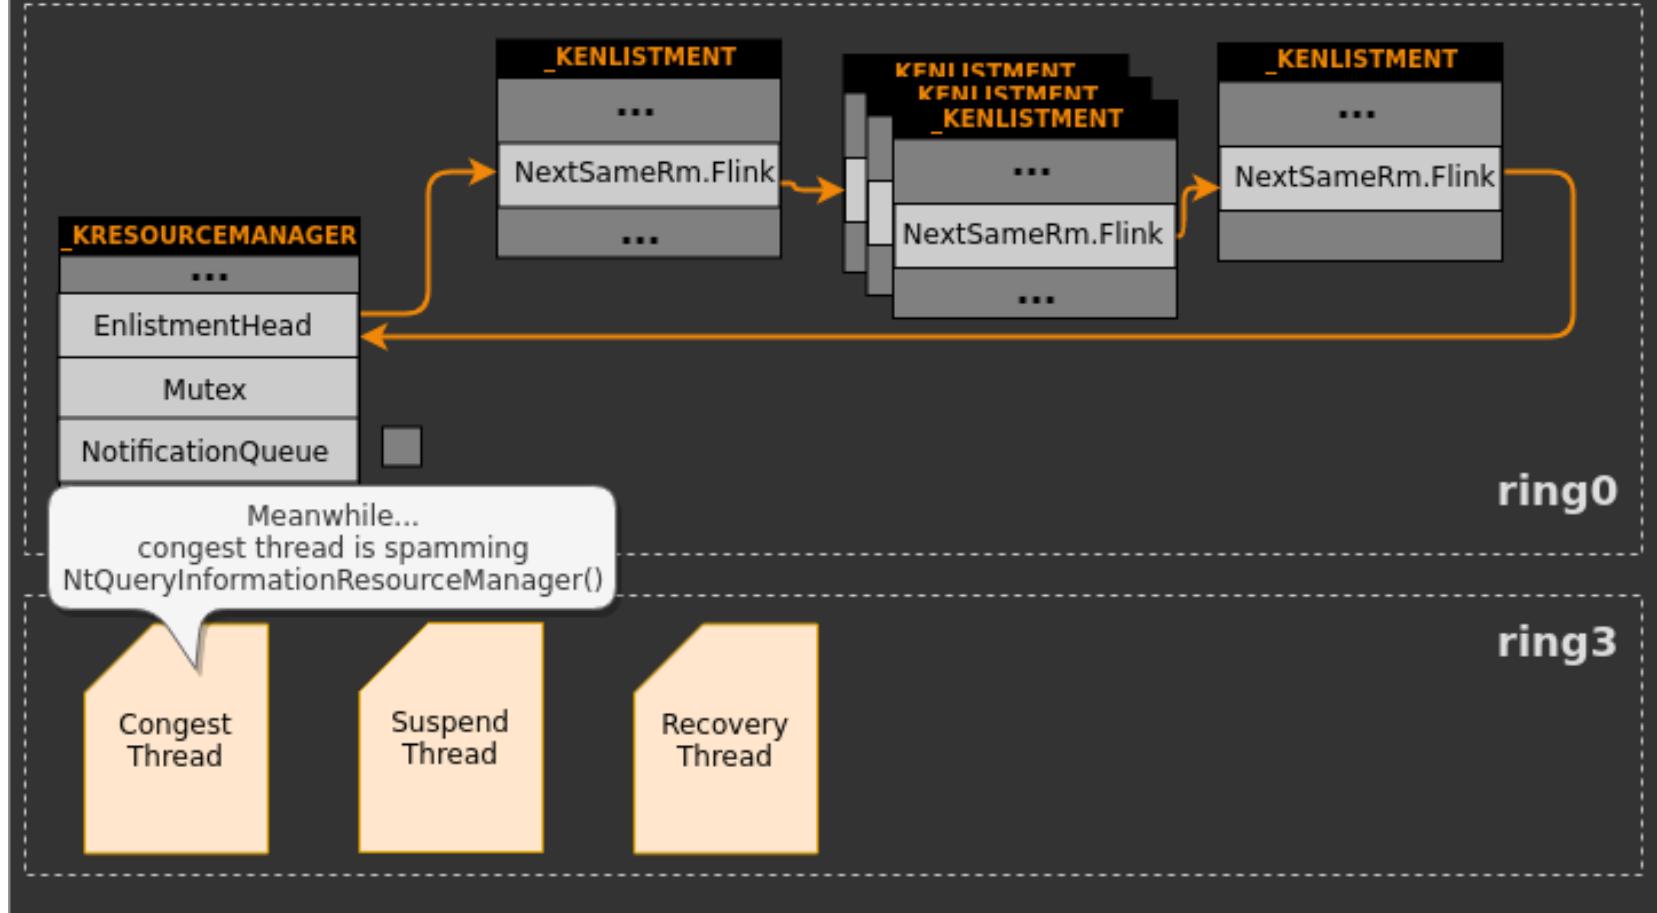
Lock congestion

TmRecoverResourceManager() looping Lock congestion #5



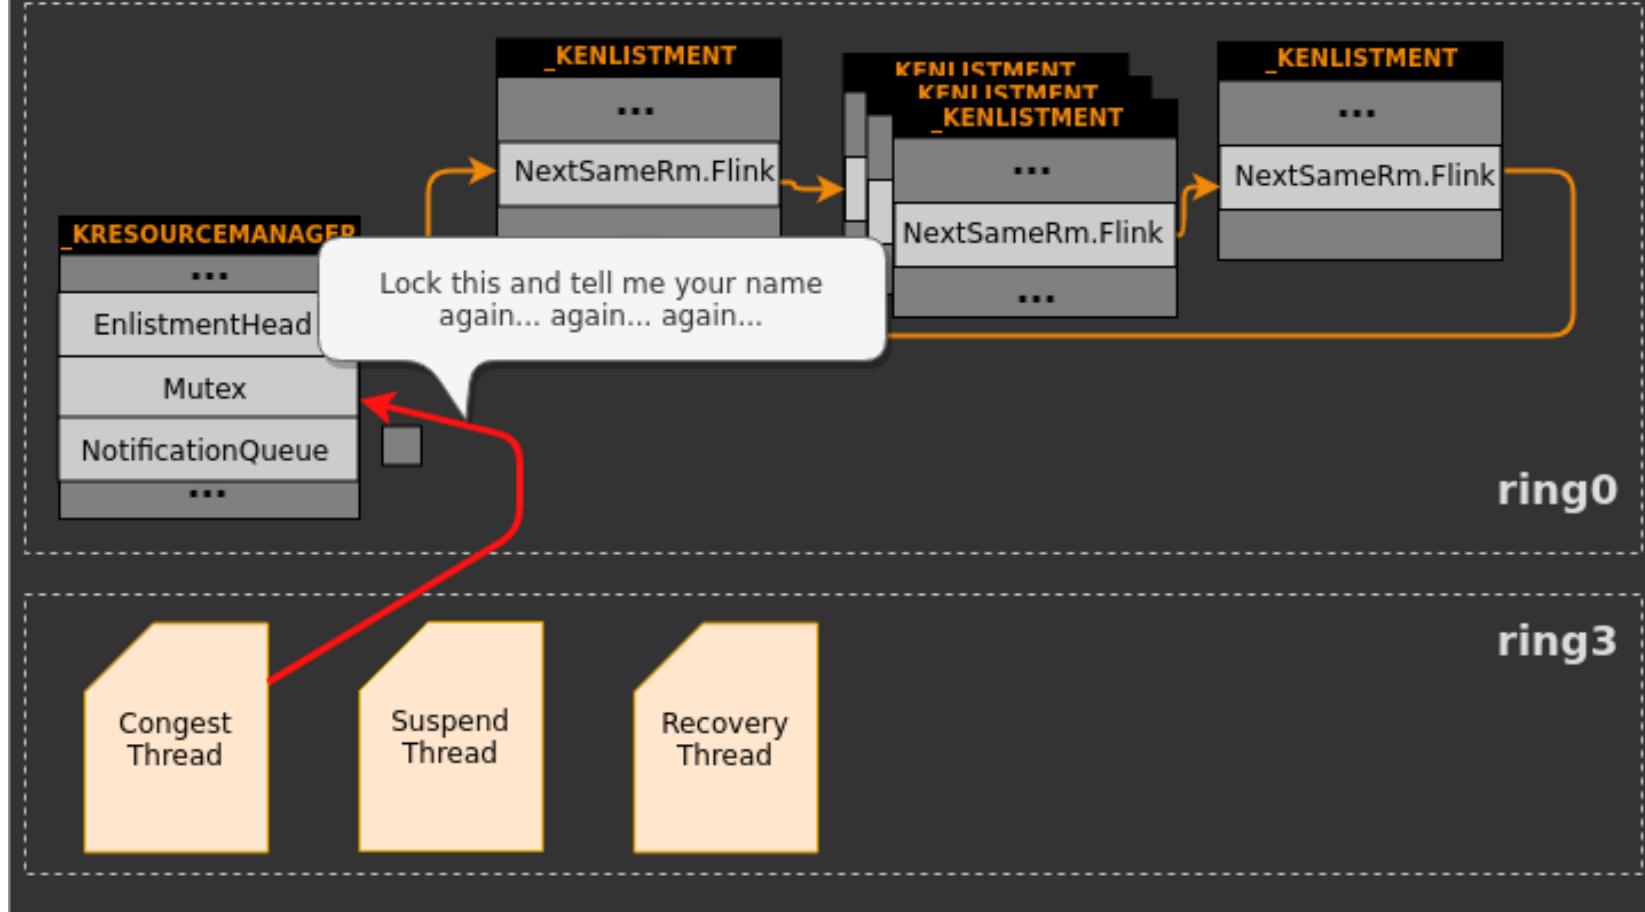
Lock congestion

TmRecoverResourceManager() looping Lock congestion #6



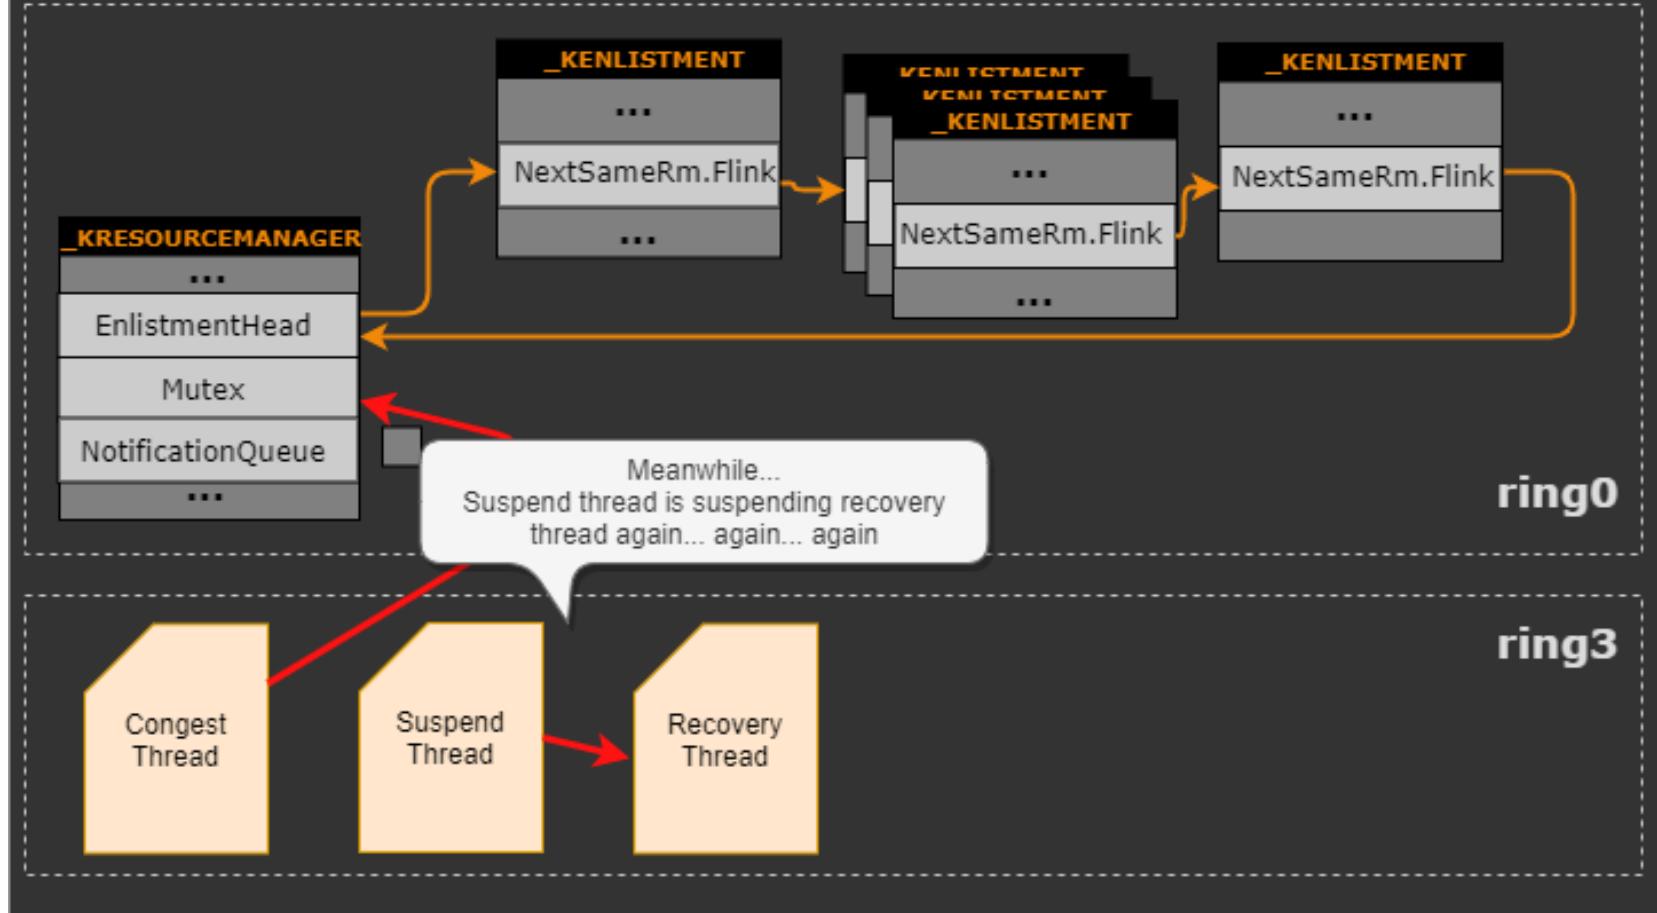
Lock congestion

TmRecoverResourceManager() looping Lock congestion #7



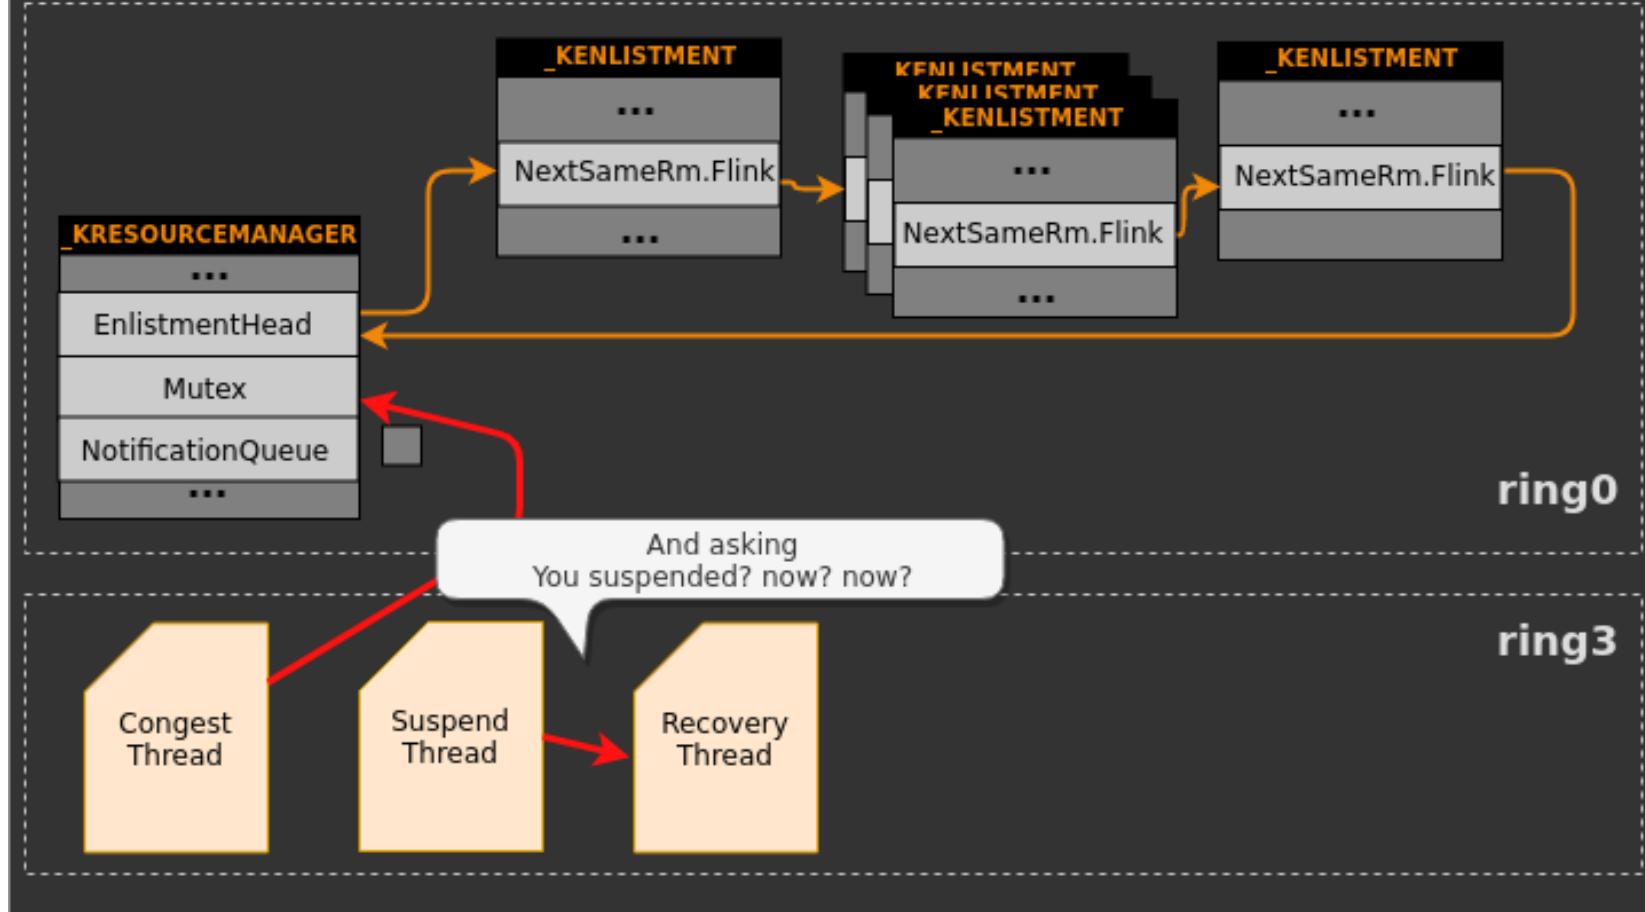
Lock congestion

TmRecoverResourceManager() looping Lock congestion #8



Lock congestion

TmRecoverResourceManager() looping Lock congestion #9

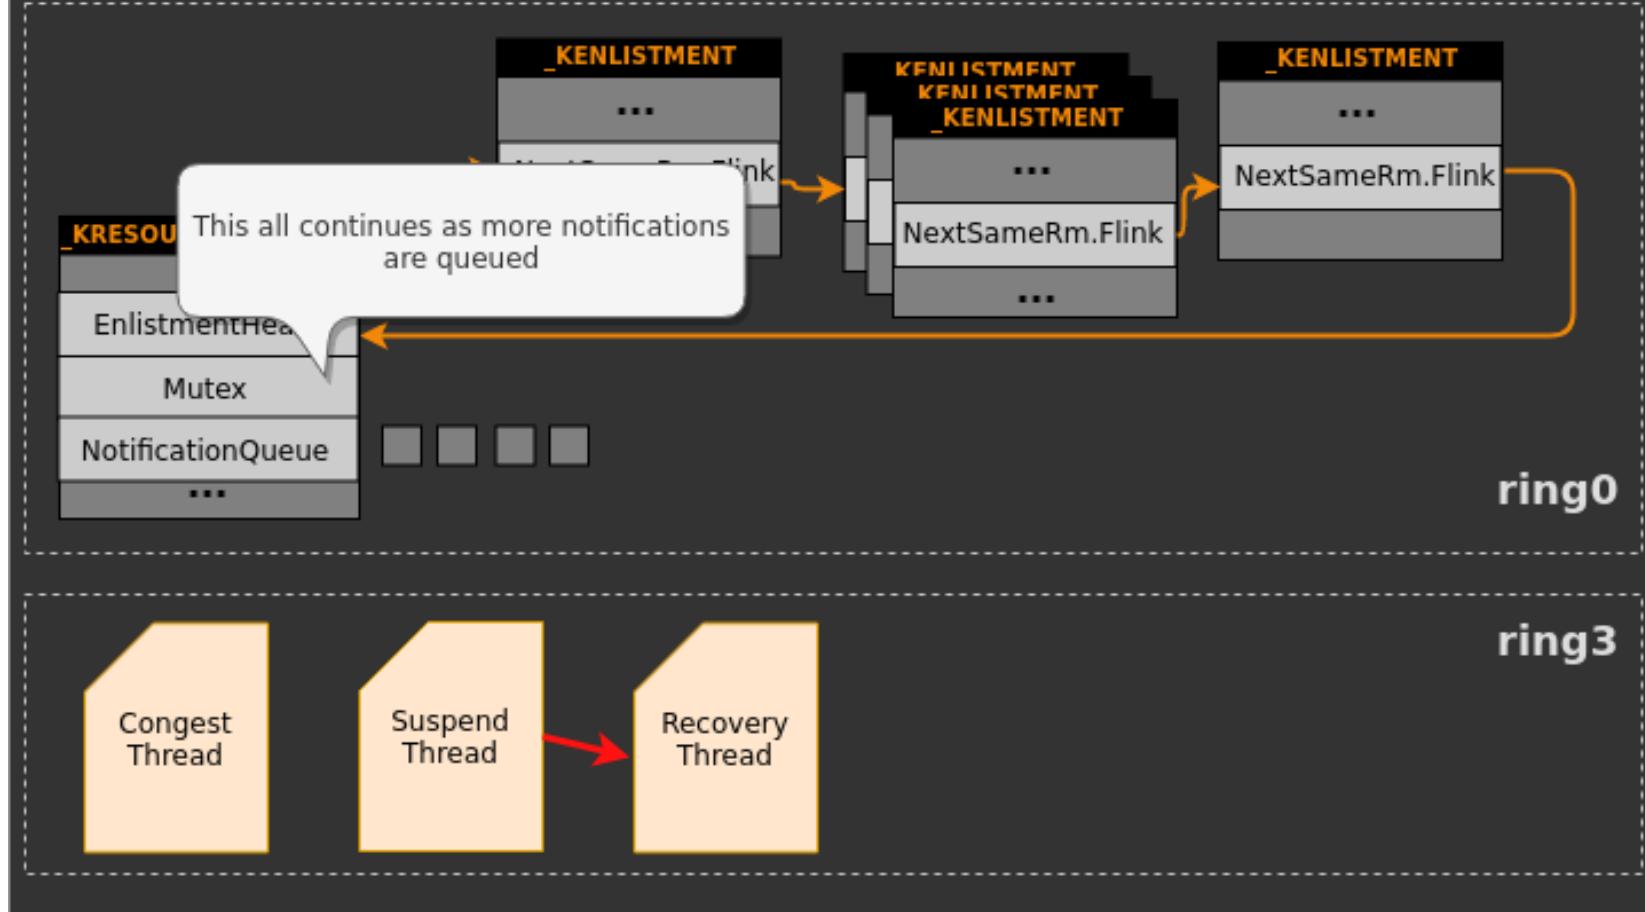


Thread suspension detection

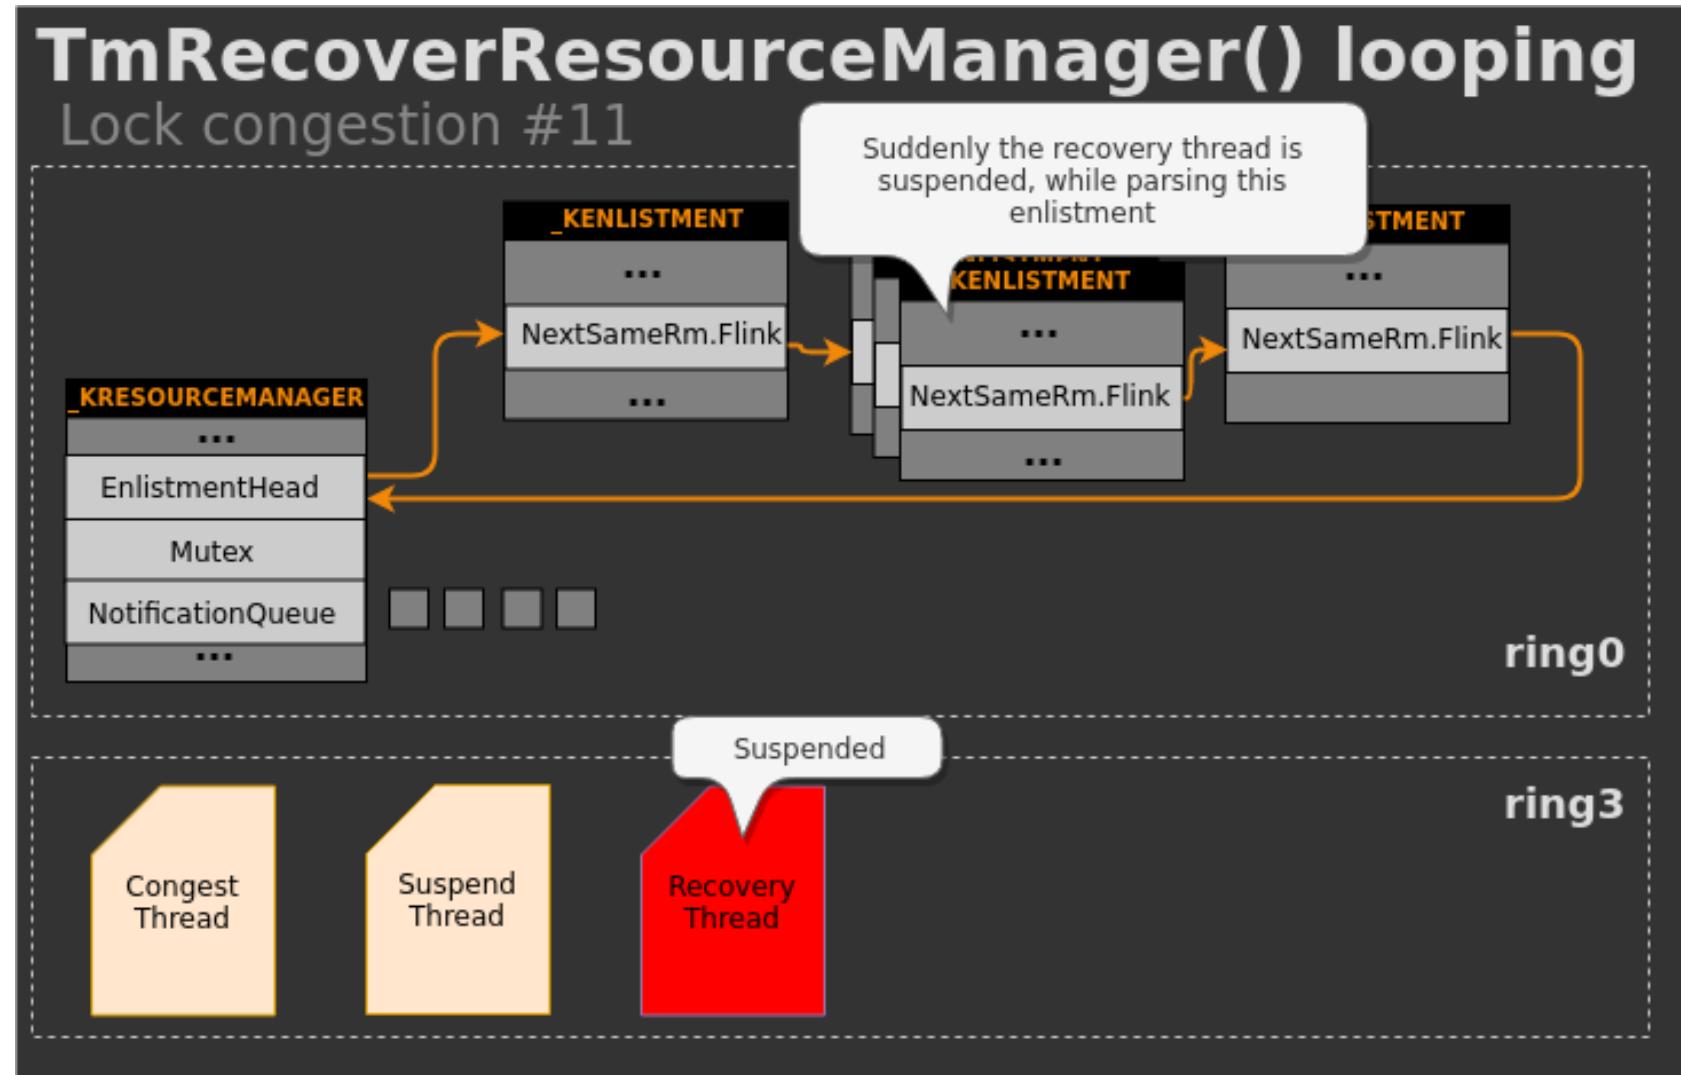
- How can you tell if a thread is suspended?
 - Use [NtQueryThreadInformation\(\)](#) to query thread
 - ThreadInformationClass of ThreadLastSyscall
 - Returns STATUS_UNSUCCESSFUL if thread is not suspended

Lock congestion

TmRecoverResourceManager() looping Lock congestion #10

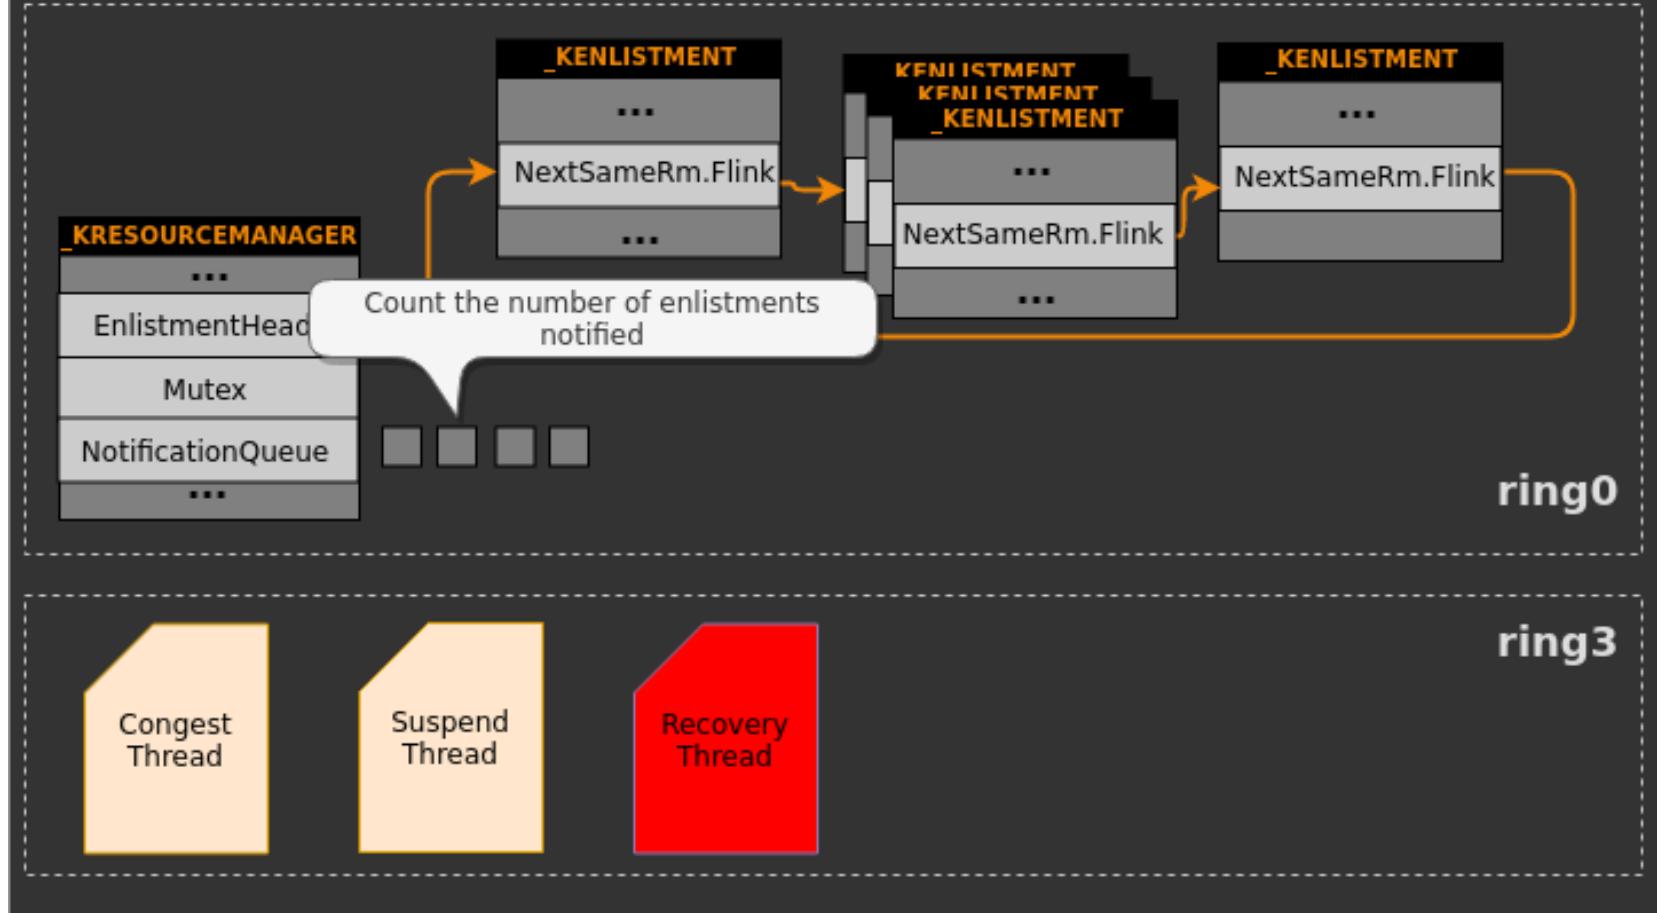


Lock congestion



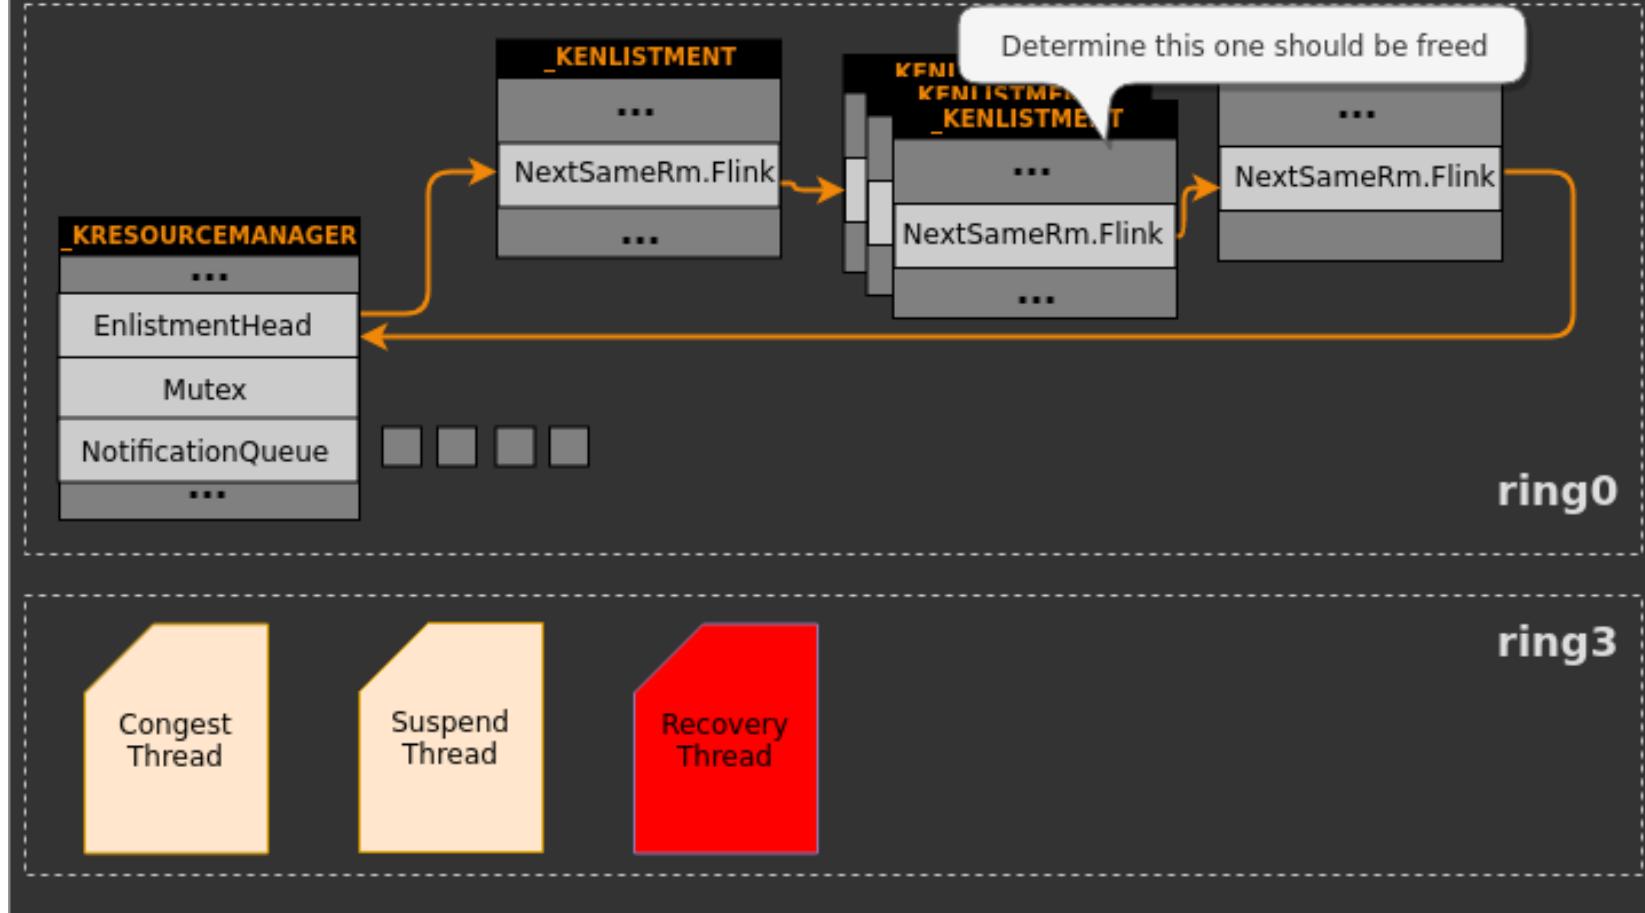
Lock congestion

TmRecoverResourceManager() looping Lock congestion #12



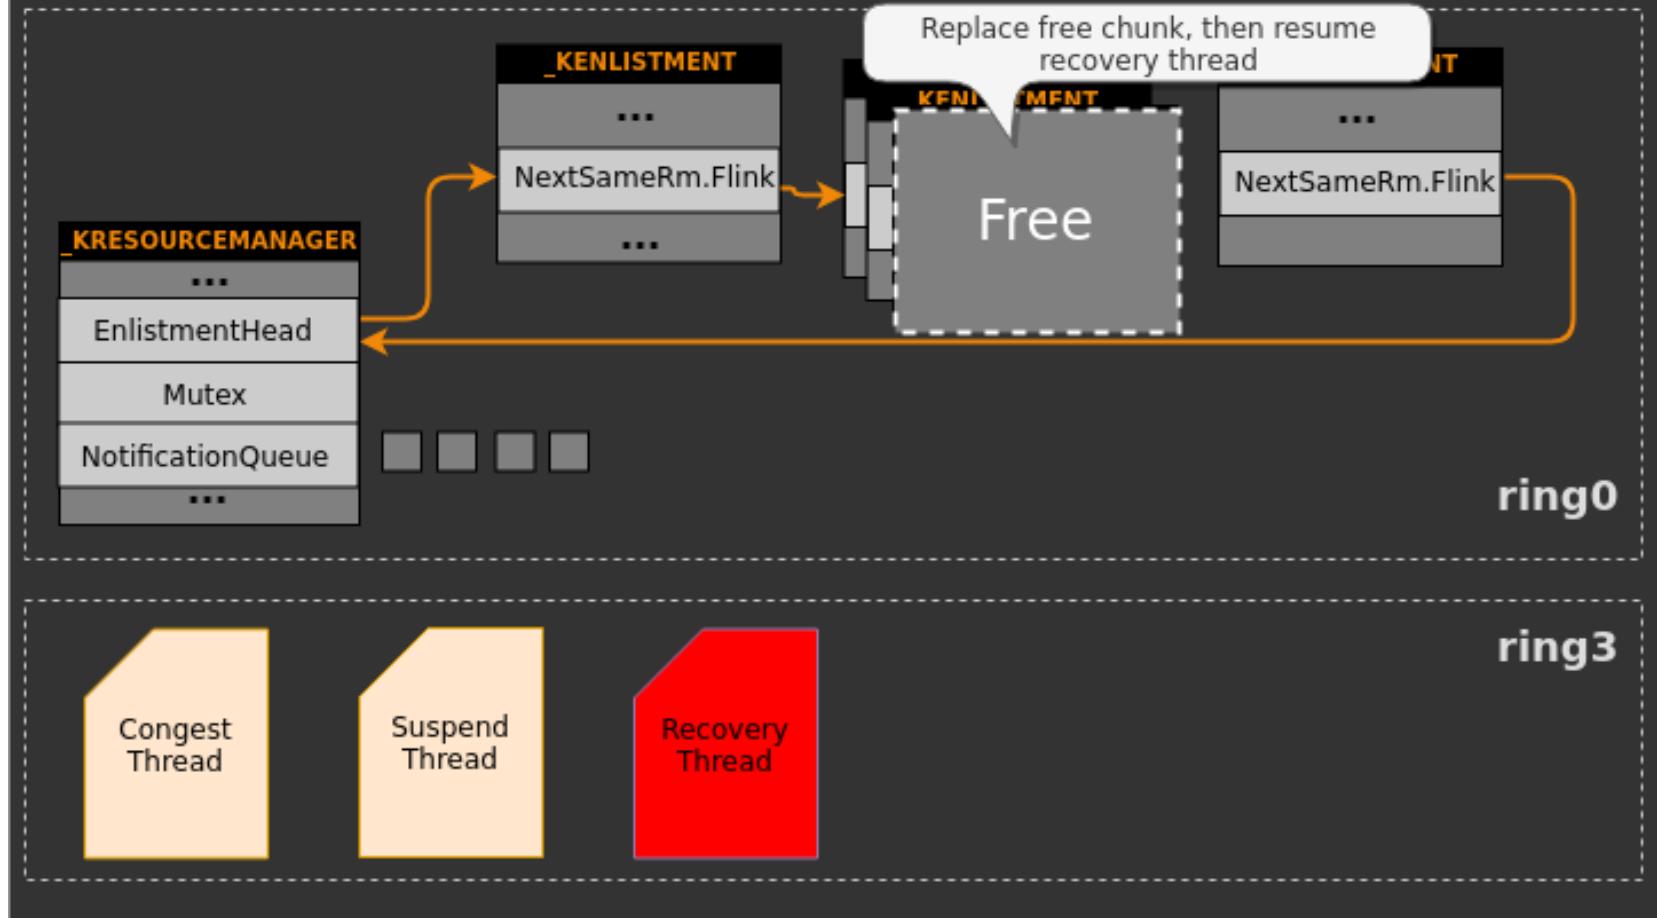
Lock congestion

TmRecoverResourceManager() looping Lock congestion #13



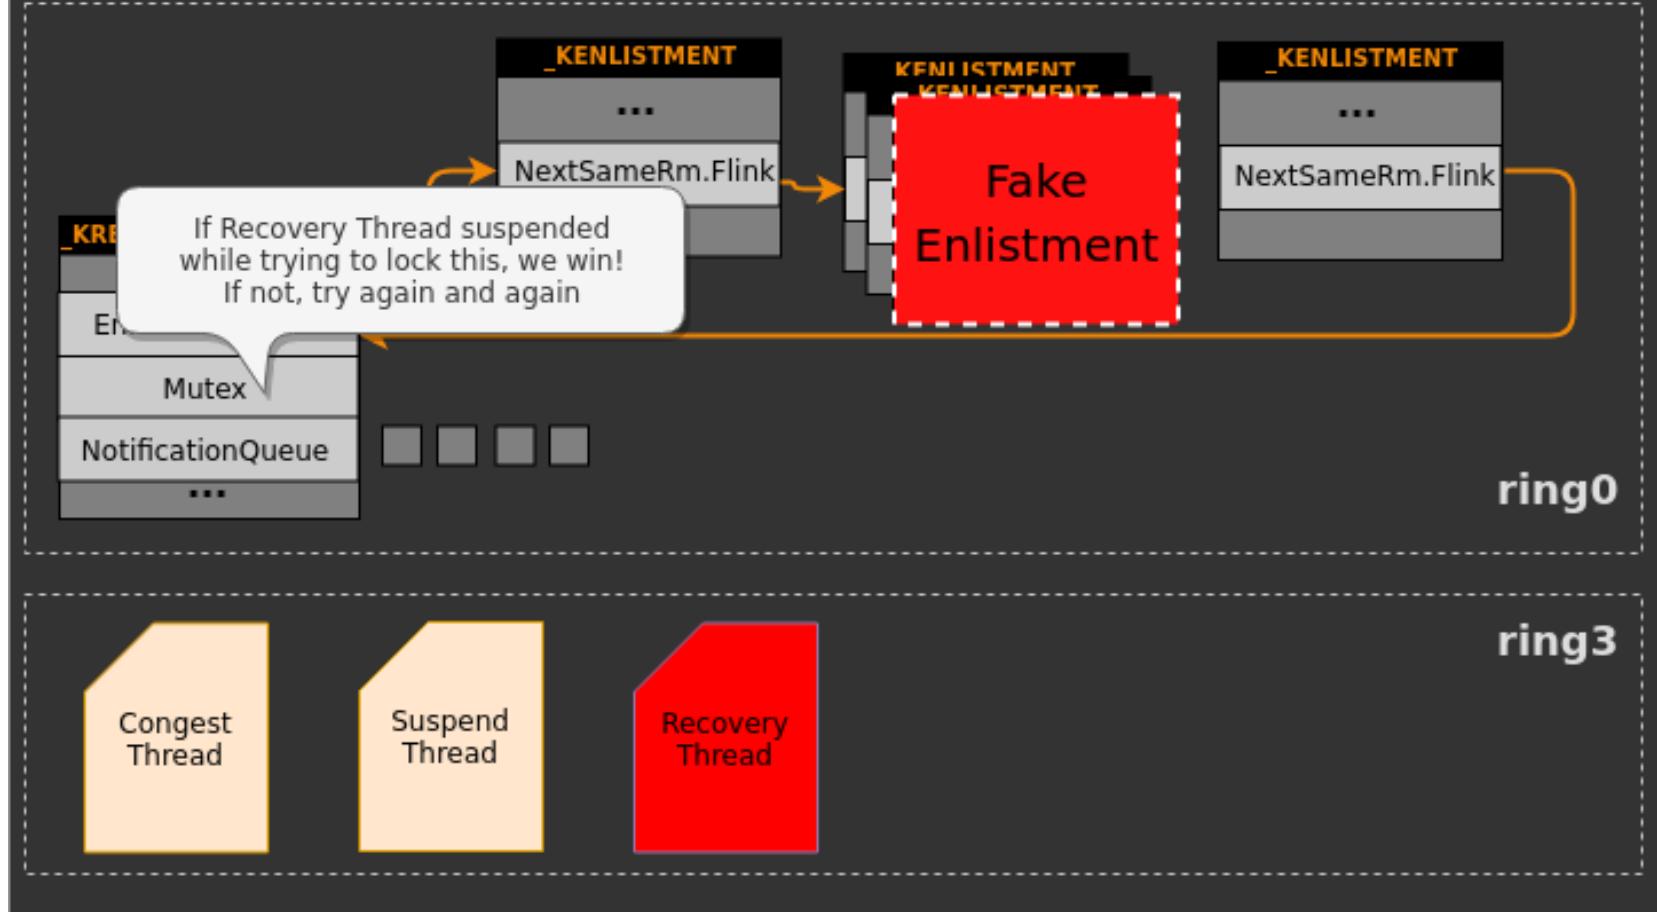
Lock congestion

TmRecoverResourceManager() looping Lock congestion #14



Lock congestion

TmRecoverResourceManager() looping Lock congestion #15



_KENLISTMENT replacement

- We know everything is on the non-paged pool
- We know the size of the _KENLISTMENT
- Non-paged pool feng shui is the obvious approach

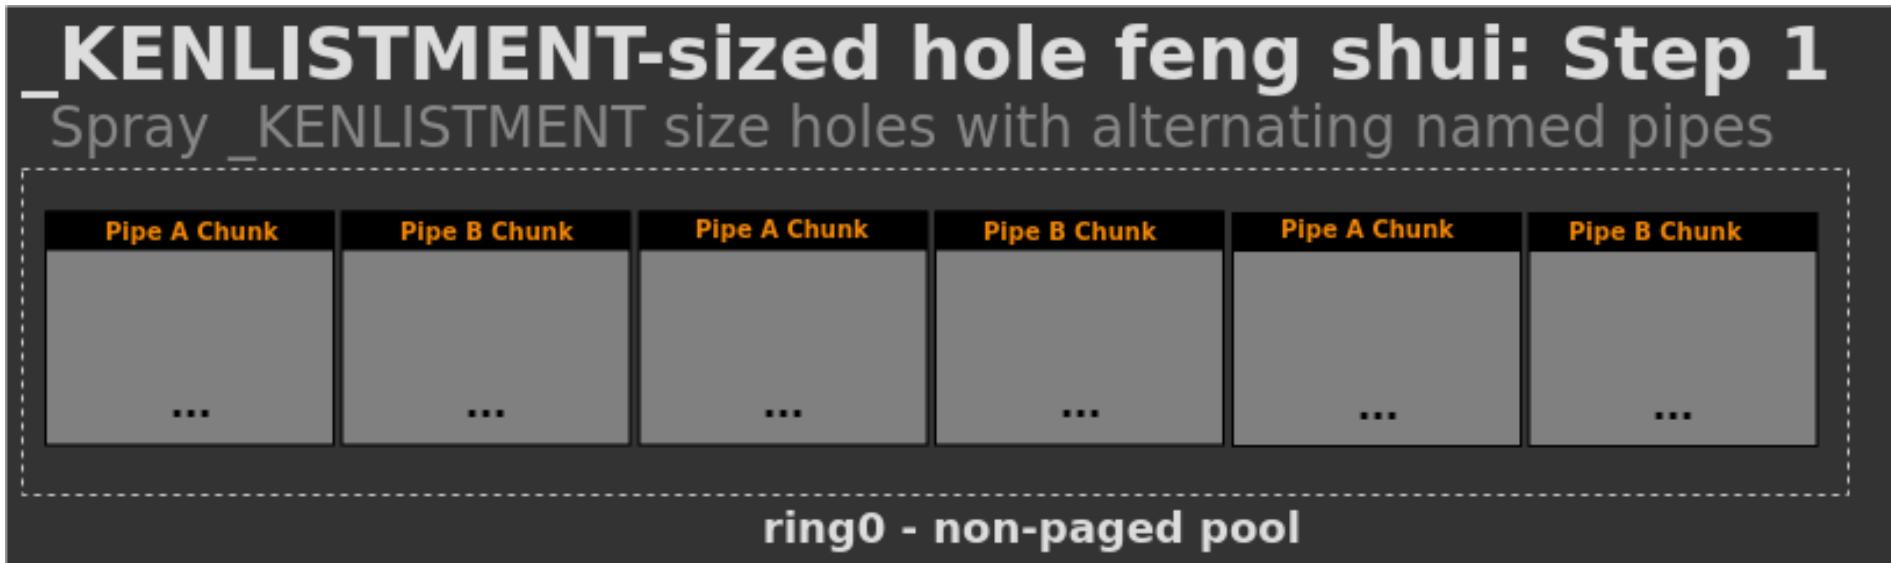
Non-Paged pool feng shui

- Widely known, not too widely shared?
 - Alex Ionescu in [2014](#)
 - Andreas Fobian from Blue Frost Security in [2020](#)

Non-Paged pool feng shui

- Named Pipe writes allocate on non-paged pool
 - Code handled by npfs.sys (tracked by NpFr pool tag)
 - Persistent until other end of pipe reads data
 - Chunk free occurs when data is read
 - Size of chunk is fully controlled
 - All data of chunk aside from DATA_ENTRY is fully controlled
 - Data prefixed with an undocumented DATA_ENTRY structure (changed among Windows versions)
 - [ReactOS](#) is best starting point (reversing/hexdump for relevant changes)

Feng shui layout #1



- As usual, want to avoid coalescing causing big holes
- Writes on alternate named pipes

Feng shui layout #2



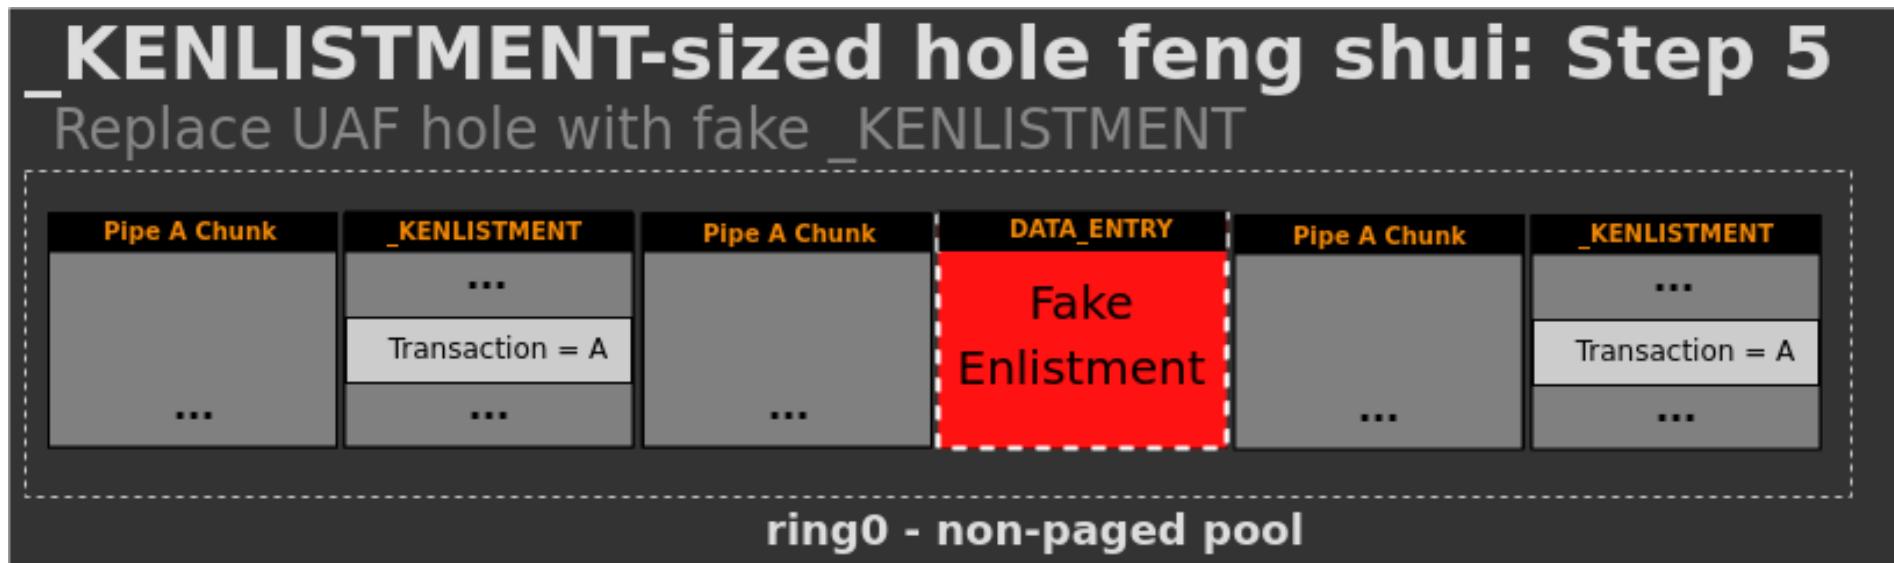
Feng shui layout #3



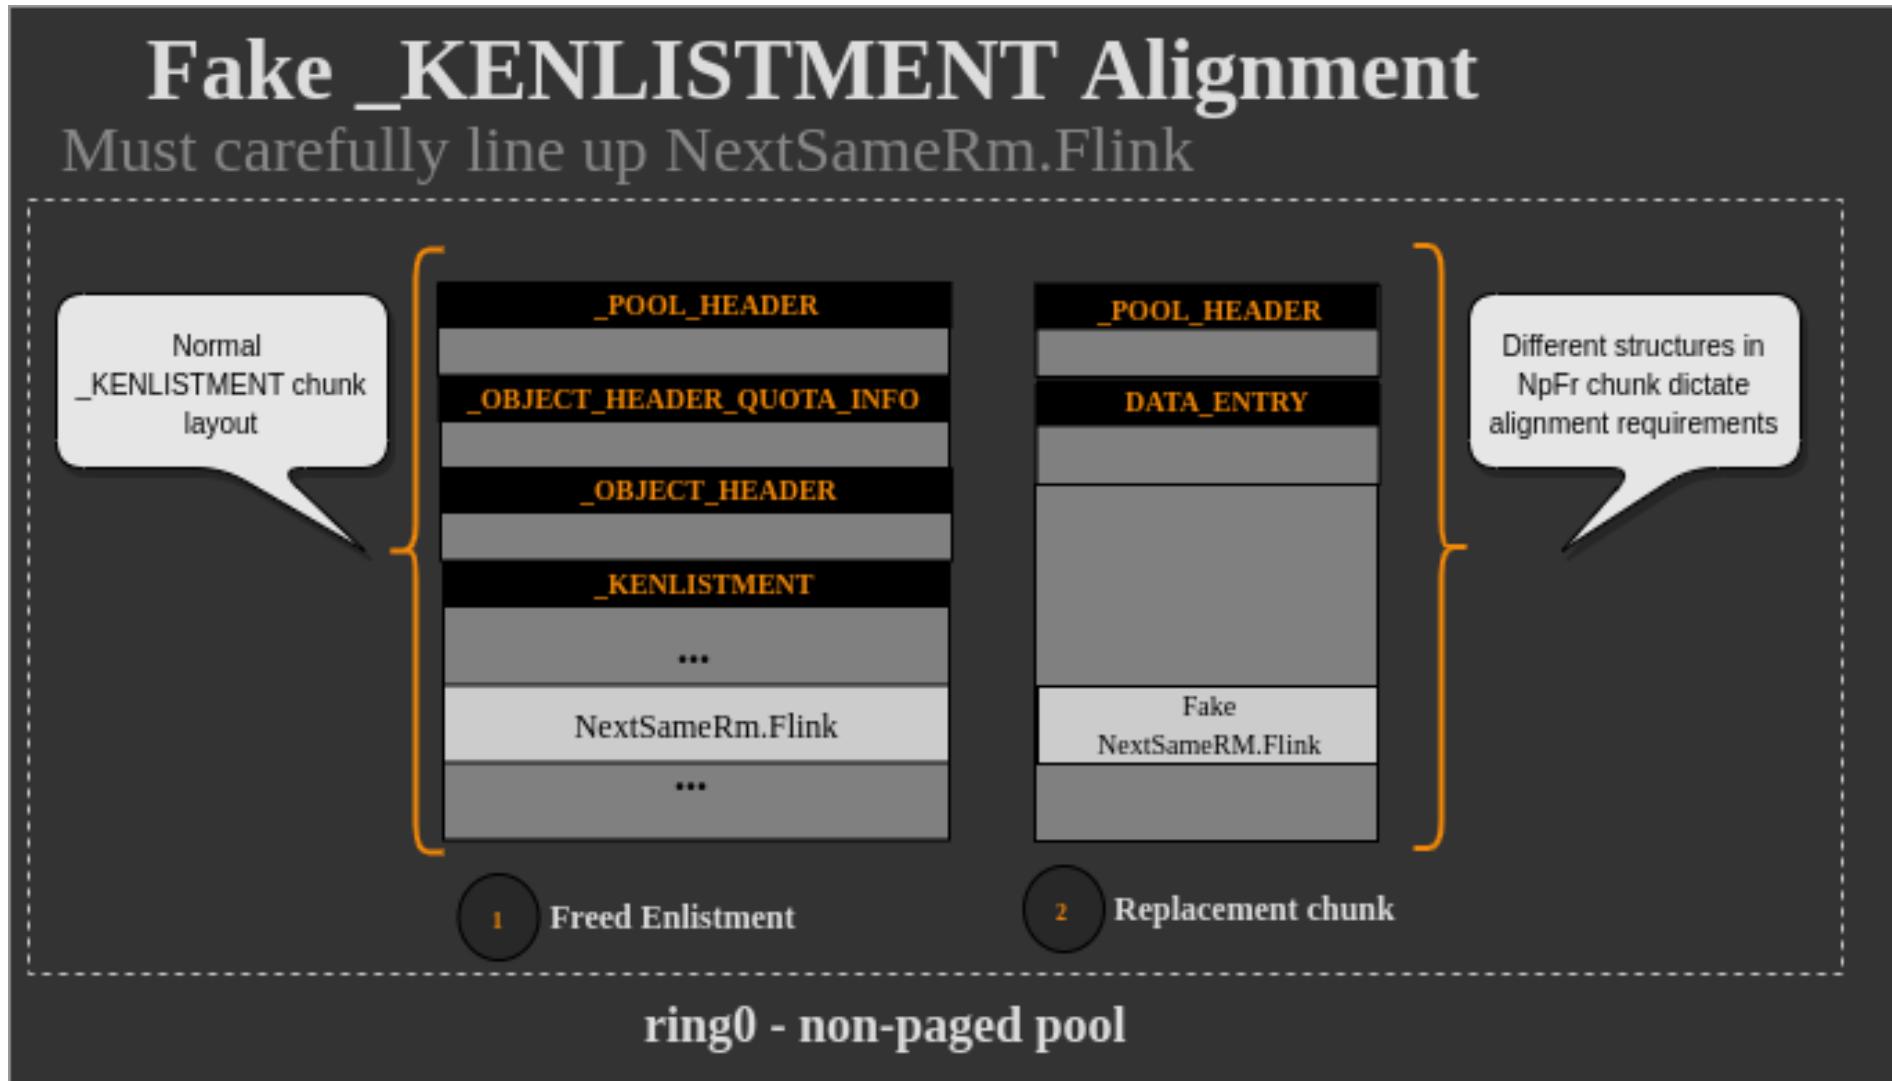
Feng shui layout #4



Feng shui layout #5



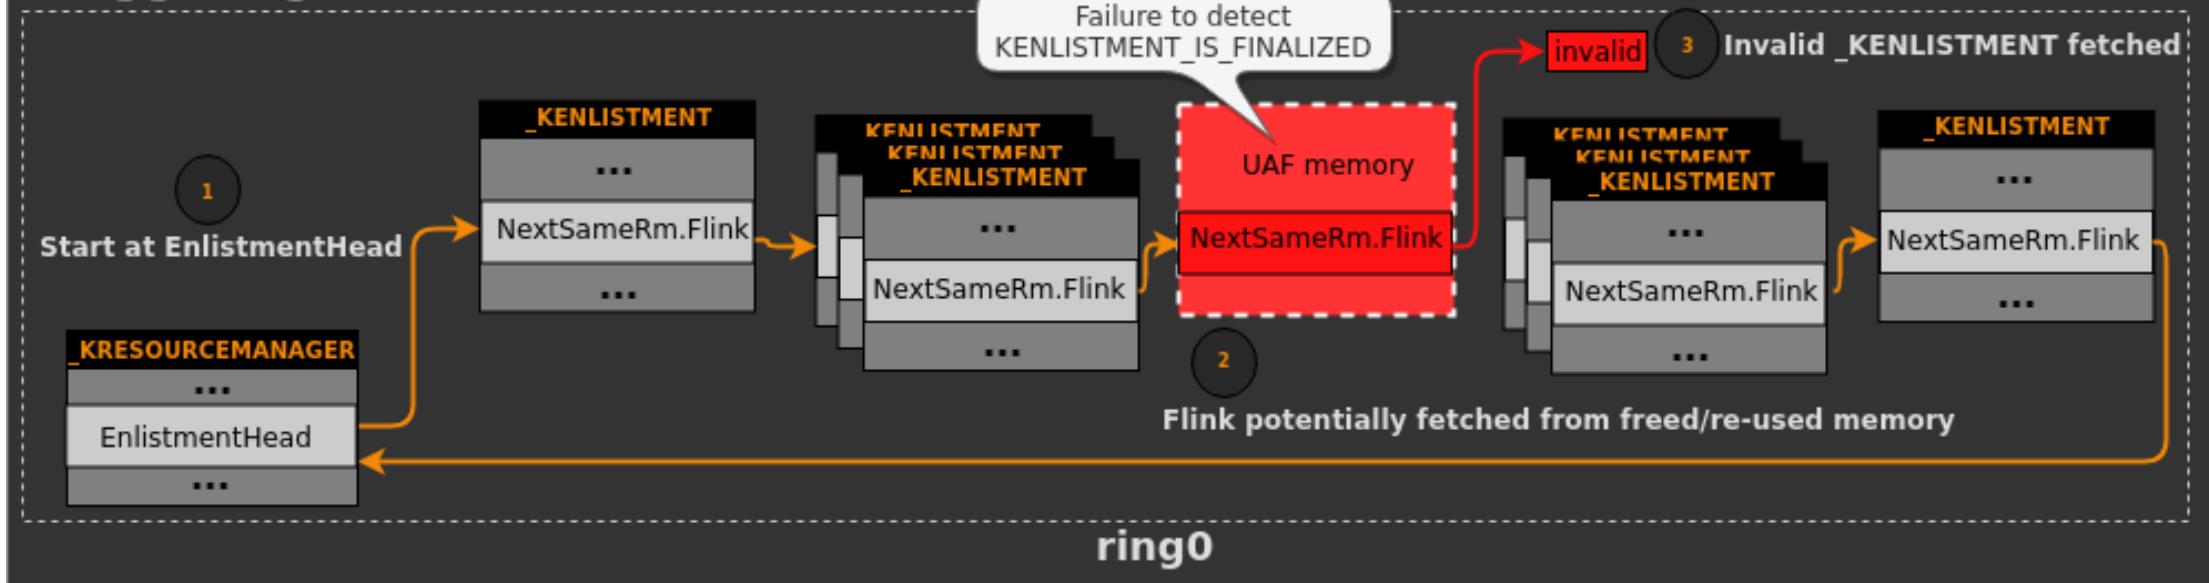
Faking a _KENLISTMENT with a named pipe chunk



End result

TmRecoverResourceManager() looping

Triggering the vulnerable condition



What about delayed free list?

- Allocator optimisation to free chunks when the delayed free list is full only
 - Since Windows Vista
 - Waits for it to contain 32 chunks
 - Delayed free list is flushed and all 32 chunks are actually freed

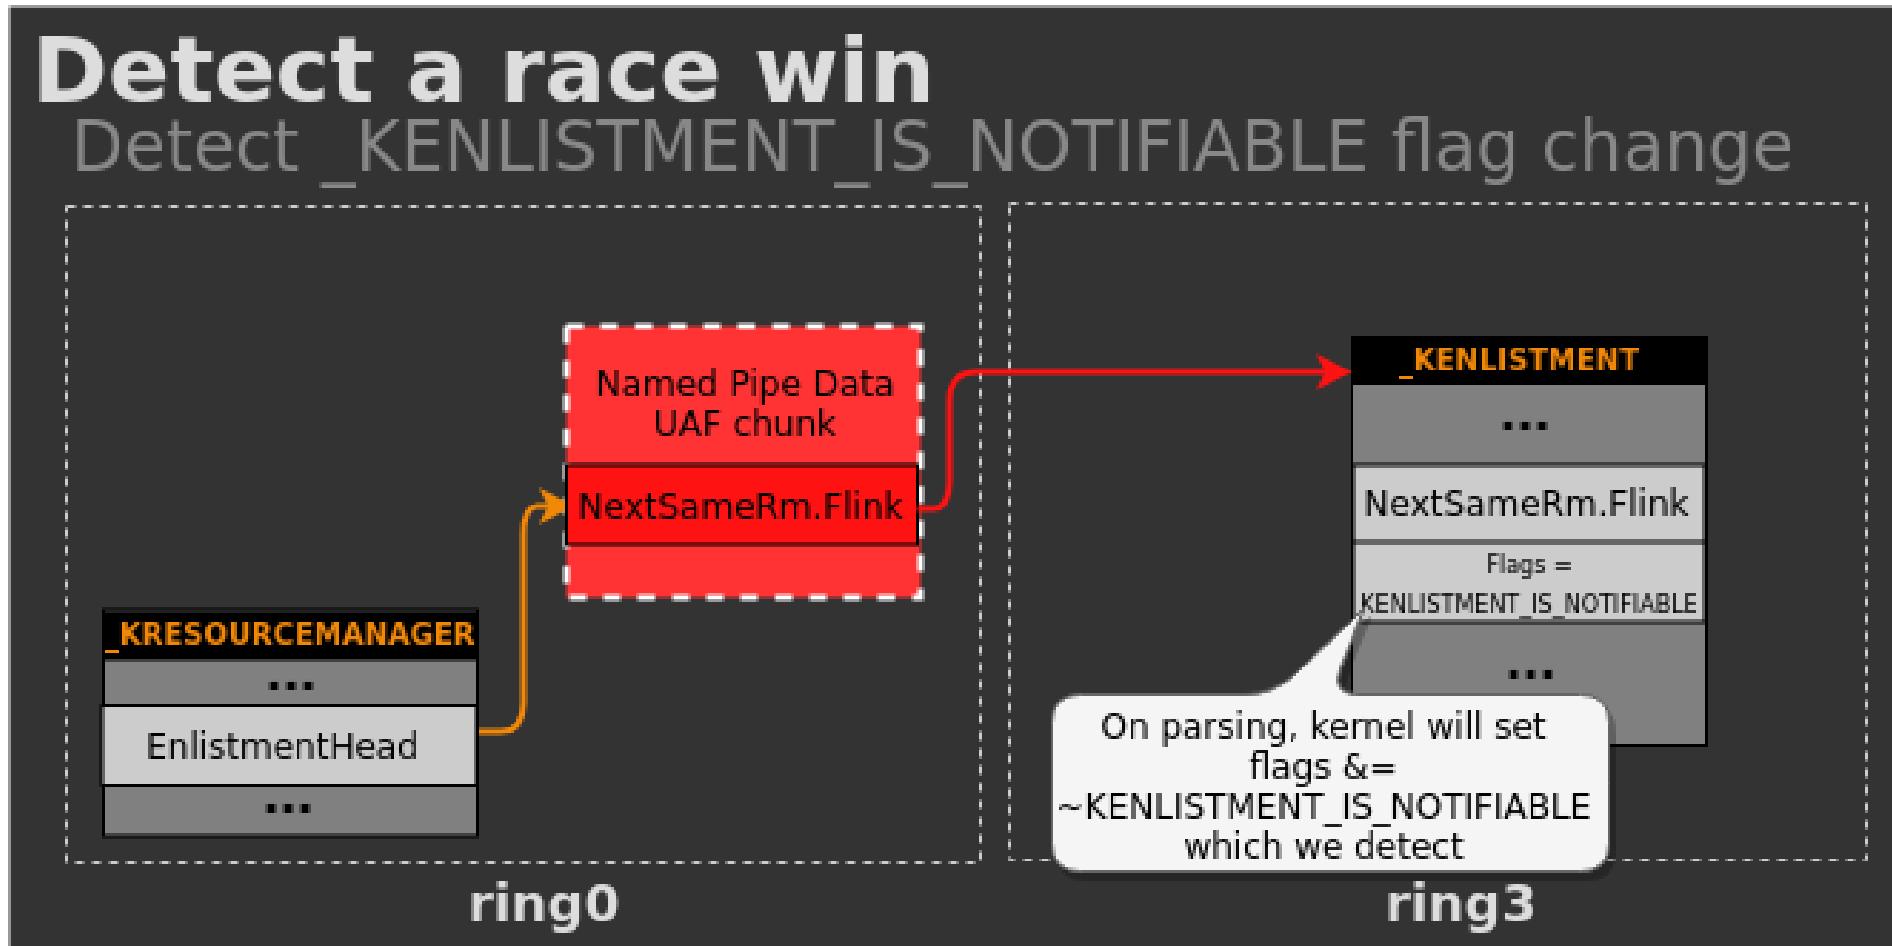
What about delayed free list?

- Turns out it does not matter for us since
 - If we win the race but UAF enlistment chunk is not effectively freed
 - Chunk content will not be replaced by anything else so NextSameRm.Flink used
 - TmRecoverResourceManager loop will just continue with following enlistment in list
 - We can win the race again later
- Can still improve triggering the UAF fast by just freeing 32 chunks to bypass delayed free

Detecting a race win

- How seize control of loop?
- No SMAP on Windows!

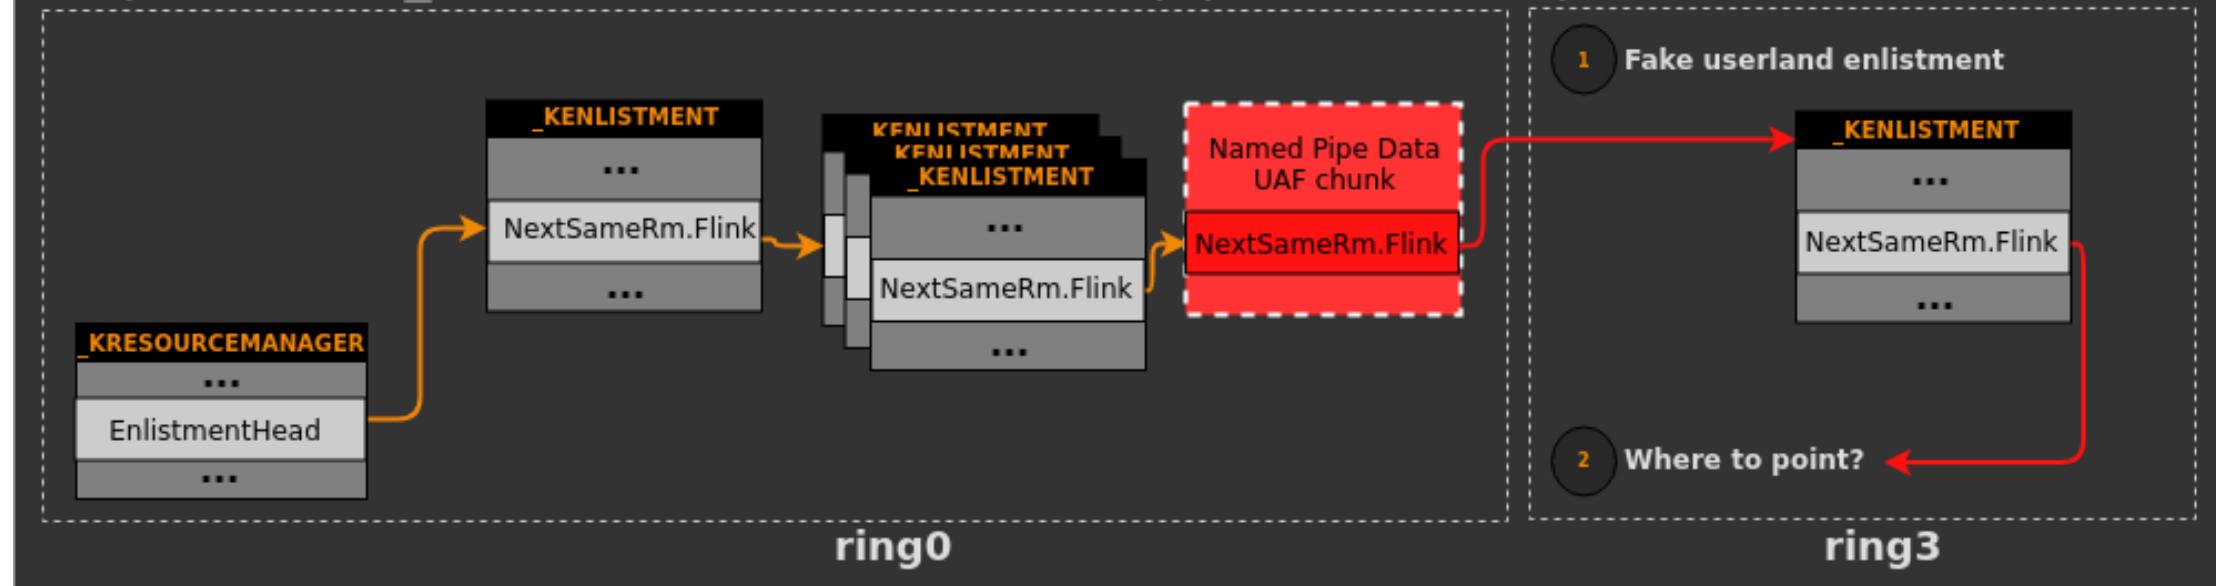
Detecting a race win



Now what?

Post-trigger Enlistment parsing

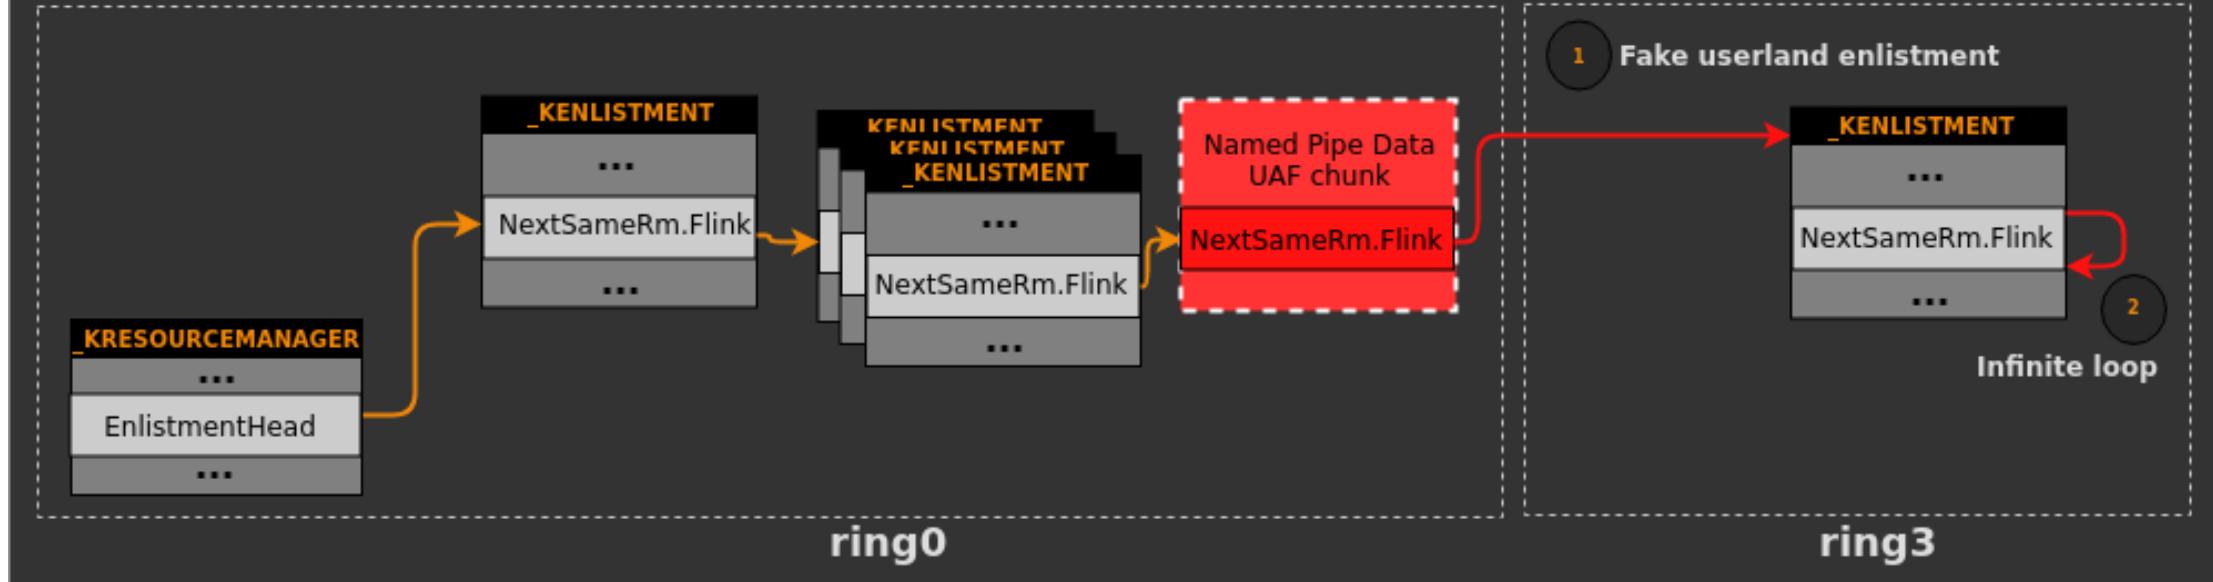
Replace free _KENLISTMENT with named pipe data, and point to userland



Trap enlistment

Trap Enlistment

Trap the TmRecoverResourceManager() inside its while loop

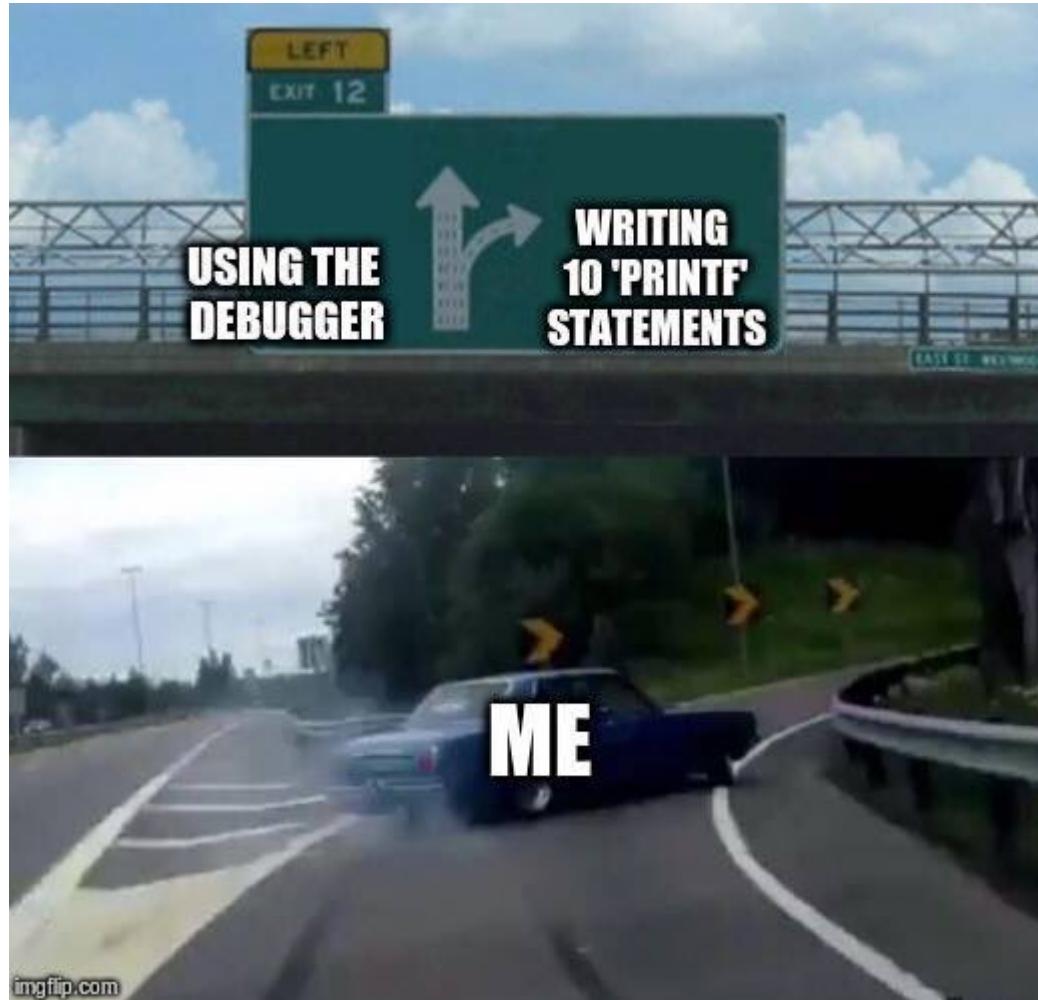


- Inject list of new enlistments into Flink when ready
- Tail of new list of enlistments can be another trap

Detecting a race win - key points

- No SMAP on Windows!
- Replacement _KENLISTMENT->NextSameRM points to yet another fake userland _KENLISTMENT
- Userland _KENLISTMENT->NextSameRM points to itself
- We refer to this as a 'trap' enlistment
- Kernel is now temporarily stuck in an infinite loop
- Kernel unsets KENLISTMENT_IS_NOTIFIABLE flag on userland enlistment
 - This modification in userland tells us we won!

Debugging a race win?

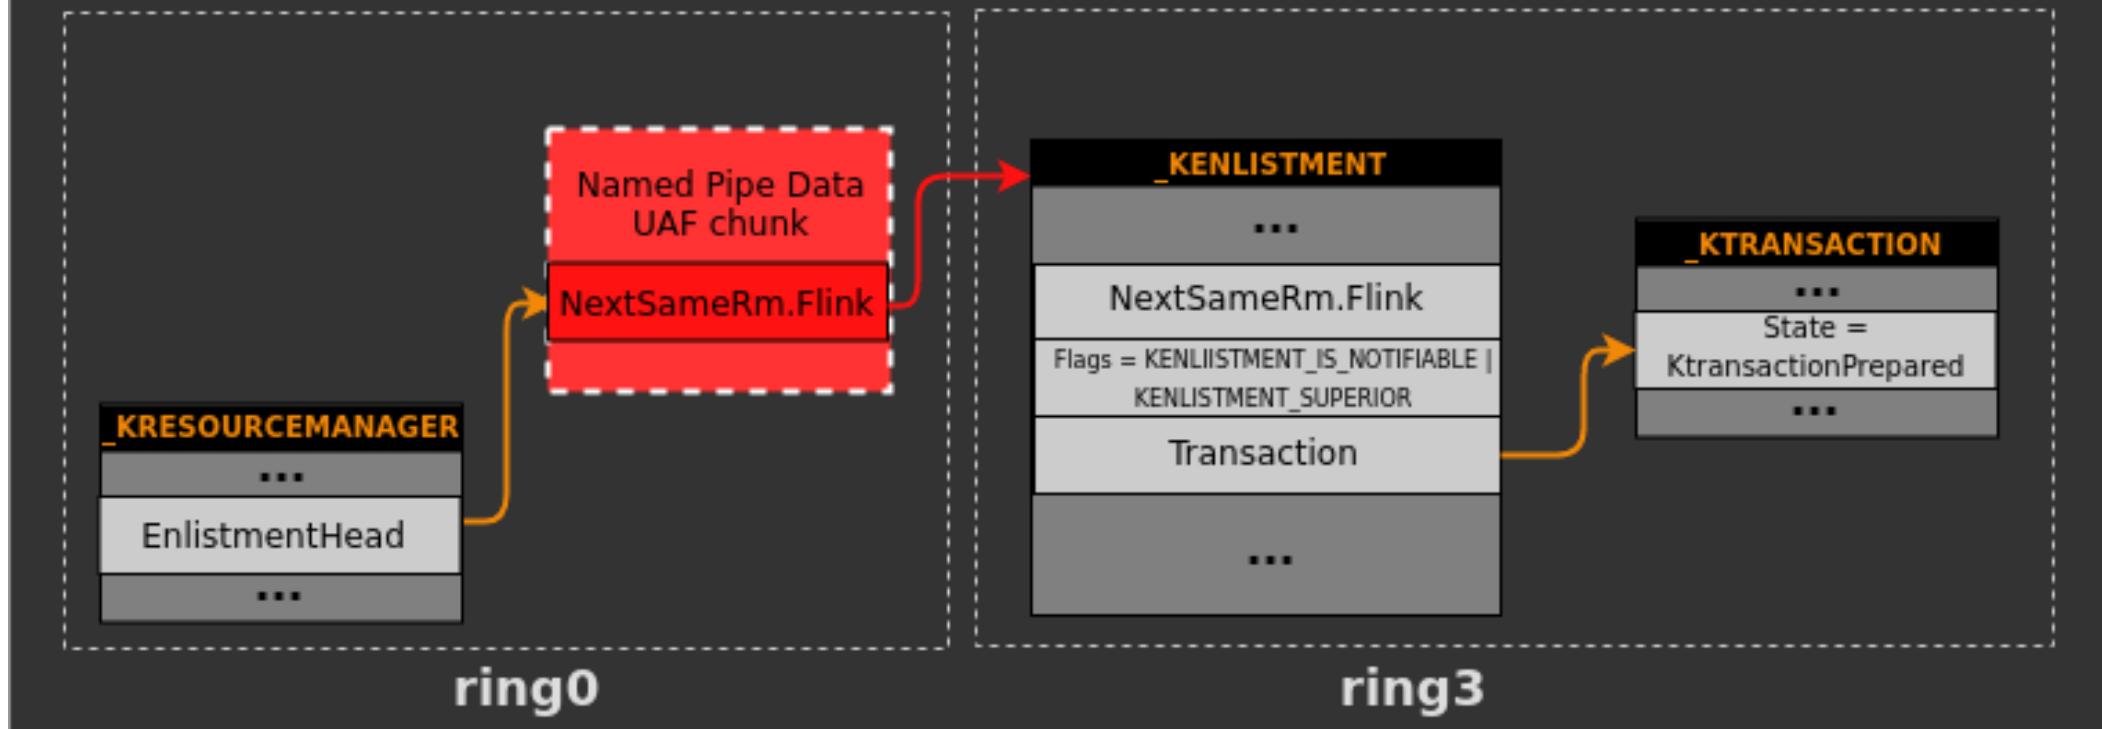


imgflip.com

Debugging a race win?

Trigger kernel debugger

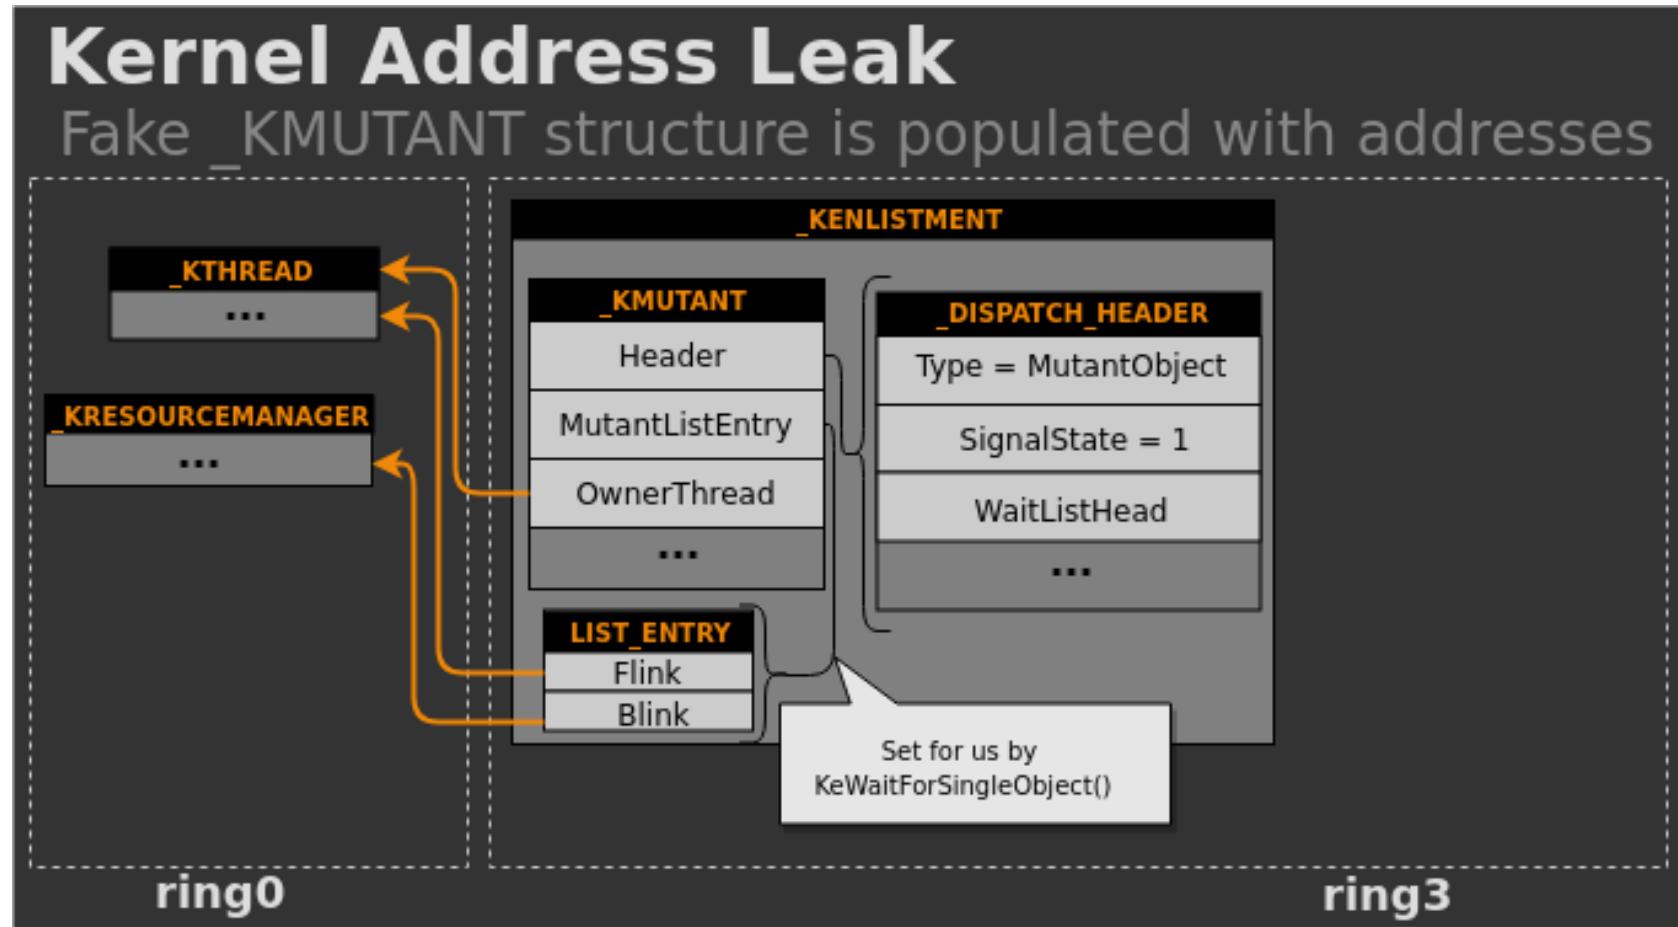
Enlistment causes breakpoint code to run



How to escape the loop?

- We have control of the loop now
- We need a write primitive of some kind
- But also need to escape the loop?

Initial kernel pointer leak



- Thank you KeWaitForSingleObject()

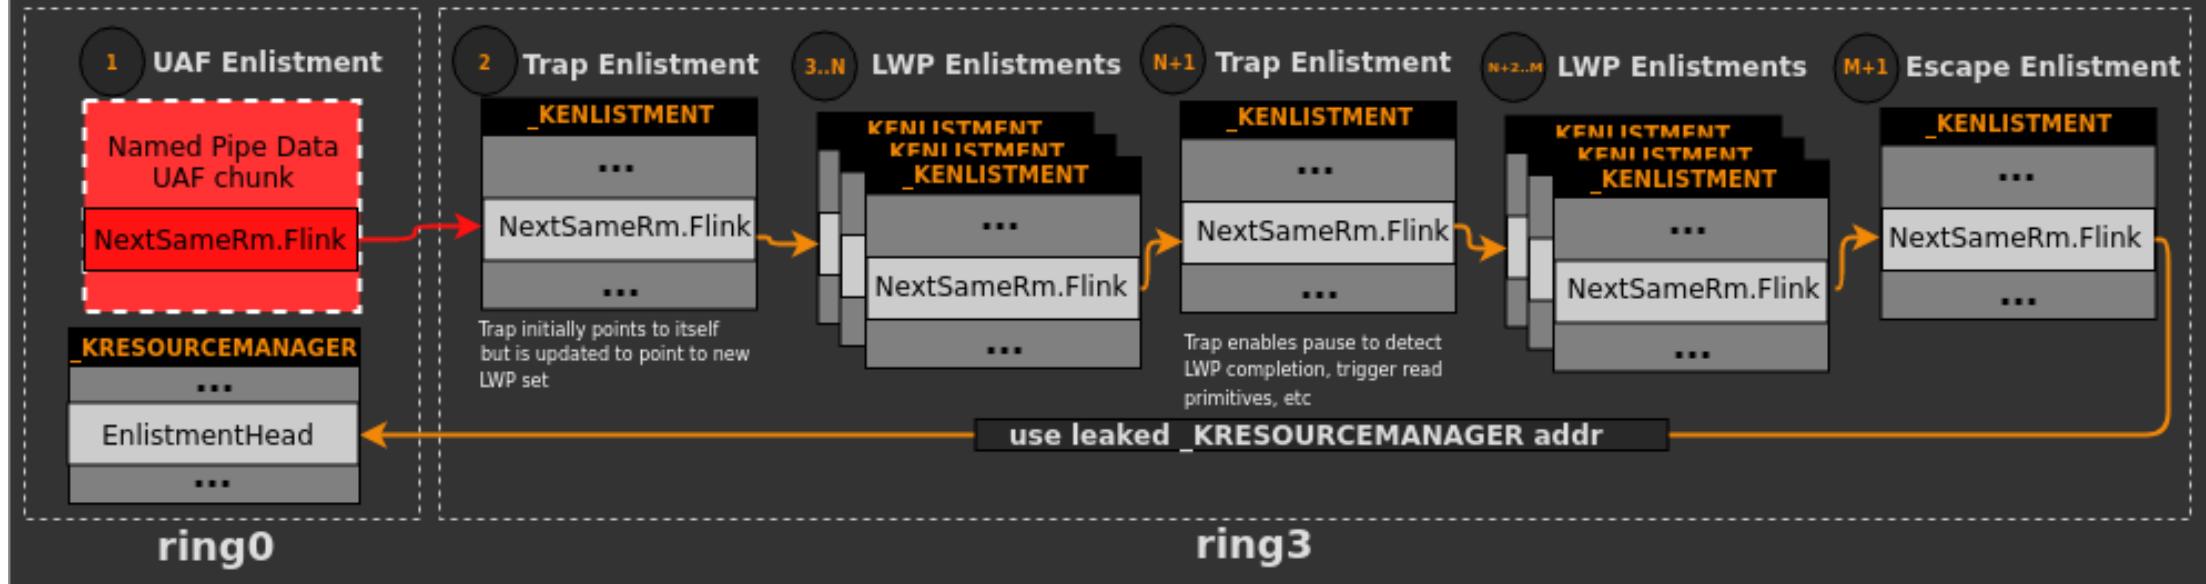
Escaping the loop

- We can now exit the loop!
- Introduce an 'escape' enlistment
- Set KENLISTMENT->NextSameRm = &_KRESOURCEMANAGER.EnlistmentHead
- Exit cleanly
- No crashes.. reproducible testing, etc.

What an escape looks like

Exploit enlistment parsing flow

UAF -> (Trap -> LWP)sxN -> Escape



- LWP = Limited Write Primitive (explained soon)

Building a write primitive

Vulnerable loop constraints

- Finding a write primitive is somewhat limited
- We are stuck inside this recovery loop
- What code paths do we follow?
- KeReleaseMutex() seemed best
 - List-based mirror-write primitives are safe unlinked after Windows 7 :(
 - Keep looking...
- Found an arbitrary increment inside KiTryUnwaitThread() call

```
if ( (OwnerThread->WaitRegister.Flags & 3) == 1 ) {  
    ThreadQueue = OwnerThread->Queue;  
    if ( ThreadQueue )  
        _InterlockedAdd(&ThreadQueue->CurrentCount, 1u);
```

- But things get complicated..

Arbitrary increment primitive

- KeReleaseMutex() - KeReleaseMutant() wrapper
 - KeReleaseMutant() - Our high level primitive function
 - KiTryUnwaitThread() - Gives us our increment primitive
 - KiProcessThreadWaitList() - Unavoidable because of increment primitive
 - KiUnlinkWaitBlocks() - Have to satisfy its attempt to unlink
 - KiReadyThread() - Unavoidable call on our fake thread
 - KiRequestProcessInSwap() - Have to satisfy early exit

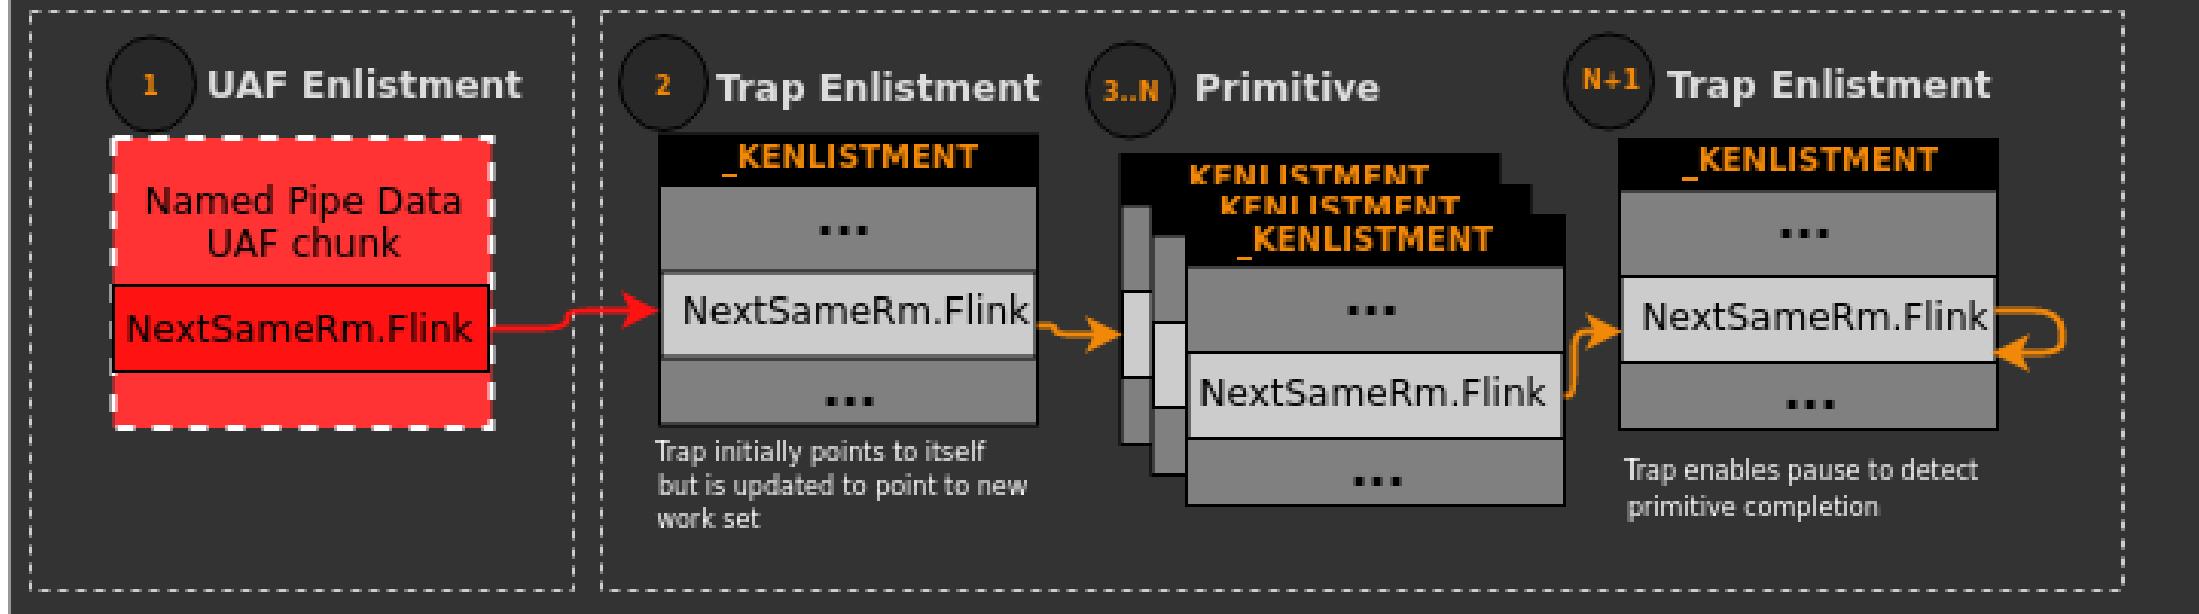
Repeatable arbitrary address increment

- Too complicated to explain in detail
- Follow up blog series covers line by line
- Positives
 - Can chain multiple increments together
 - Effectively an arbitrary write primitive
- Negatives
 - Need to know the starting contents of the address being written to
 - Some risks related to running at DISPATCH_LEVEL

Primitive injection at a glance

Detect primitive completion with traps

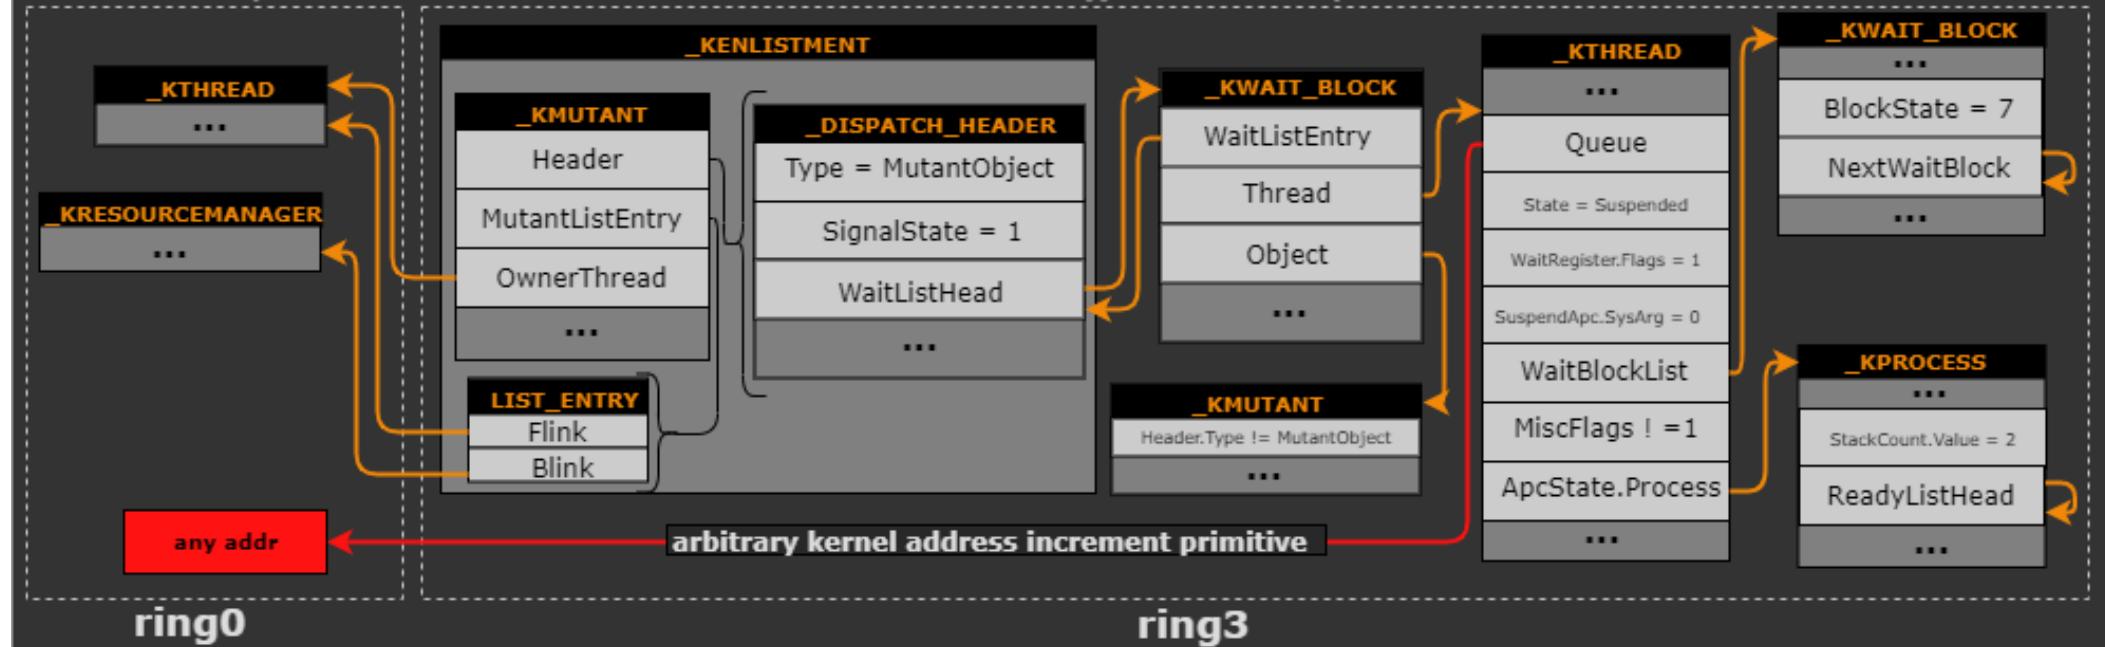
UAF -> Trap -> Primitive -> Trap



What does our increment primitive look like?

Limited Write Primitive (LWP) Enlistment

Arbitrary increment via KeReleaseMutex() - Prerequisite values



- Lots of constraints
- Some requirements change across OS versions

Me running the LWP primitive

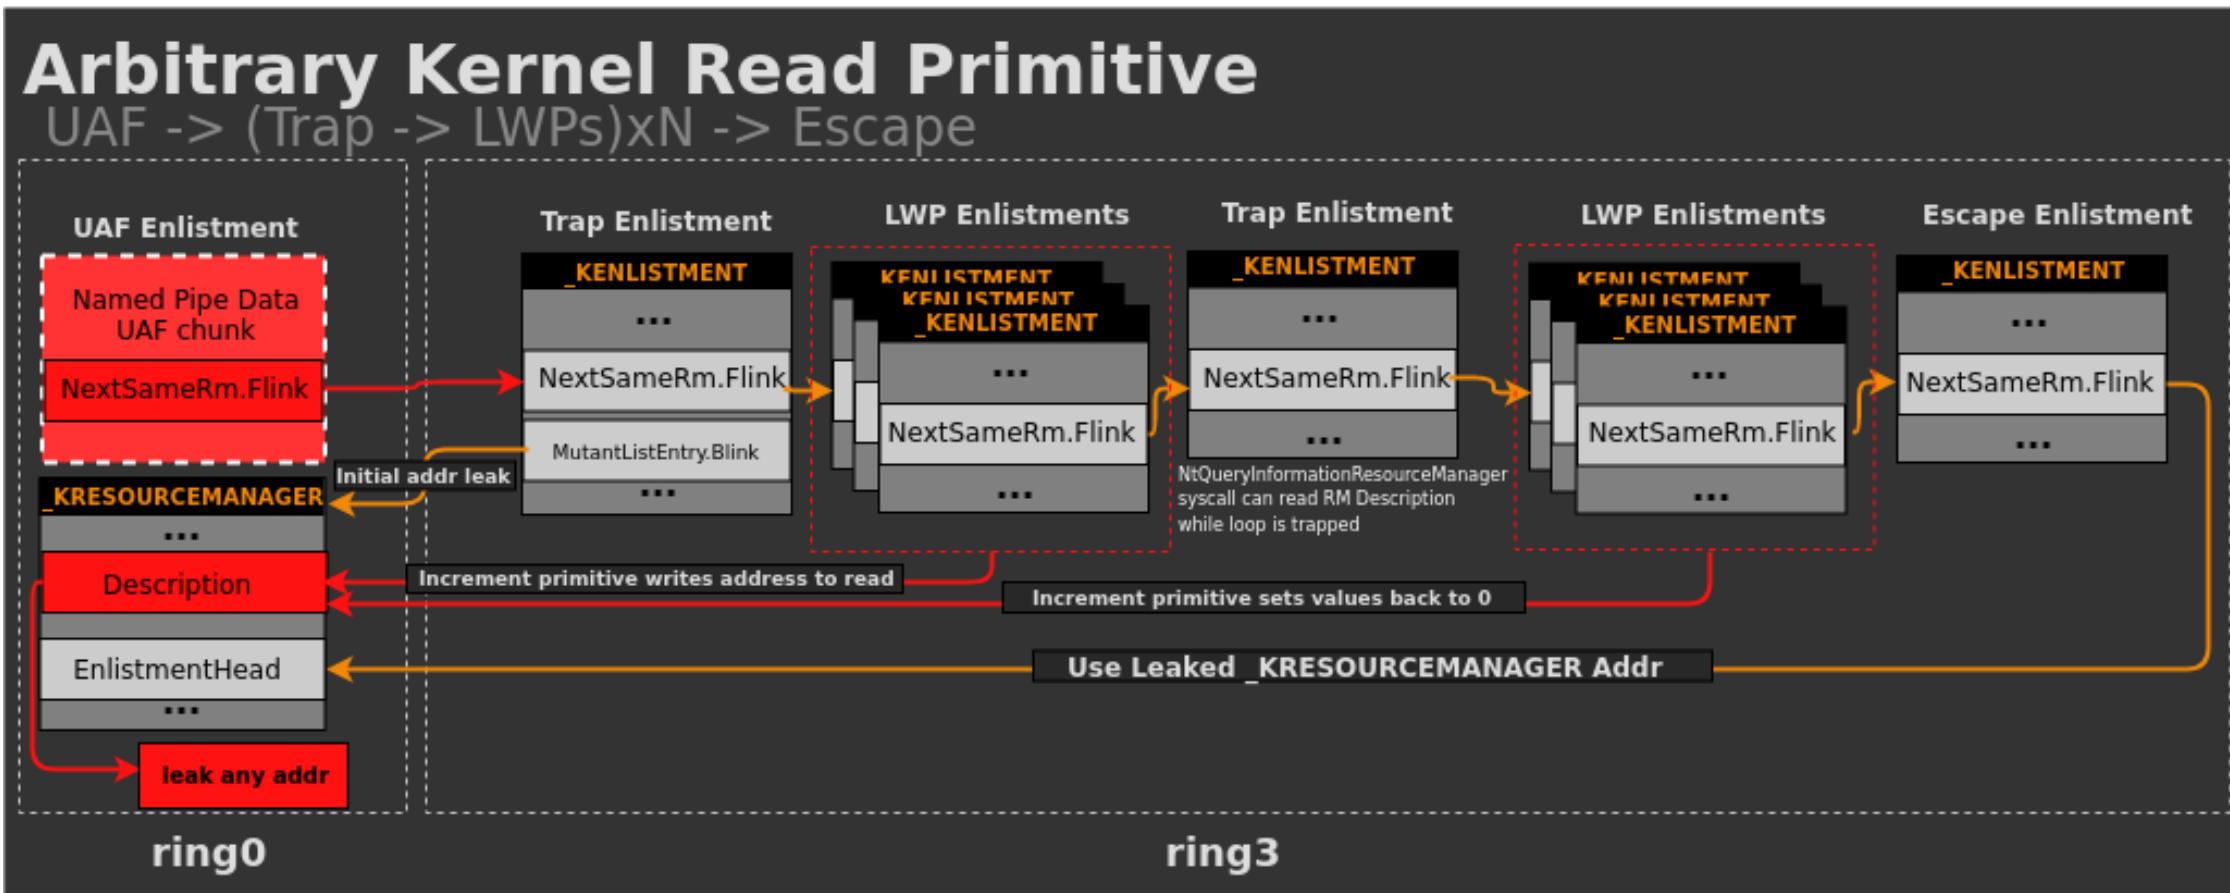


Building an arbitrary kernel read primitive

What to do?

- We have an arbitrary write as long as we know original value
- We know where _KRESOURCEMANAGER is
- We can avoid setting a Description field when creating the RM
 - Means we know _UNICODE_STRING Length and Name are both zero
- Point anywhere we want
- Call NtQueryInformationResourceManager syscall to get description
- Rinse and repeat

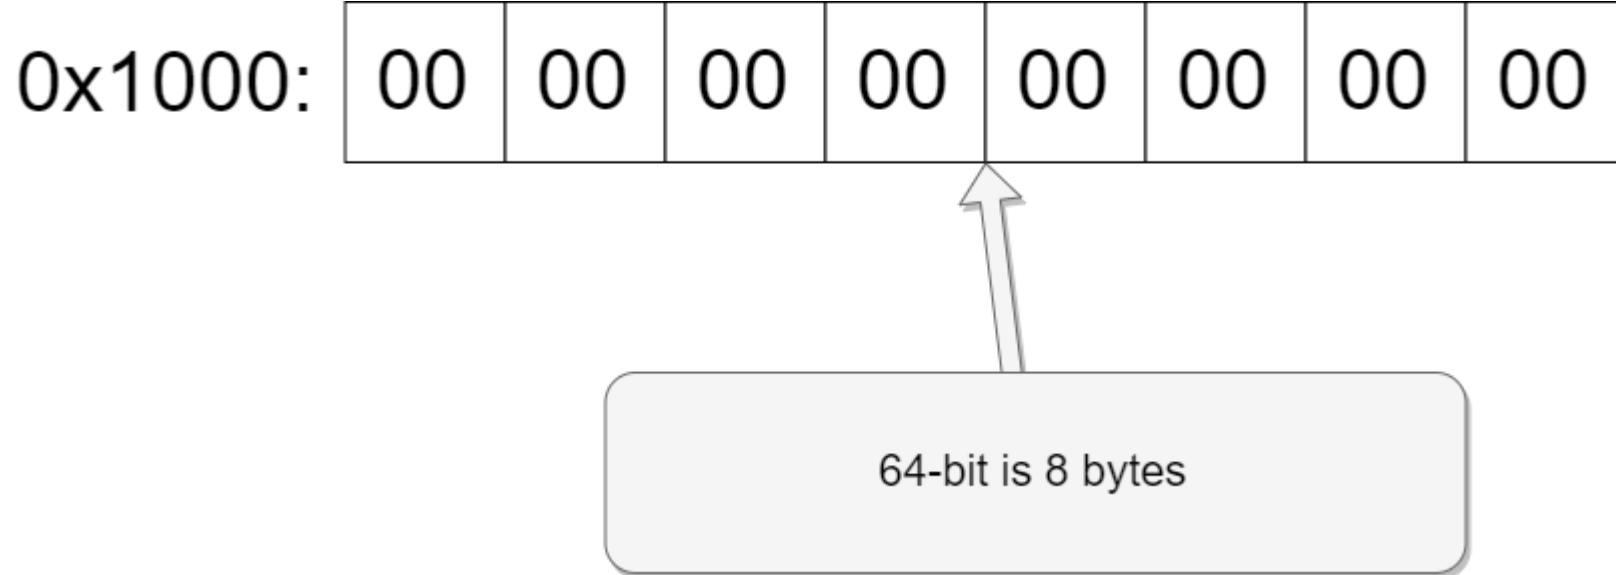
What does our read primitive look like?



How is that possible in practice?

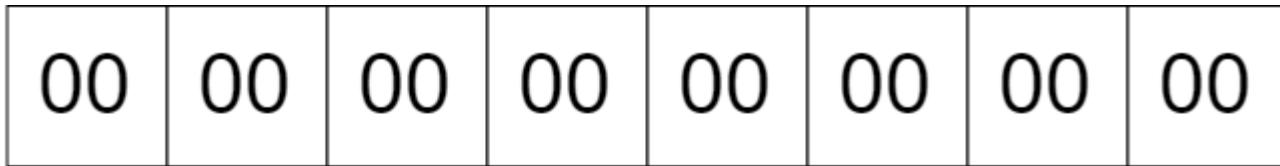
- What if we need to get a really big value?
- Or need to wrap around to get a smaller value?
- Are we going to need 2^{64} increments?
 - Fortunately no, there is a trick ($\text{max } 2*(256*8) = 2048$ writes)
 - Split the increments across 8 different addresses
 - Each byte of the 64-bit value incremented 256 or less times

Efficient use of increment



Efficient use of increment

0x1000:



Let's say we want to write 0x1234
without doing 0x1234 increments

Efficient use of increment

0x1000:

34	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----



We increment 0x34 times
at address 0x1000

Efficient use of increment

0x1000:

34	12	00	00	00	00	00	00
00	00	00	00	00	00	00	00



We increment 0x12 times
at address 0x1001

Efficient use of increment

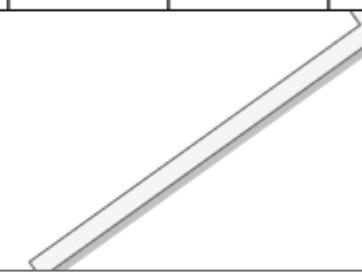
0x1000: 00 00 00 00 00 00 00 40 00



What happens if we want to
then change the MSB from
0x40 to 0x3F?

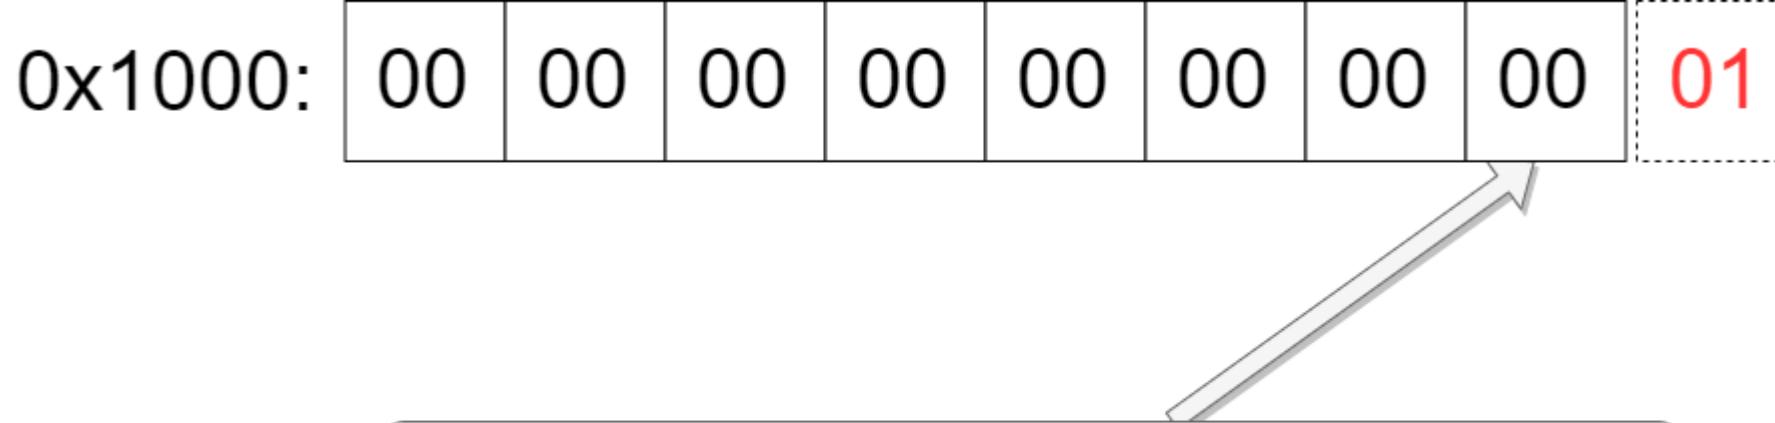
Efficient use of increment

0x1000: 00 00 00 00 00 00 00 FF 00



If we use the increment at 0x1007, and increment 0xC1 times, when we cross from 0xFF to 0x00, the adjacent byte will be incremented by 1

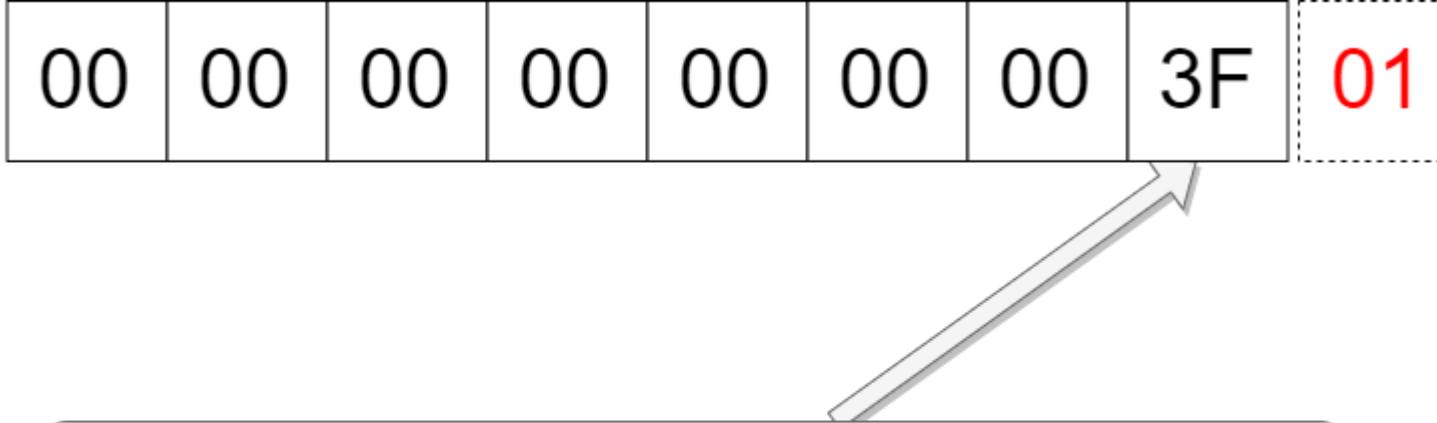
Efficient use of increment



If we use the increment at 0x1007, and increment 0xC1 times,
when we cross from 0xFF to 0x00, the adjacent byte will be
incremented by 1 => corruption

Efficient use of increment

0x1000: 00 00 00 00 00 00 00 3F 01



We reach wanted 0x3F after 0xFF increments,
but the adjacent byte is corrupted

Efficient use of increment

0x1000:

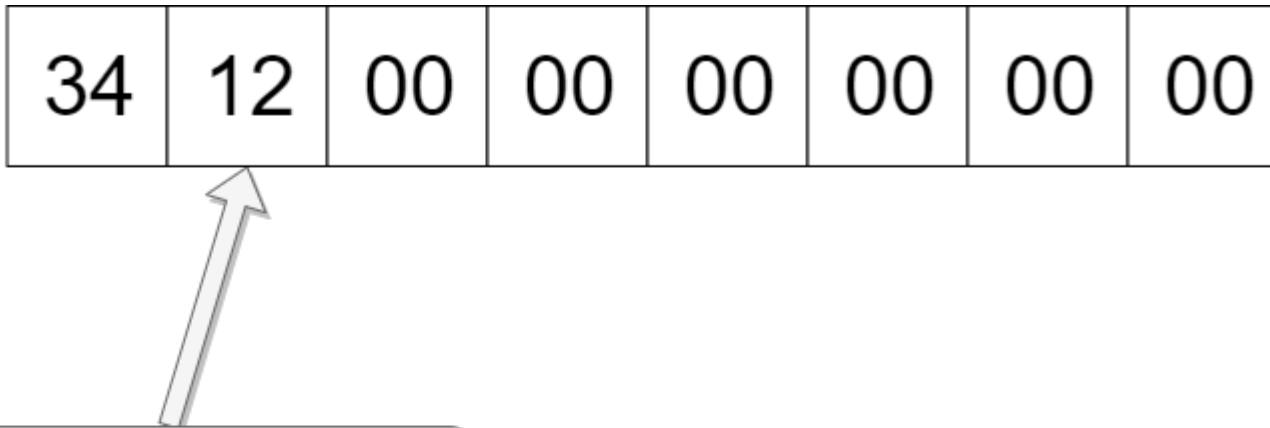
34	12	00	00	00	00	00	00
----	----	----	----	----	----	----	----



To avoid that, we reset the full
64-bit value to zero

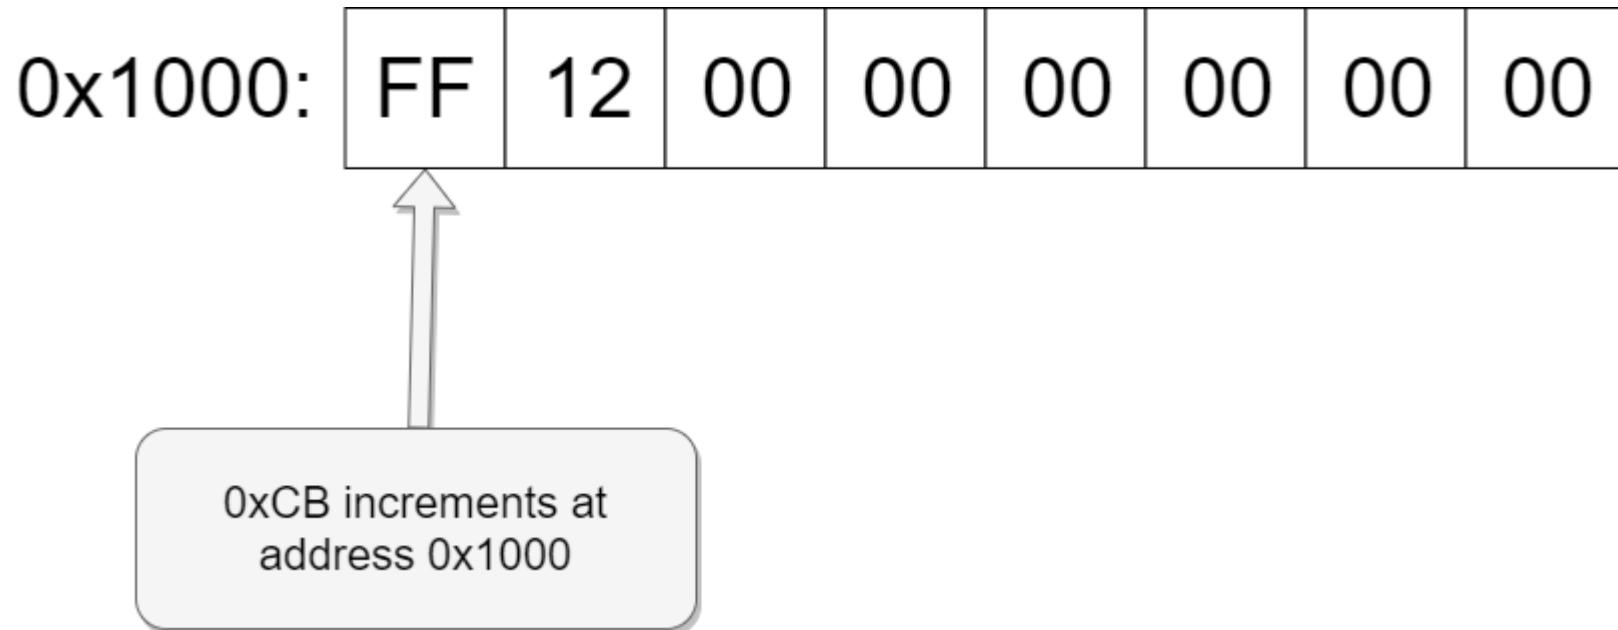
Efficient use of increment

0x1000: 34 12 00 00 00 00 00 00



We first set every byte to FF with misaligned increments

Efficient use of increment



Efficient use of increment

0x1000: FF FF 00 00 00 00 00 00



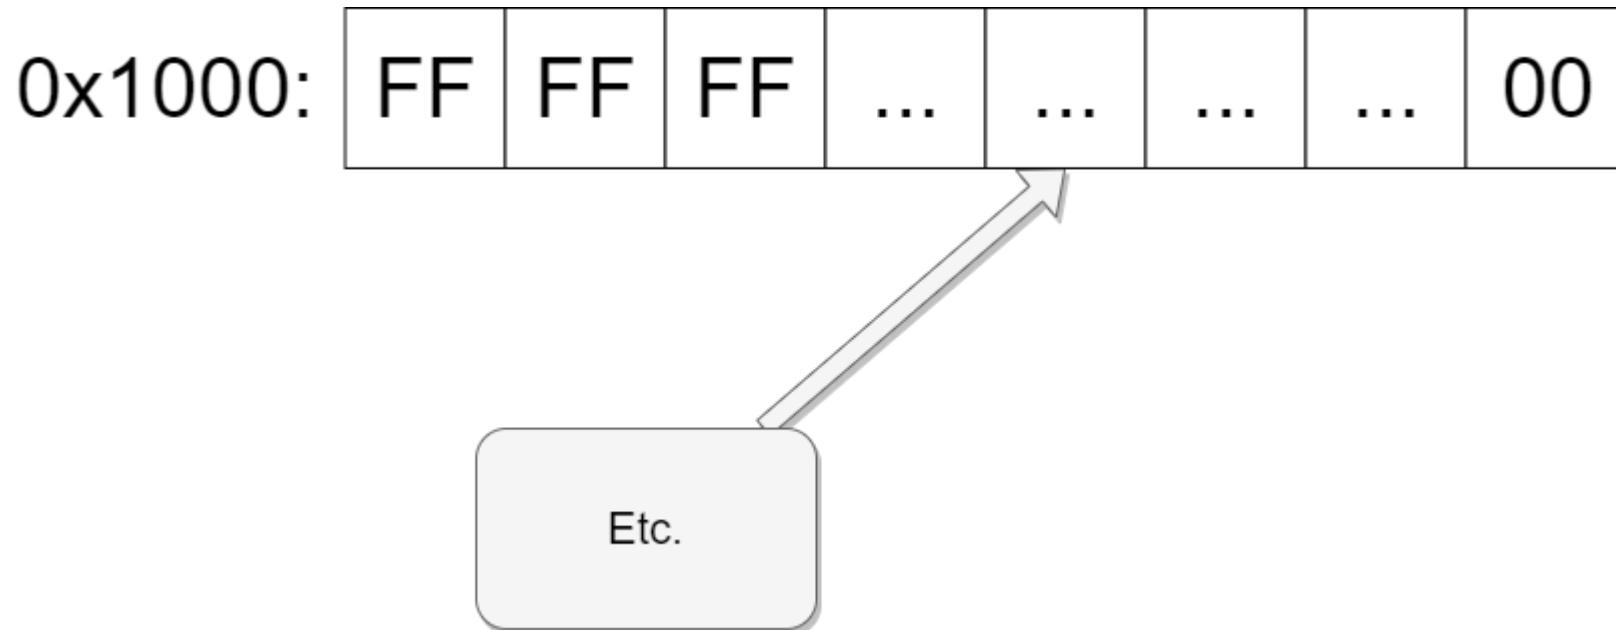
0xED increments at
address 0x1001

Efficient use of increment

0x1000: FF FF FF 00 00 00 00 00

0xFF increments at address 0x1002

Efficient use of increment

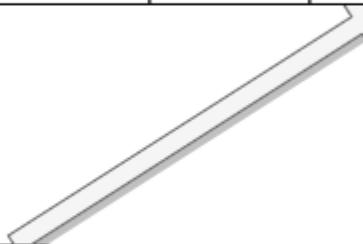


Efficient use of increment

0x1000:

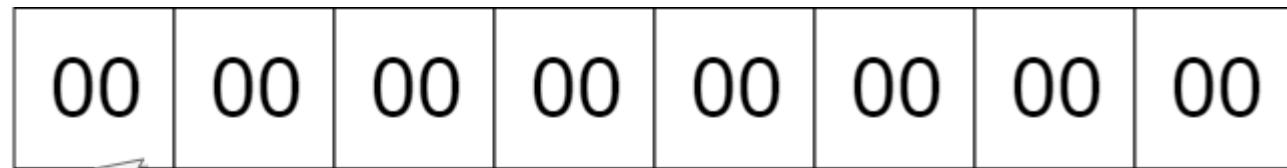


0xFF increments at
address 0x1007



Efficient use of increment

0x1000:



One increment at aligned address 0x1000 with no unwanted side effect

How is that possible in practice?

- Max $2*(256*8) = 2048$ writes
 - < $256*8$ writes to reset value to zero
 - < $256*8$ writes to set value to anything

Privilege escalation

Data only attack - Using the increment primitive

- We can trigger the increment primitive indefinitely
- Use the increment write primitive to enable an arbitrary read primitive
- Use the read primitive to read SYSTEM token
- Use the write primitive to adjust our EPROCESS token to SYSTEM
- Caveats: If EPROCESS token is read during our slow adjustment, we BSOD
 - If Task Manager is running
 - If Process Explorer is running

Exploiting Windows 10 1809 (RS5) x86/x64

- Data only attack
- Bypassing kernel CFG wasn't investigated (yet)
 - But primitives should make it doable
- Only major x64 and x86 differences is structure sizes and offset
 - Except for the following thing to come...
- Relatively easy to port to all versions back to Vista

Bonus - BlueHat Shanghai May 2019

Bonus - The "invisible" paper

- Turns out Kaspersky presented on this in May 2019 at [BlueHat Shanghai 2019](#)
 - Boris Larin (@oct0xor) and Anton Ivanov (@antonivanovm)
 - Explains some of what we just described
- Found after we got accepted to speak at POC2019
 - win32k syscall filter search keywords found it by accident
 - Searching CVE-2018-8611 or KTM did not

Me reading the paper



Bonus - The "invisible" paper

- Actually quite happy in the end we never saw it!
- Most interesting highlight
 - 0day exploit used multiple different approaches from us

Bonus - Race winning detection

- 0day didn't use same trap enlistment approach to detect race win
- Used Event Notification object to trap kernel on KeWaitForSingleObject()
 - Swap object type after detection
 - Modified mutex allows write 0 primitive (similar code path to ours)
 - Positives
 - It's interesting to see a different approach
 - Negatives
 - Must modify every mutex that gets touched by loop
 - More complicated than our primitive

Bonus - Write primitive: No increment, write 0 only

- 0day didn't use the increment primitive either!
- Abused an earlier write 0 in same KeReleaseMutex() code path
 - Writes a sizeof(void *) 0 value to any address
 - Least significant bit must already be 0 to avoid deadlock
 - Positives
 - Reduced setup complexity to reach write primitive
 - Negatives
 - Doesn't work on all OS versions
 - Vista x64 due to code differences
 - Vista/7 x86 because whole pointer required to be 0 in the first place (different macro)
 - Situationally less powerful primitive

Bonus - What to write with 0?

- 0day targeted KTHREAD.PreviousMode field
 - First documented by Tarjei Mandt in 2011
 - Misaligned write to this field allows setting to 0
 - Unrestricted NtReadVirtualMemory() and NtWriteVirtualMemory().
 - Arbitrary kernel read/write
 - Positives:
 - Super powerful (64-bit)
 - Possibly first in-the-wild use?
 - Negatives:
 - Doesn't work on x86 due to PreviousMode being reset to UserMode based on CS when syscall entered

Demo

Conclusion

- Quite reliably exploitable race condition leading to UAF
- Very interesting and fun to exploit
- Should be usable to bypass most kernel mitigations (if necessary)
 - KASLR, SMEP, CFG, etc.
- Our approach differed significantly from 0day
 - Both methods have a lot of value!
- Tons of details still missing
 - Follow up 5 part blog series coming

Questions?

- Cedric Halbronn - [@saidelike](https://twitter.com/saidelike), cedric.halbronn@nccgroup.com